

ALPS:

An Adaptive Learning, Priority OS Scheduler for Serverless Functions

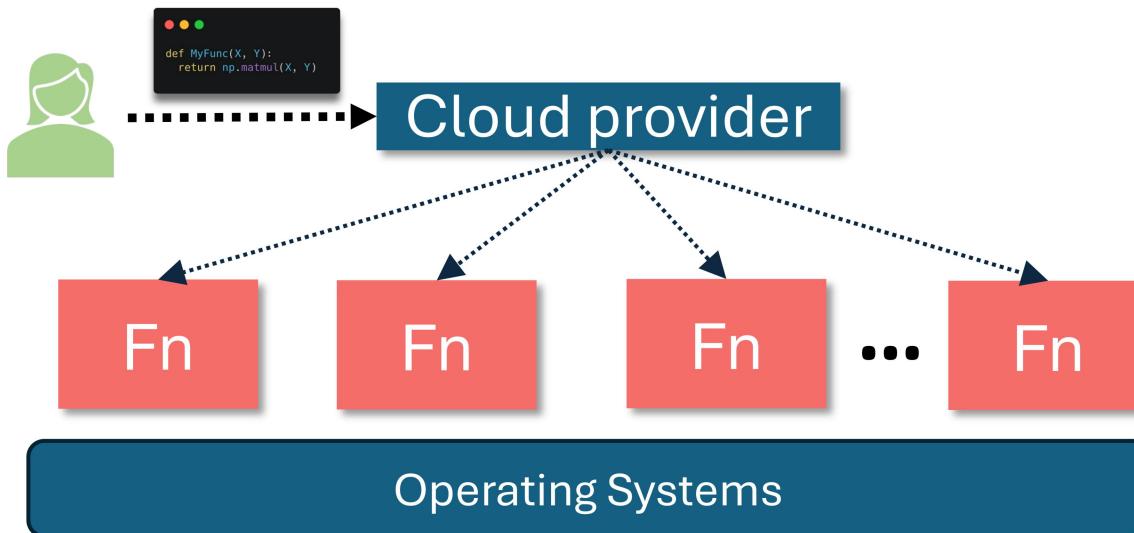
Yuqi Fu¹, Ruizhe Shi², Haoliang Wang³, Songqing Chen², Yue Cheng¹



Serverless Computing

Function-as-a-Service (FaaS): Cloud function as a basic deployment unit.

Main benefits: **Scalability; Pay as you go; DevOps cost**



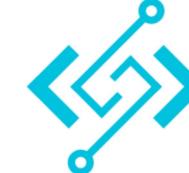
AWS
Lambda



Azure
Functions



Google
Cloud
Functions

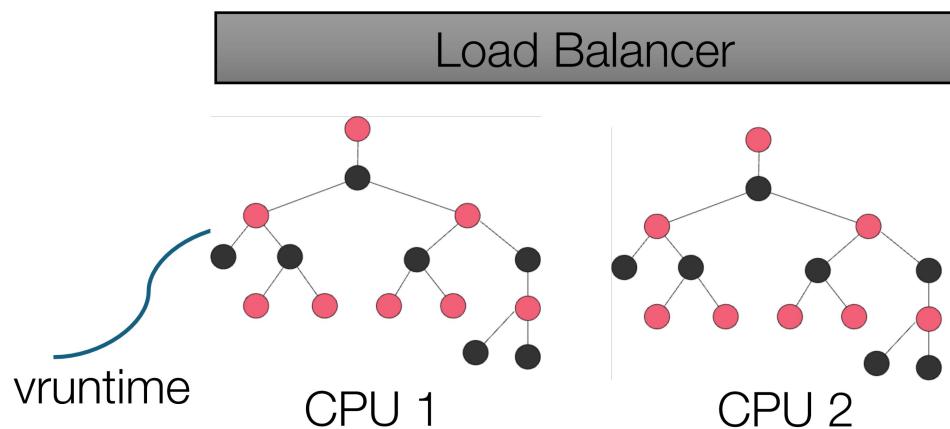


Alibaba
Function
Compute

Current Scheduler Limitations

CFS (Completely Fair Scheduler)

- Proportional-share time slice
- Default Linux scheduler

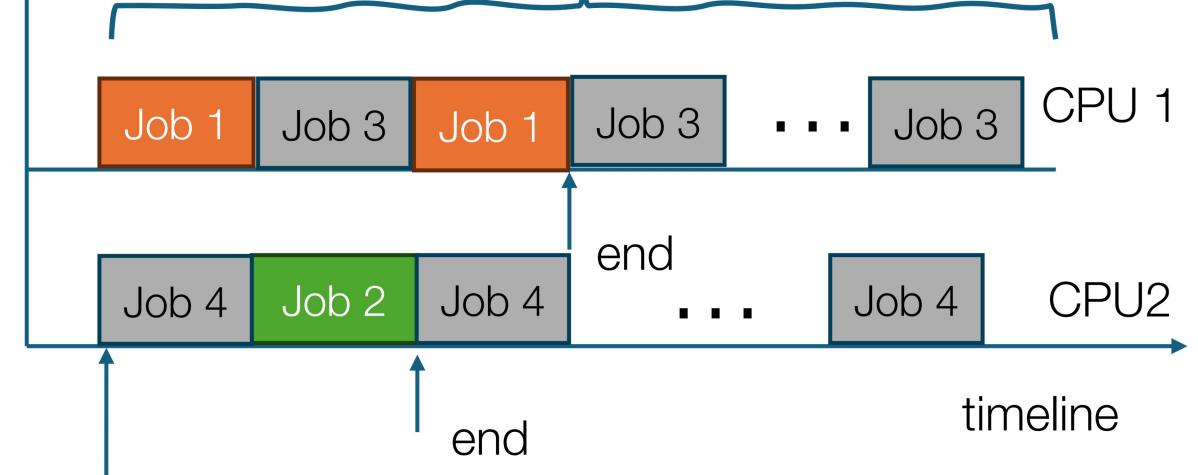


Short jobs: 1, 2

Long jobs: 3, 4



Fairness respect to vruntime



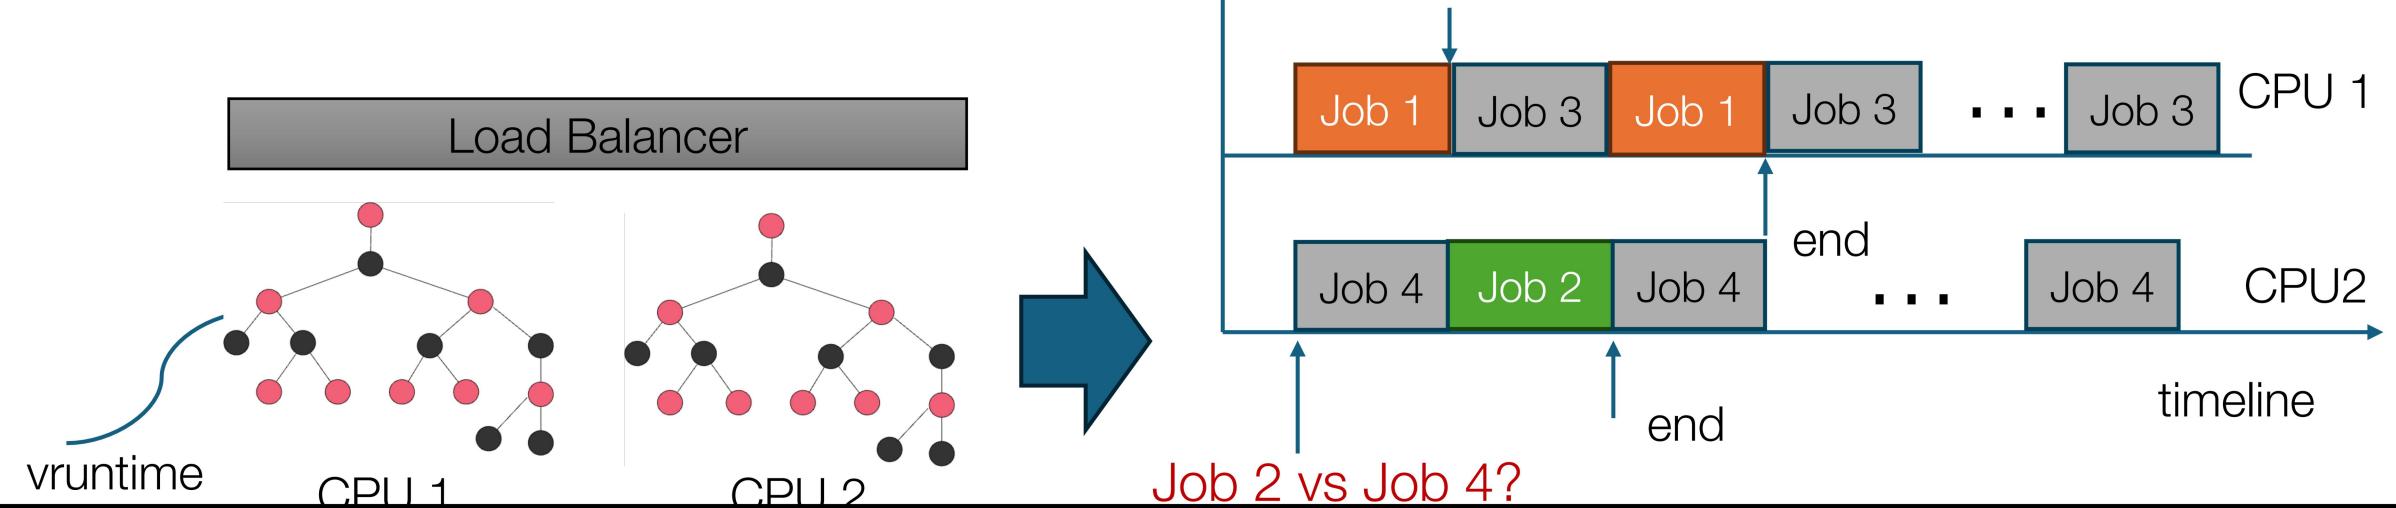
Context switch &
Update vruntime &
Find next task

Current Scheduler Limitations

CFS (Completely Fair Scheduler)

- Proportional-share time slice
- Default Linux scheduler

Short jobs: 1, 2
Long jobs: 3, 4



- **Implication: CFS lacks application workload intelligence**

Approximating SRPT

SRPT (Shortest Remaining Processing Time)

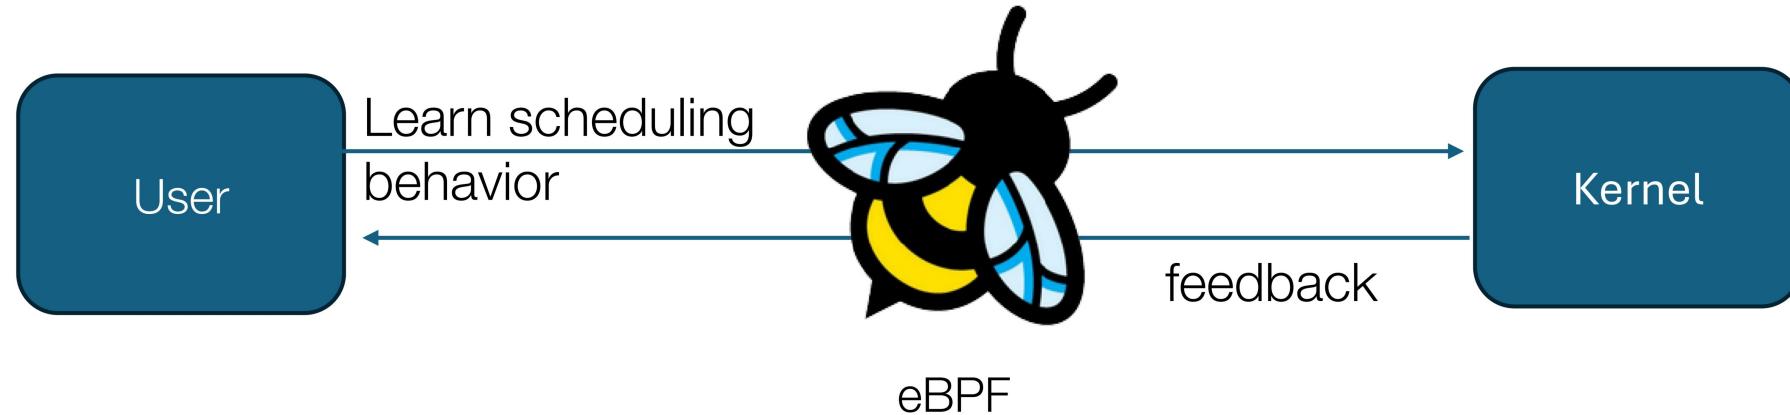
- provably **optimal** for average performance
- **impractical** in online scheduling
- cause **CPU resource starvation** for long-term jobs

- **Implication:** SRPT is impractical and causes starvation for long jobs

Outline

- Motivation
- ALPS design and implementation
- Evaluation

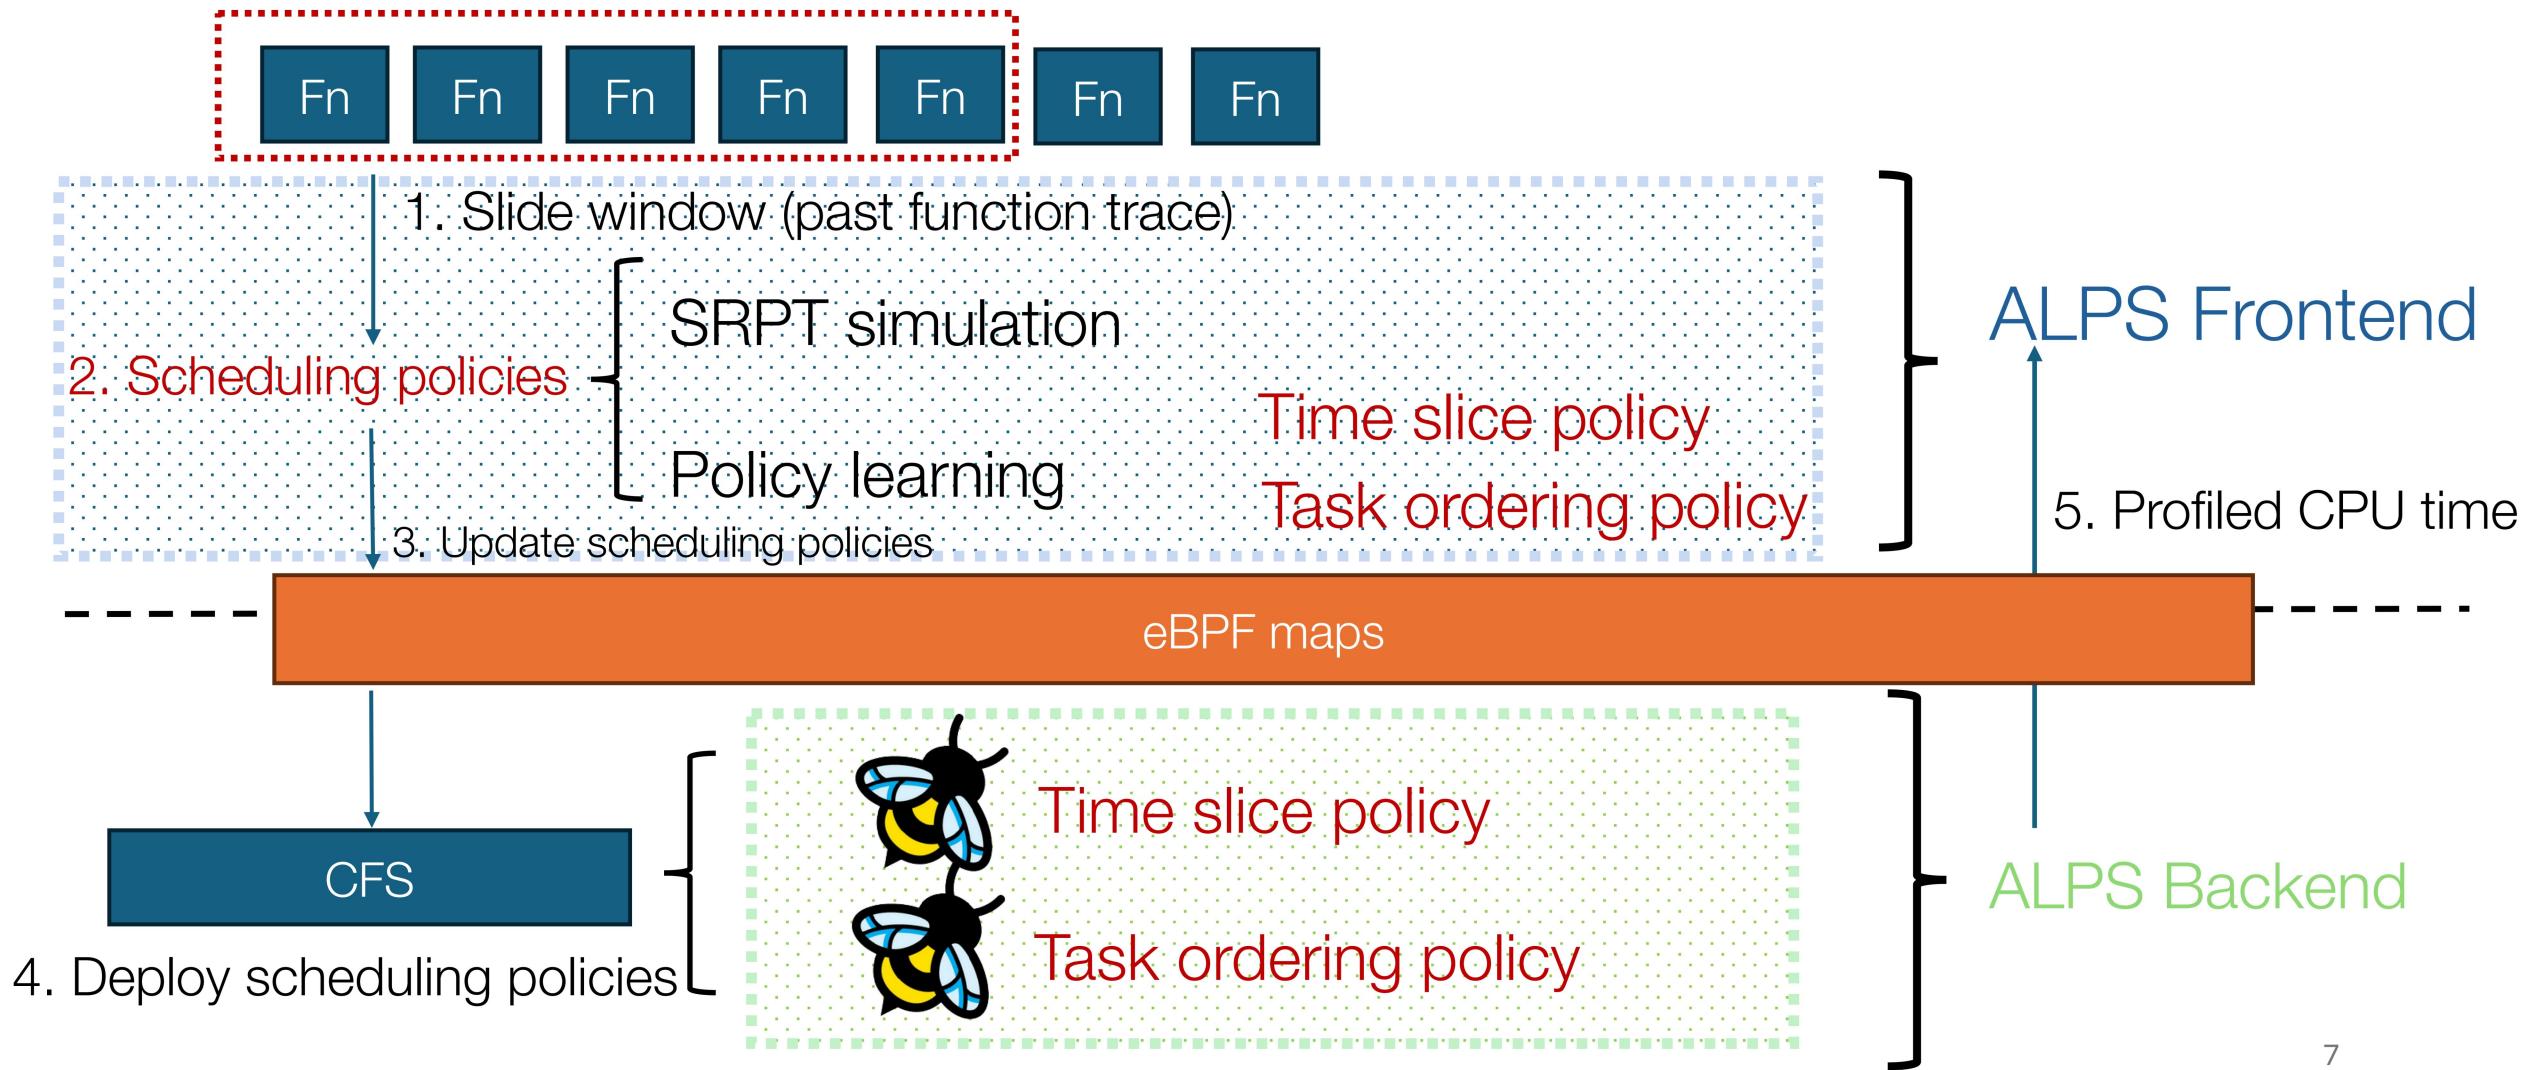
Birds' eye view of ALPS design



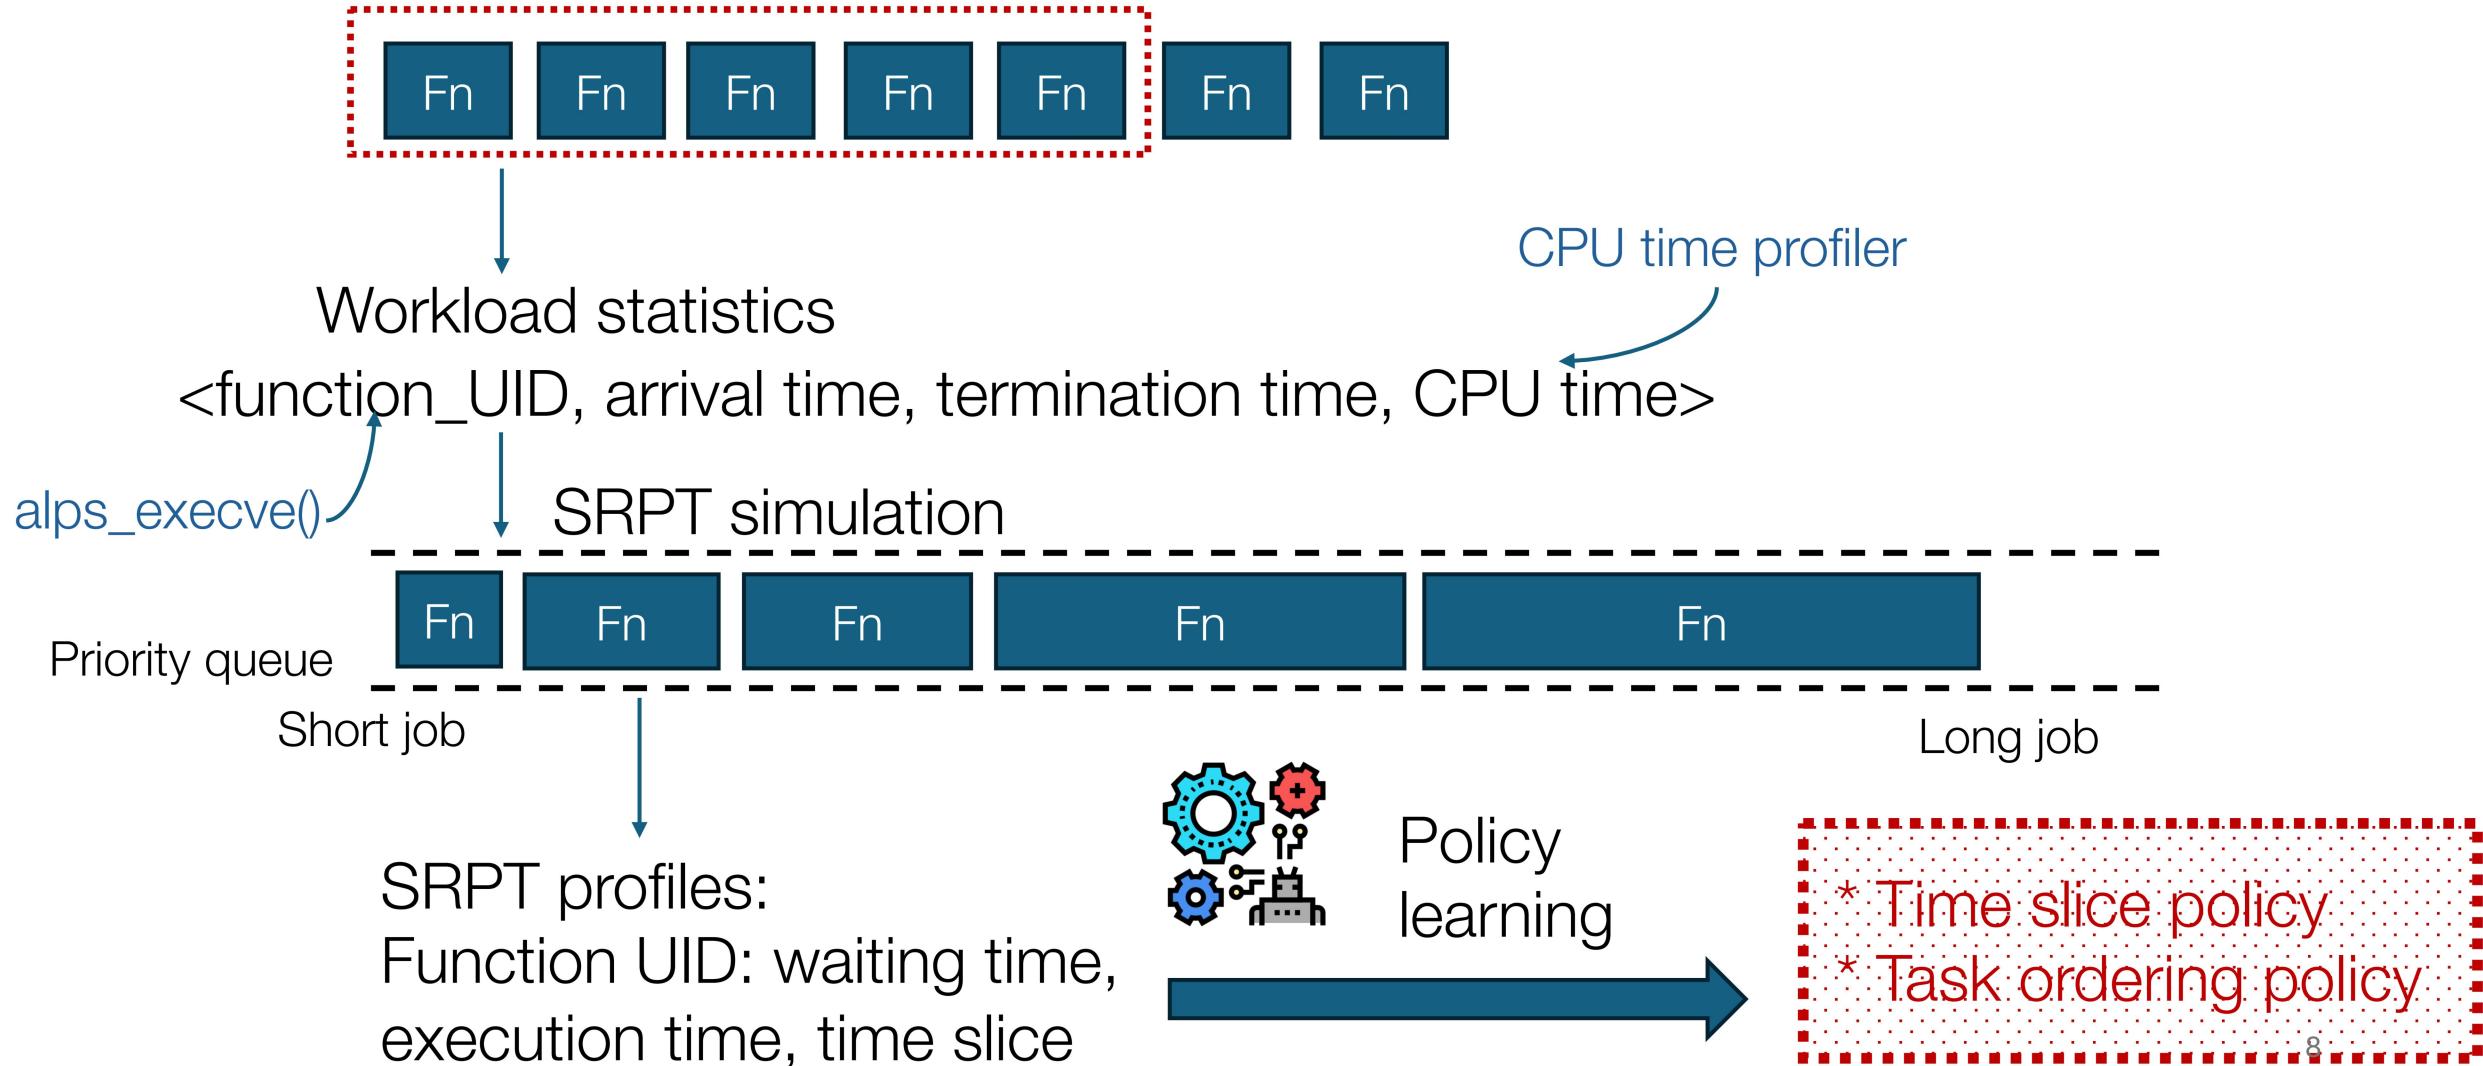
ALPS has a novel OS scheduler architecture that decouples a user-space frontend and kernel-space backend

- **Policy:** Time slice policy & time ordering policy
- **Mechanism:** Leveraging eBPF for user-kernel-space communication

ALPS Design

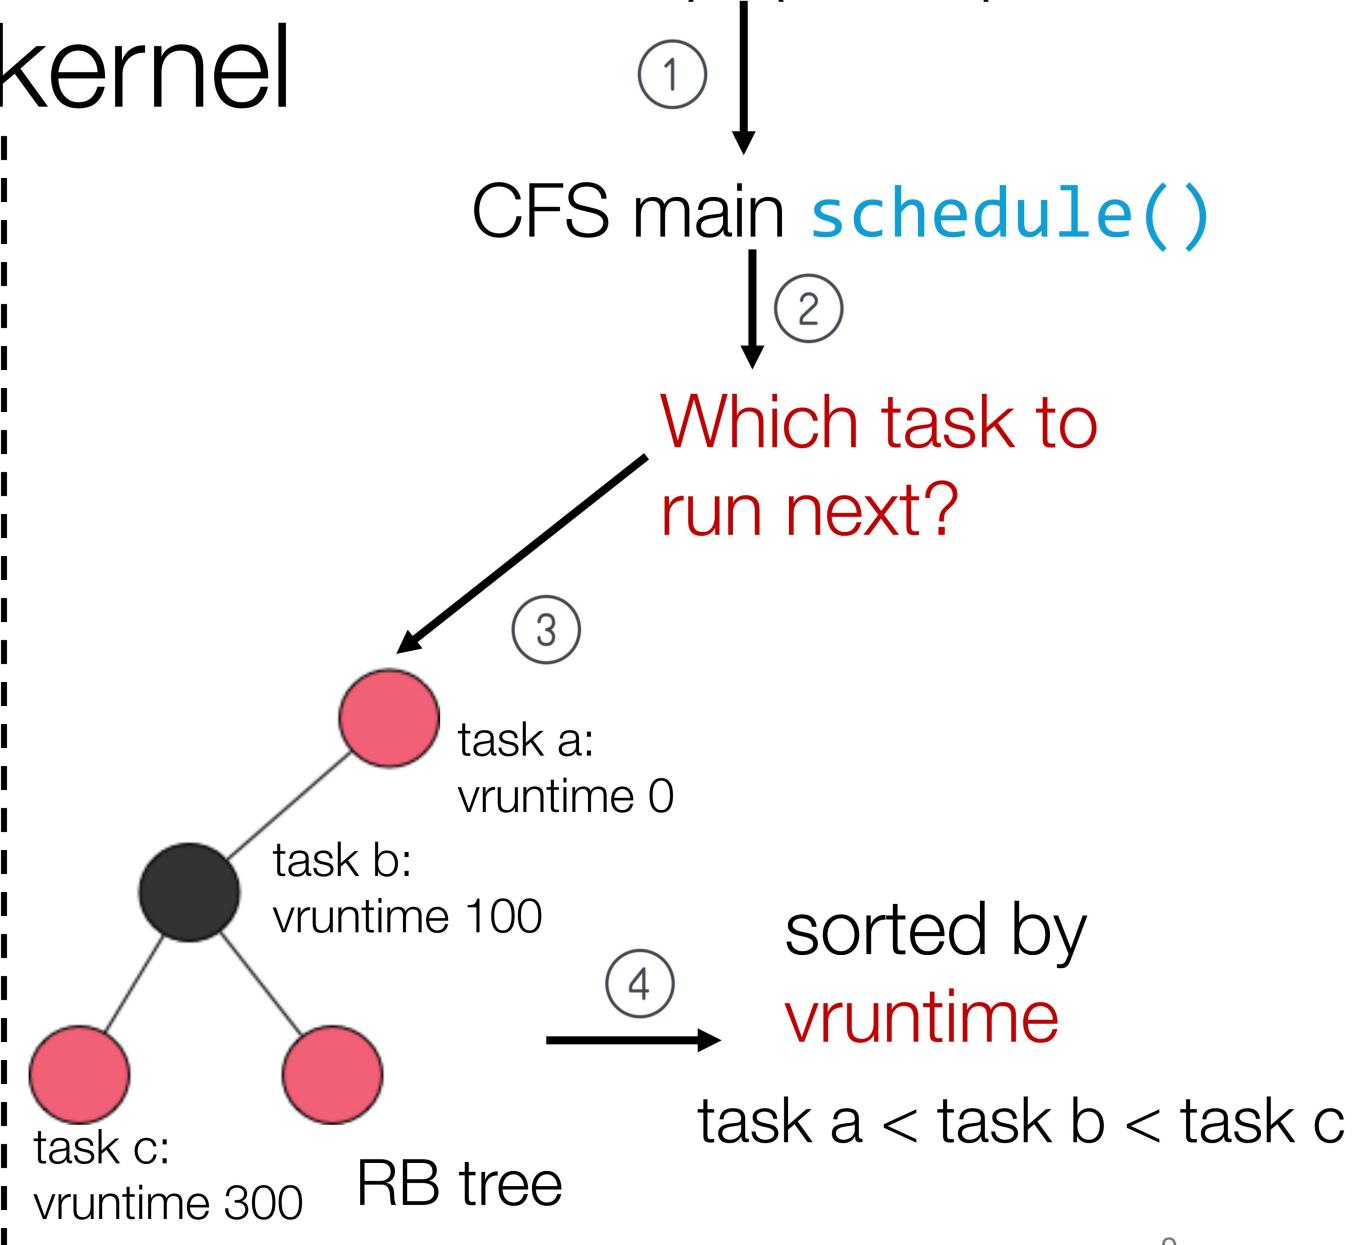


ALPS frontend: SRPT simulation



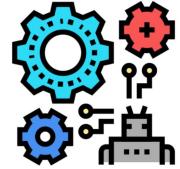
CFS workflow in the kernel

User space | Kernel space tick interrupt/preemption/...



ALPS' time slice policy

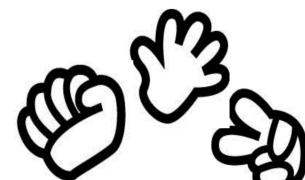
SRPT simulation profiles



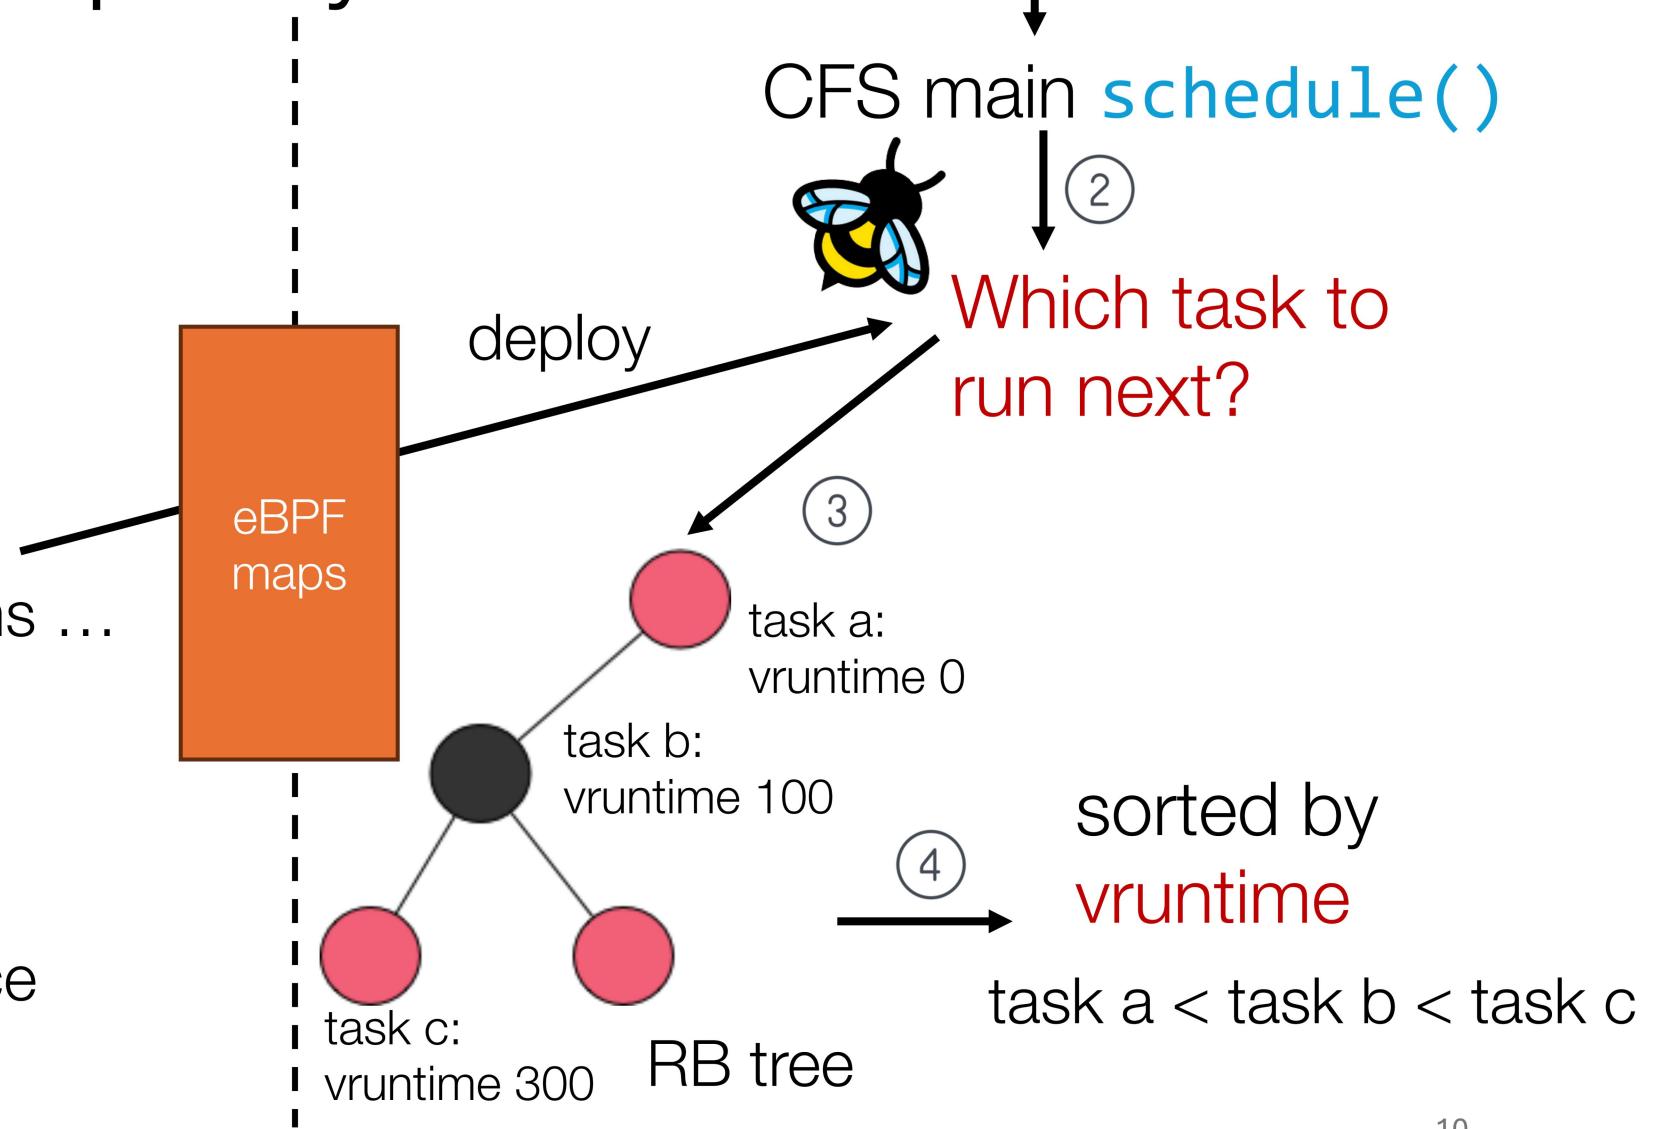
Policy learning

Expected time slice
upper bound for function
e.g. task a: 40ms; task b: 8ms ...

Current time slice

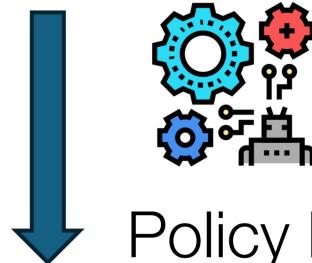


upper bound time slice



ALPS' time slice policy

SRPT simulation profiles

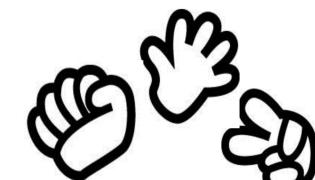


Policy learning

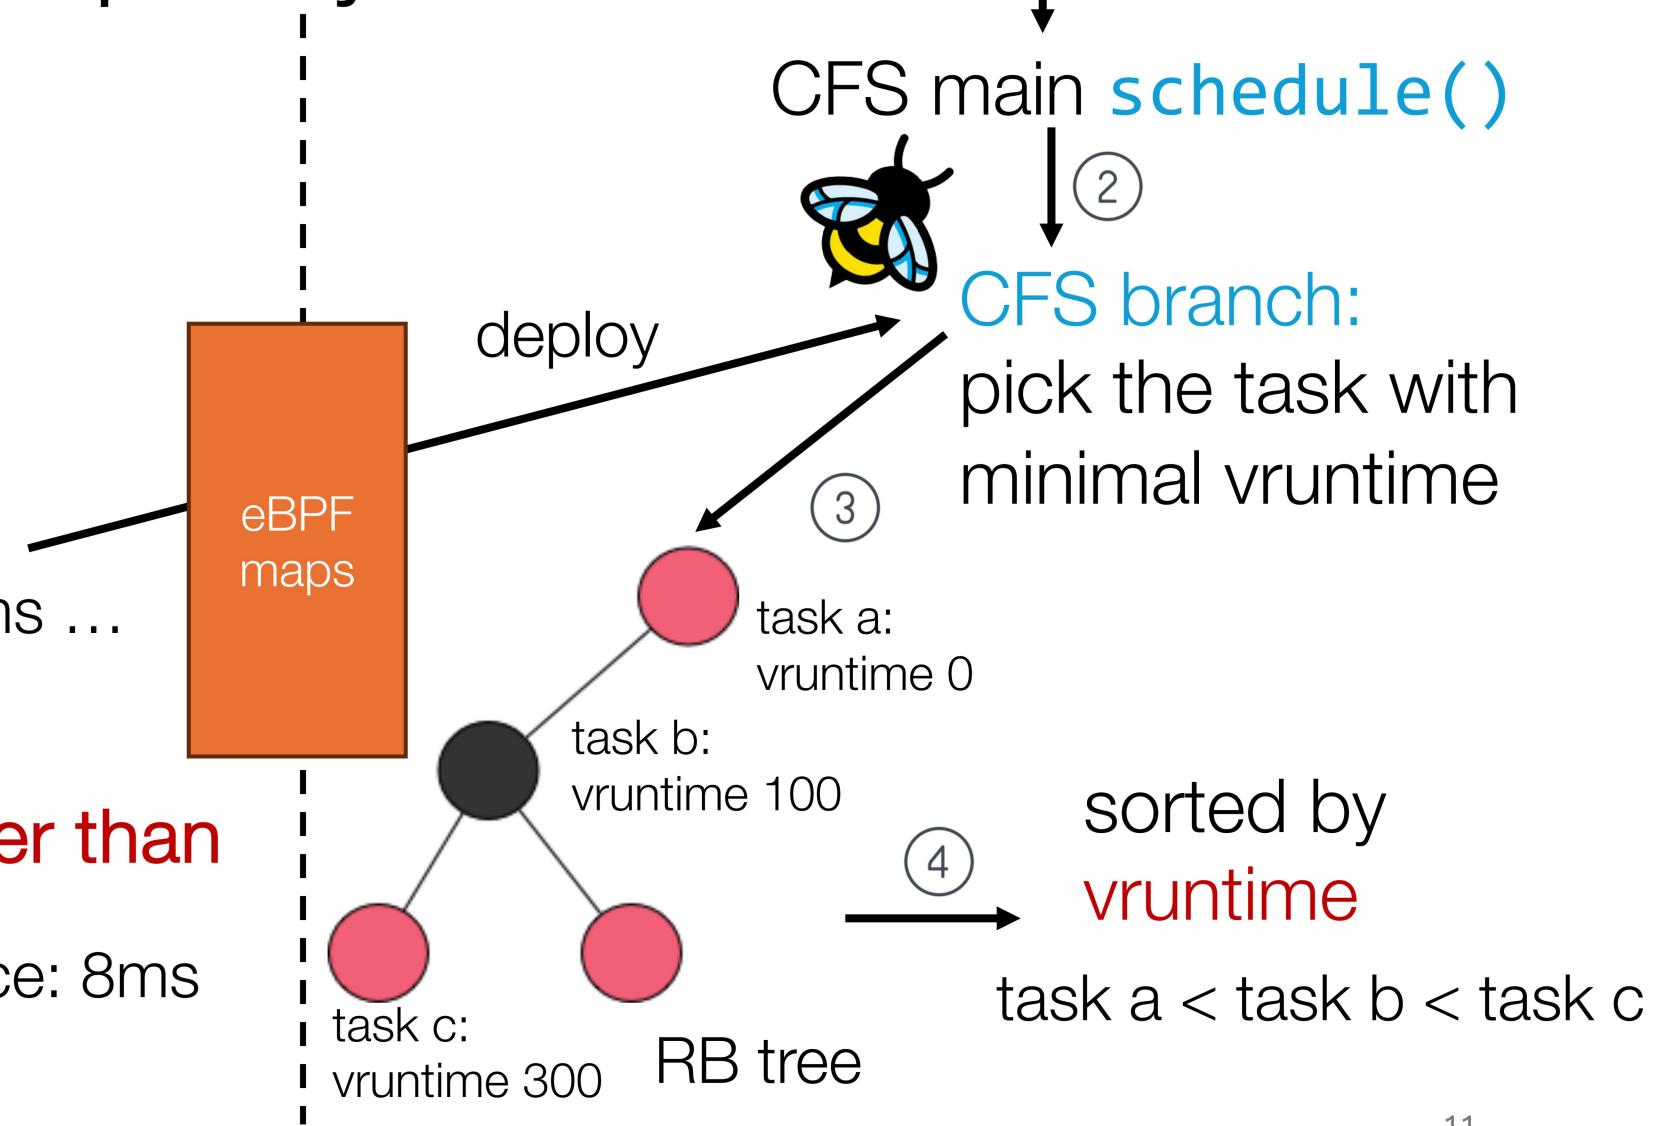
Expected time slice
upper bound for function

e.g. task a: 40ms; task b: 8ms ...

Current time slice: 40 ms

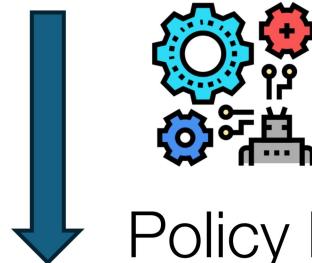


Greater than
upper bound time slice: 8ms



ALPS' time slice policy

SRPT simulation profiles



Policy learning

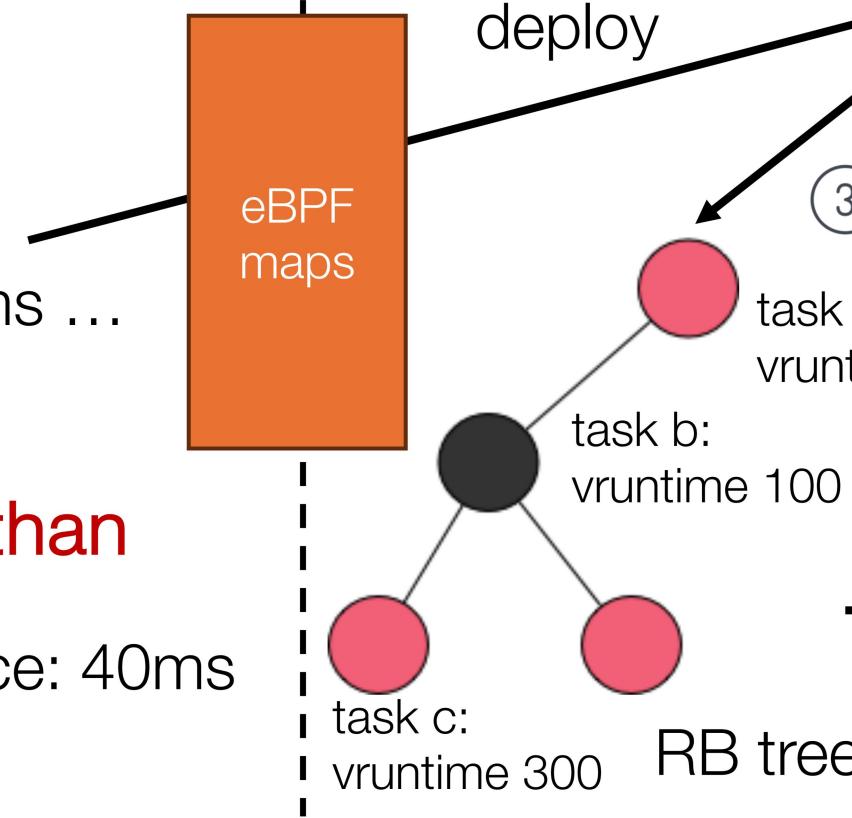
Expected time slice
upper bound for function

e.g. task a: 40ms; task b: 8ms ...

Current time slice: 8 ms



Less than
upper bound time slice: 40ms



tick interrupt/preemption/...

①

CFS main **schedule()**



②

ALPS branch:

grant additional
time slice to current
task

③

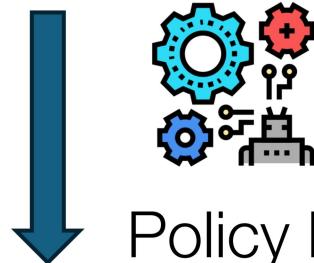
④

sorted by
vruntime

task a < task b < task c

ALPS' task ordering policy

SRPT simulation profiles

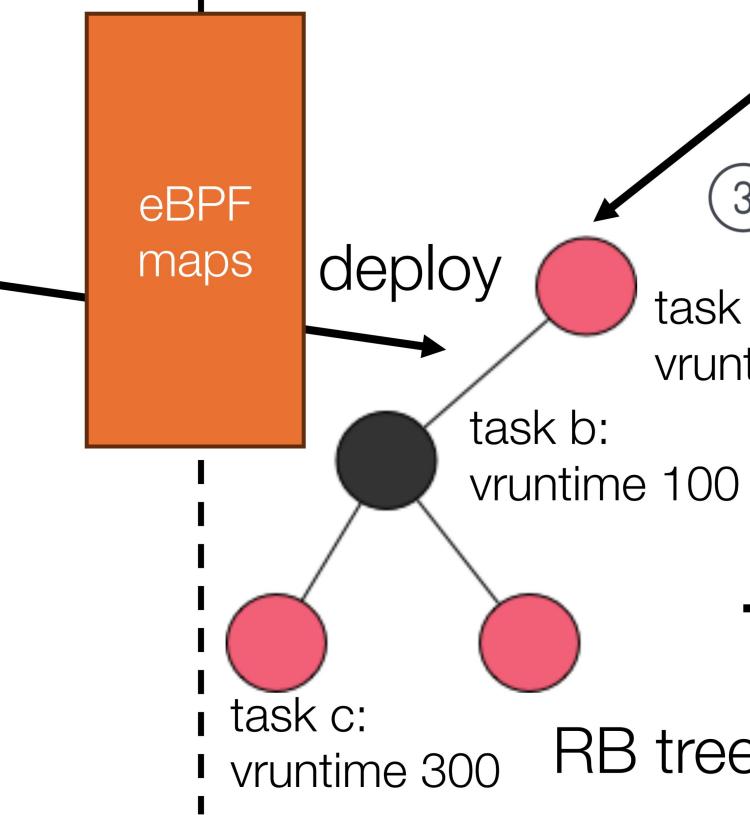


Policy learning

Expected ALPS task priority
for function

e.g. task a: 2; task b: 1 ...

ALPS priority overrides CFS
vruntime



CFS main `schedule()`

Which task to
run next?

③

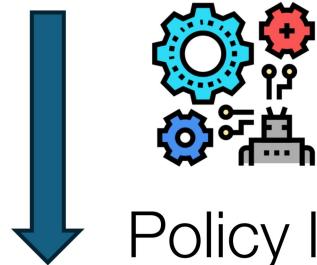
④

sorted by
vruntime

task a < task b < task c

ALPS' task ordering policy

SRPT simulation profiles

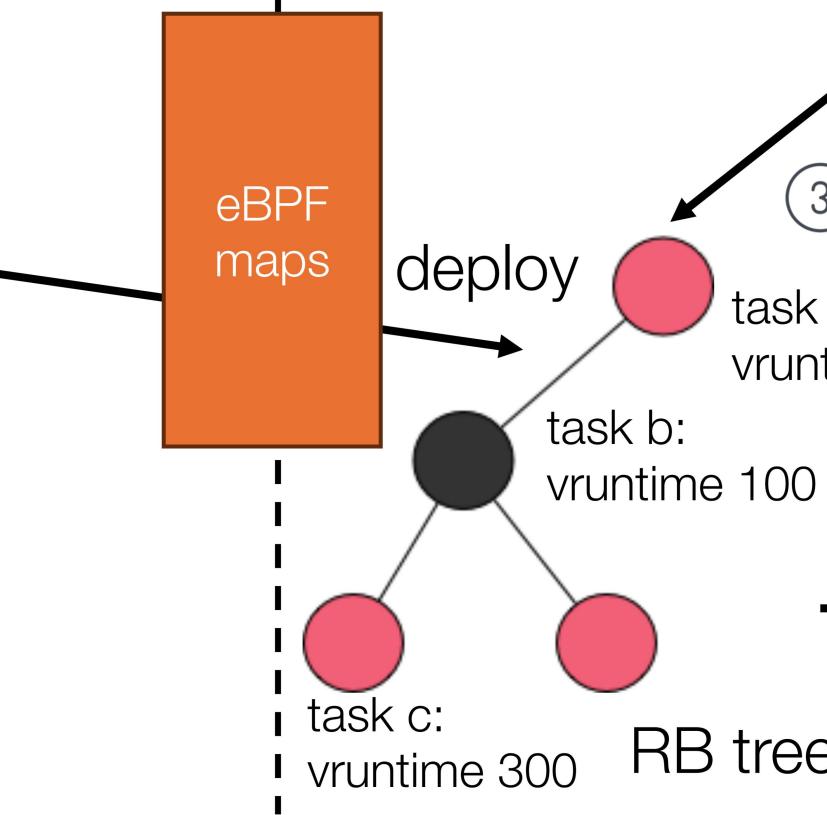


Policy learning

Expected **ALPS task priority** for function

e.g. task a: 2; task b: 1 ...

ALPS priority overrides CFS vruntime



CFS main `schedule()`

(1)

Which task to run next?

(2)

task a:

task b:

task c:

vruntime 0

vruntime 100

vruntime 300

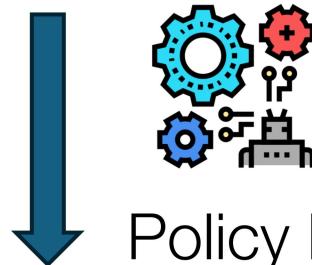
RB tree

(3)

sorted by
ALPS priority
task b < task a < task c

ALPS' task ordering policy

SRPT simulation profiles

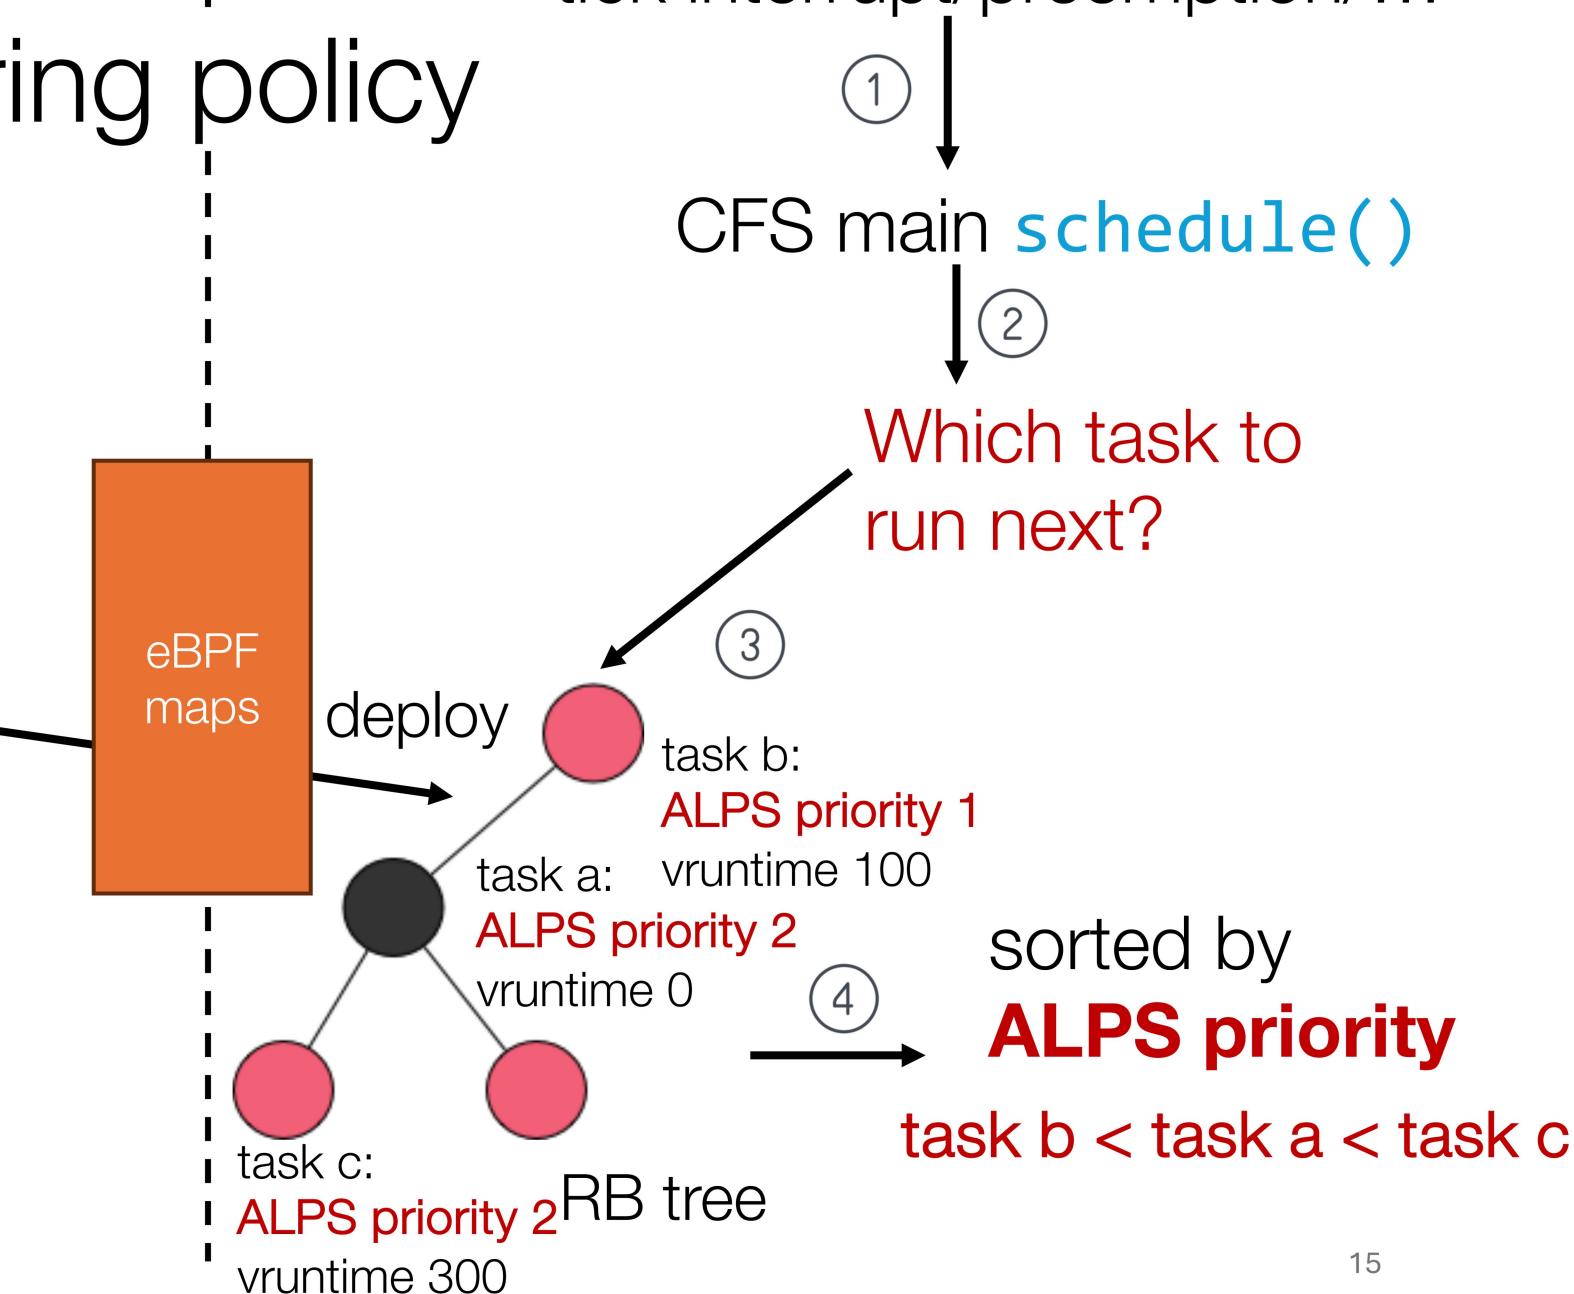


Policy learning

Expected ALPS task priority
for function

e.g. task a: 2; task b: 1 ...

ALPS priority overrides CFS
vruntime



Outline

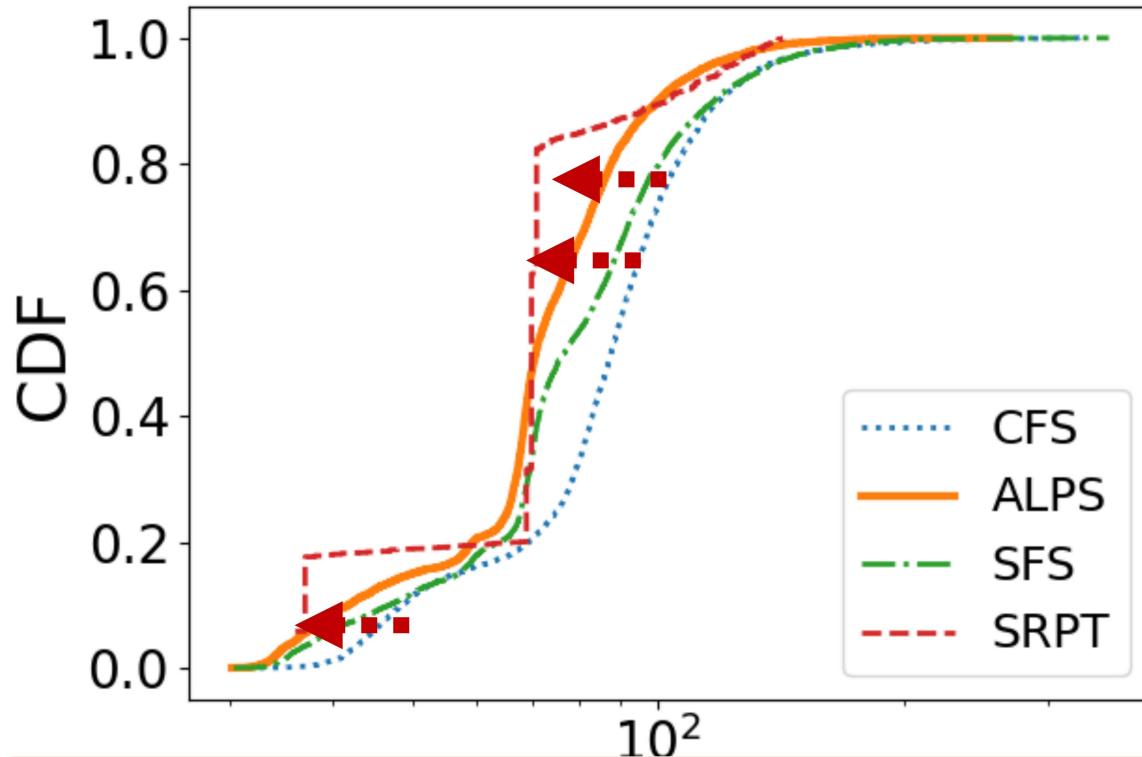
- Motivation
- ALPS design and implementation
- Evaluation

Experimental methodology

- Platform:
 - OpenLambda & Docker
 - We modified **135** LoC in OpenLambda and **223** LoC in Docker
- Host machine: bare-metal with **56** CPUs and **256** GB RAM
- OS: Ubuntu 22.04.1 LTS
- Production workload traces:
 - Huawei Cloud Functions trace [SoCC'23]
 - Azure Functions trace [ATC'20]

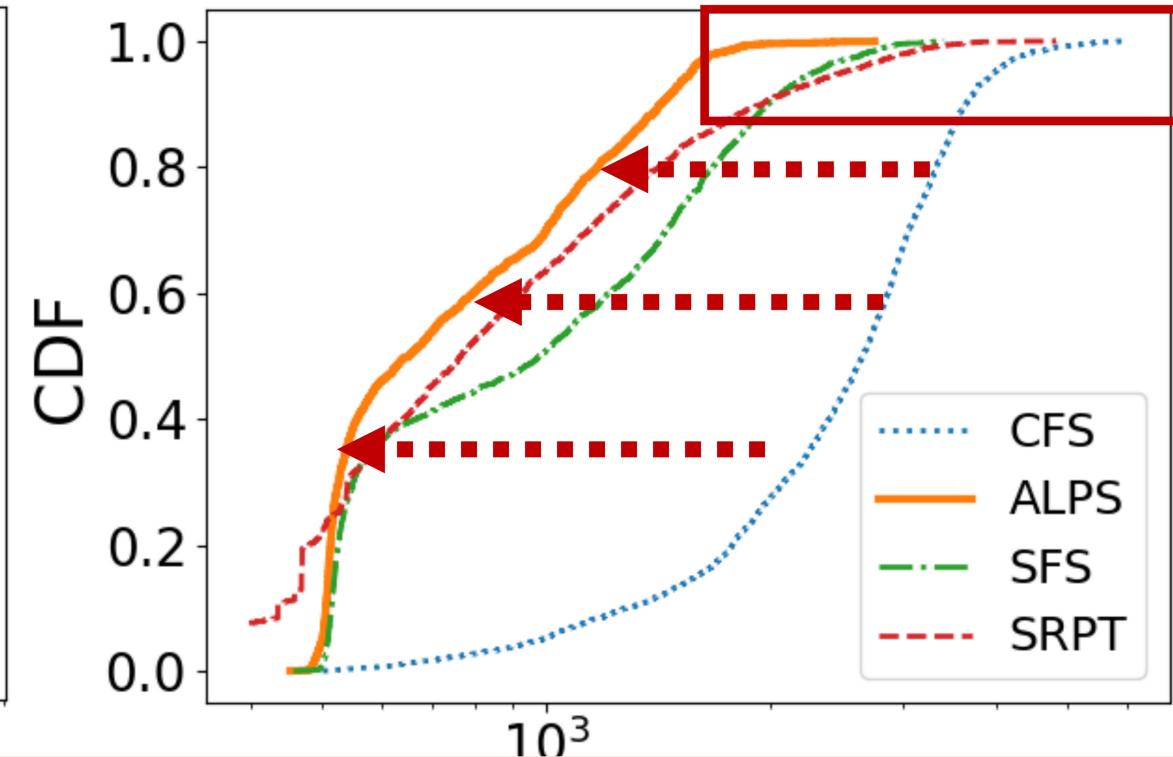
End-to-End Performance: Azure Workload

Short Functions



Shorter tail latency

Long Functions



ALPS achieves shorter execution durations compared to CFS

Conclusion

- ALPS continuously learns FaaS **workload intelligence** from user-space SRPT simulation.
- We built a **prototype** of ALPS into a user-space frontend and a kernel-space backend atop Linux CFS using customized **eBPF functions** and hooks.
- Extensive evaluation shows that ALPS improves the performance for both short functions and long functions compared to CFS.