

Enabling Tensor Language Model to Assist in Generating High-Performance Tensor Programs for Deep Learning

Authors: Yi Zhai¹, Sijia Yang², Keyu Pan³, Renwei Zhang², Shuo Liu¹
Chao Liu and Zichun Ye², Jianmin Ji¹, Jie Zhao⁴,
Yu Zhang and Yanyong Zhang¹

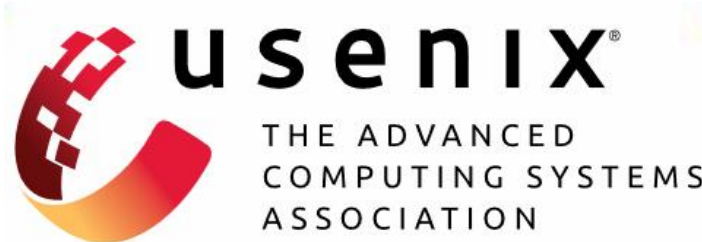
Reporter: Hangshuai He

¹ University of Science and Technology of China,

² Huawei Technologies Co., Ltd.

³ ByteDance Ltd.

⁴ Hunan University



Primary research interests:

- Program Languages, Computer Systems, Parallel Computing
- Smart IoT, Smart Sensing, Smart Unmanned Systems for Sensing



Recent Papers:

- [CGO' 25] *GoFree: Reducing Garbage Collection via Compiler-inserted Freeing*
- [TCAD' 25] *PauliForest: Connectivity-Aware Synthesis and Pauli-Oriented Qubit Mapping for Near Term Quantum Simulation*
- [OOPSLA' 24] *MEA2: a Lightweight Field-Sensitive Escape Analysis with Points-to Calculation for Golang*
- [DAC' 24] *Crop: An Analytical Cost Model for Cross-Platform Performance Prediction of Tensor Programs*

Corresponding Author:

Professor Yu Zhang

IEEE Fellow Yanyong Zhang

[Lab for Intelligent Networking and Knowledge Engineering](#)



- *Background*
- *System Design*
- *Evaluation*
- *Conclusion*



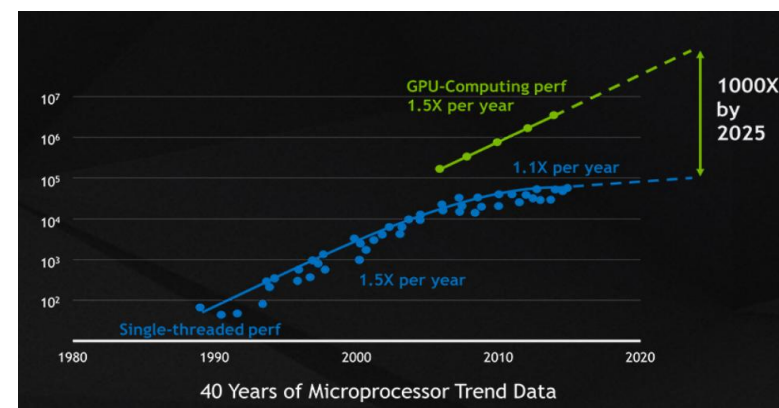
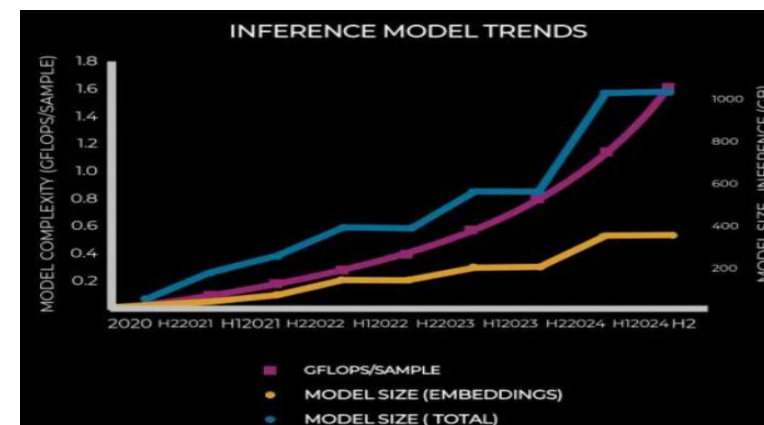
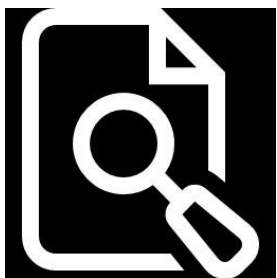
- ***Background***
- *System Design*
- *Evaluation*
- *Conclusion*

Background

The Growing Gap: Exponential DL Demand vs. Linear Hardware Scaling



Scale
complexity



Background

Low-latency execution of workloads is *imminent*



cuDNN

oneDNN



[OSDI' 18] *TVM: An Automated End-to-End Optimizing Compiler for Deep Learning*

[TOG' 19] *Learning to optimize **halide** with tree search and random programs*

[ASPLOS' 20] *Flextensor: An automatic schedule exploration and optimization framework for tensor computation on heterogeneous system*

[ATC' 19] *Optimizing {CNN} model inference on {CPUs}*

Background

kernel libraries is *costly*, tensor compilers are becoming popular



cuDNN



oneDNN



[OSDI' 18] *TVM: An Automated End-to-End Optimizing Compiler for Deep Learning*

[TOG' 19] *Learning to optimize **halide** with tree search and random programs*

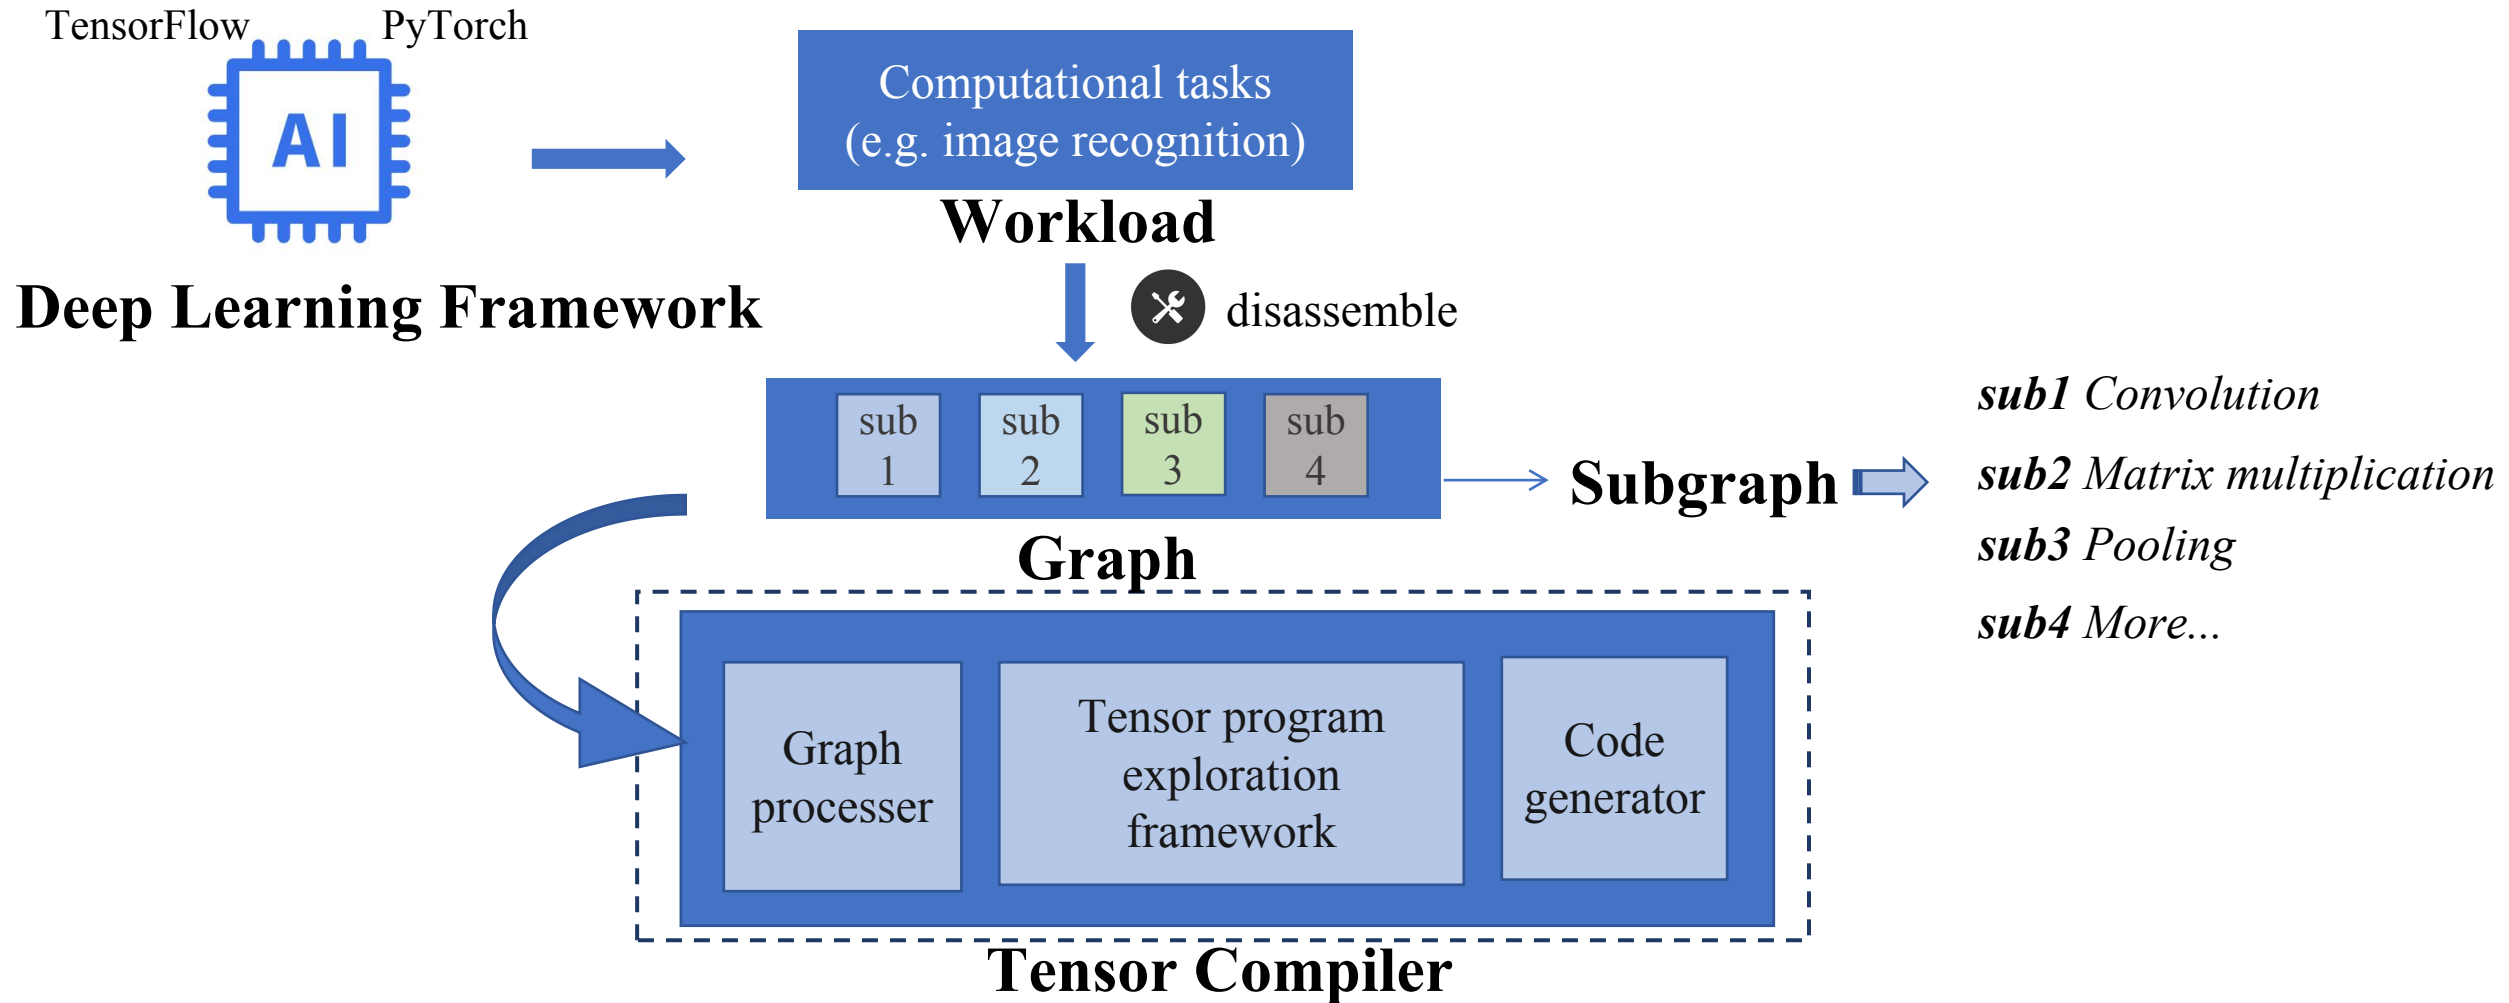
[ASPLOS' 20] *Flextensor: An automatic schedule exploration and optimization framework for tensor computation on heterogeneous system*

[ATC' 19] *Optimizing {CNN} model inference on {CPUs}*

Background

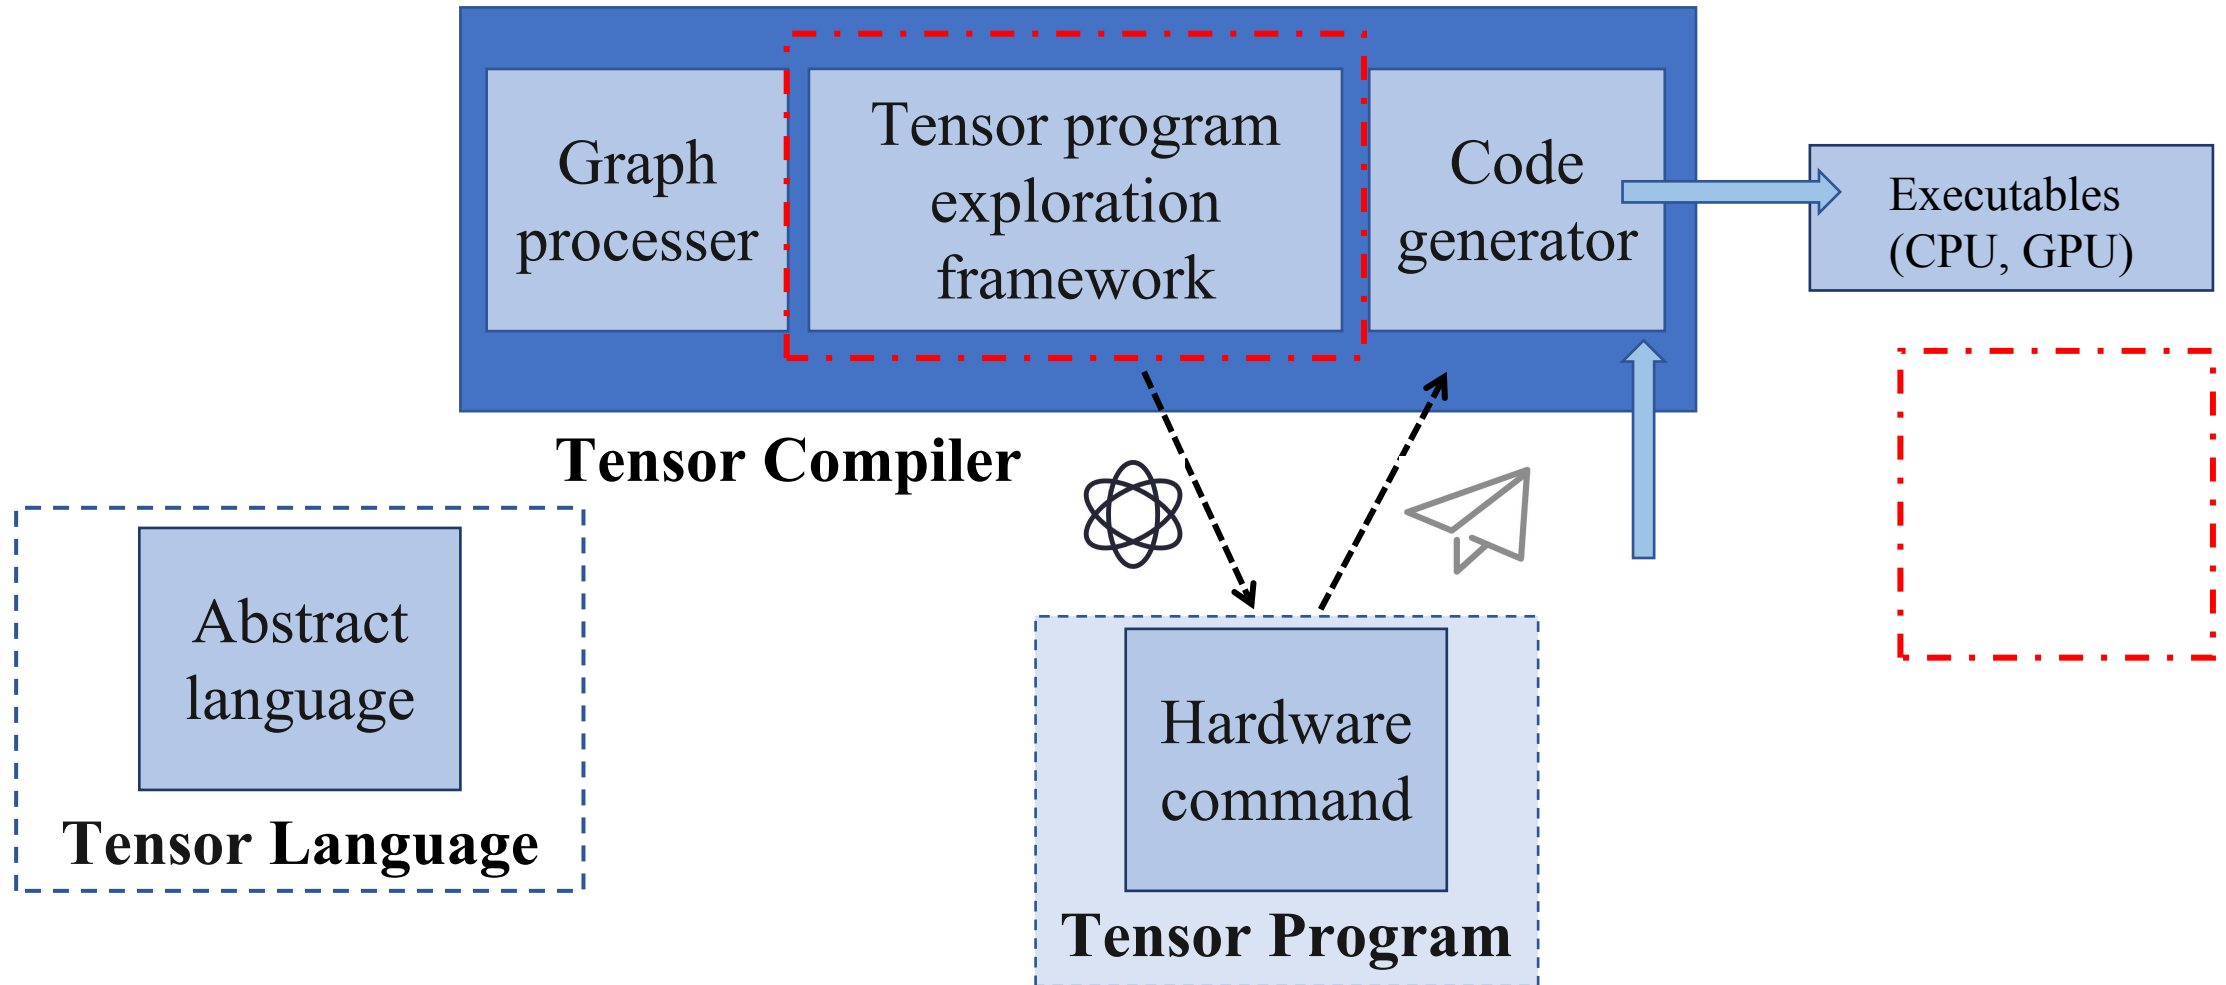


Things about *tensor*



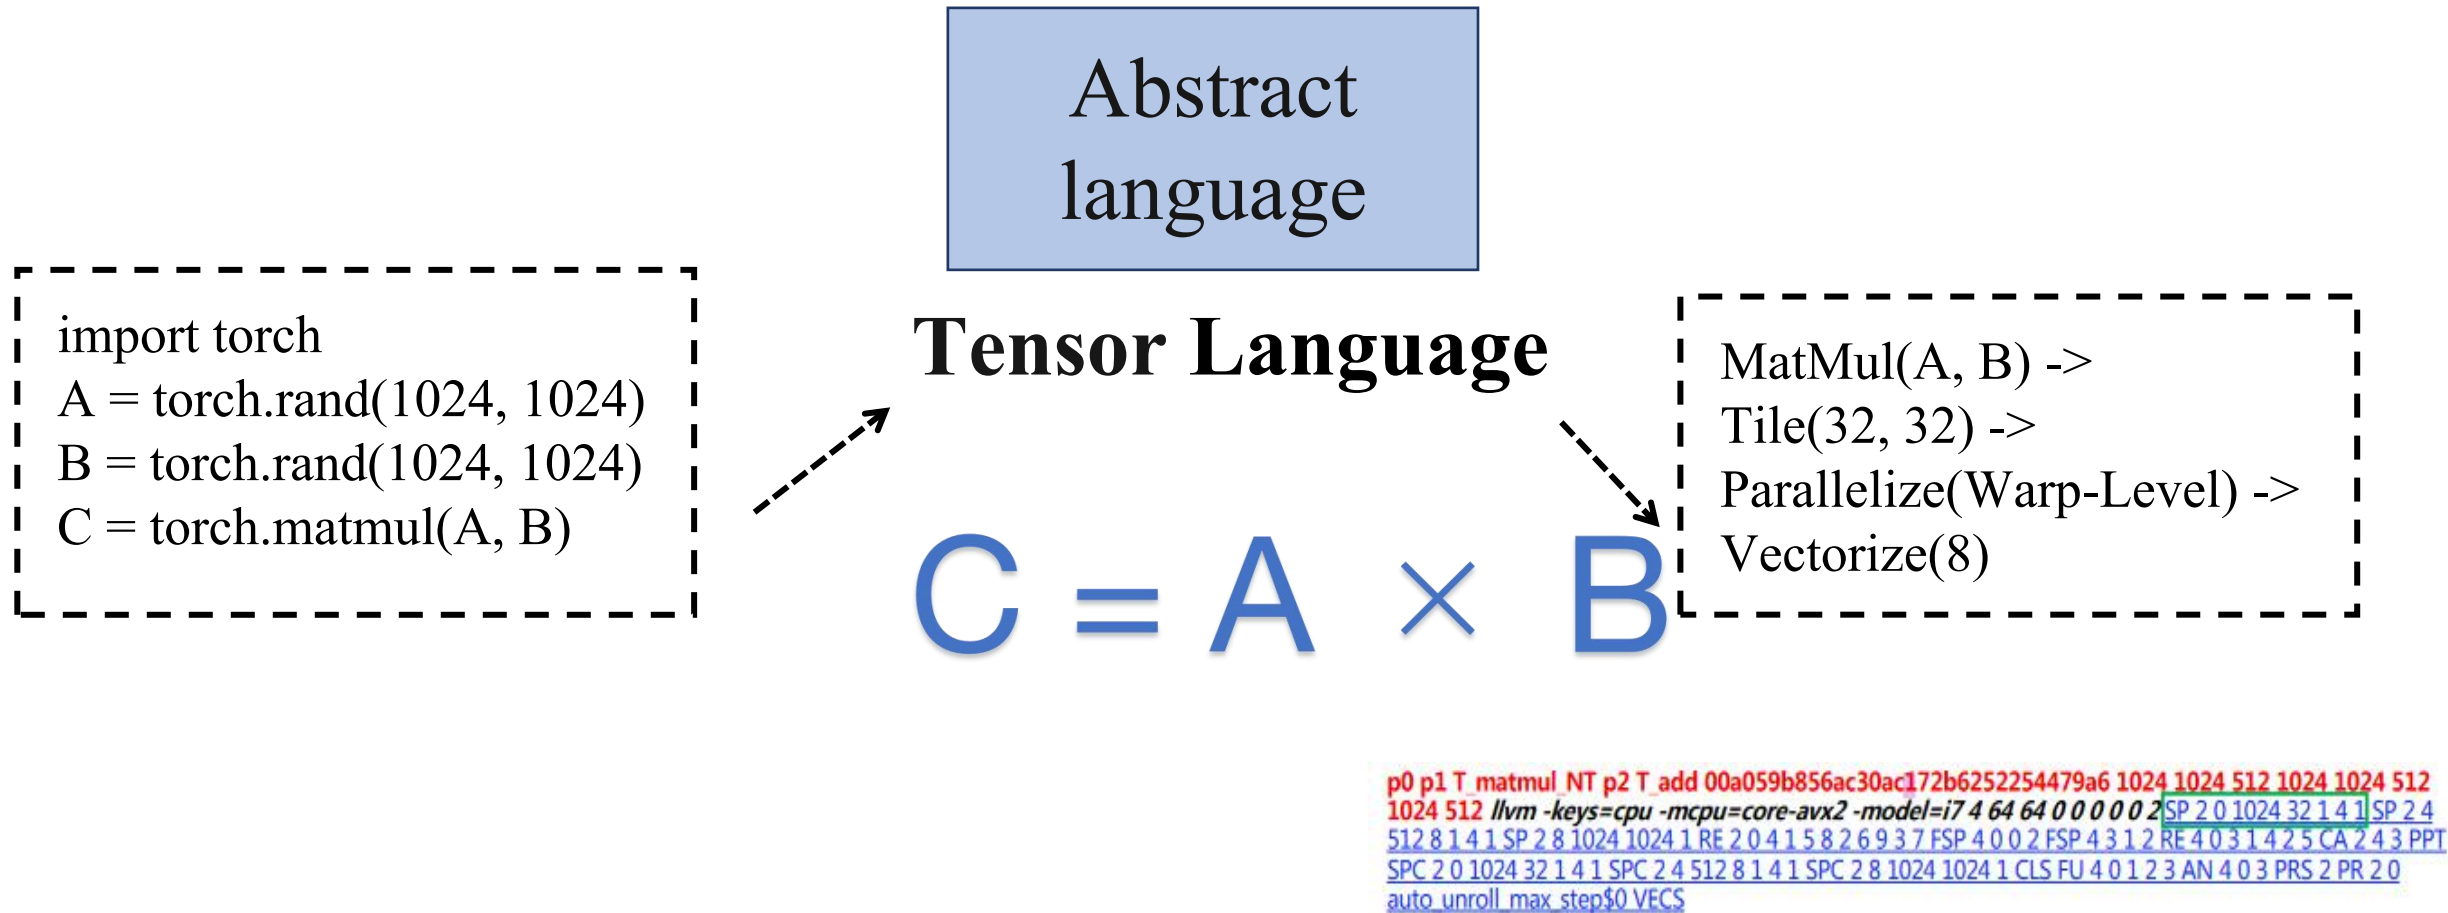
Background

Things about *tensor*



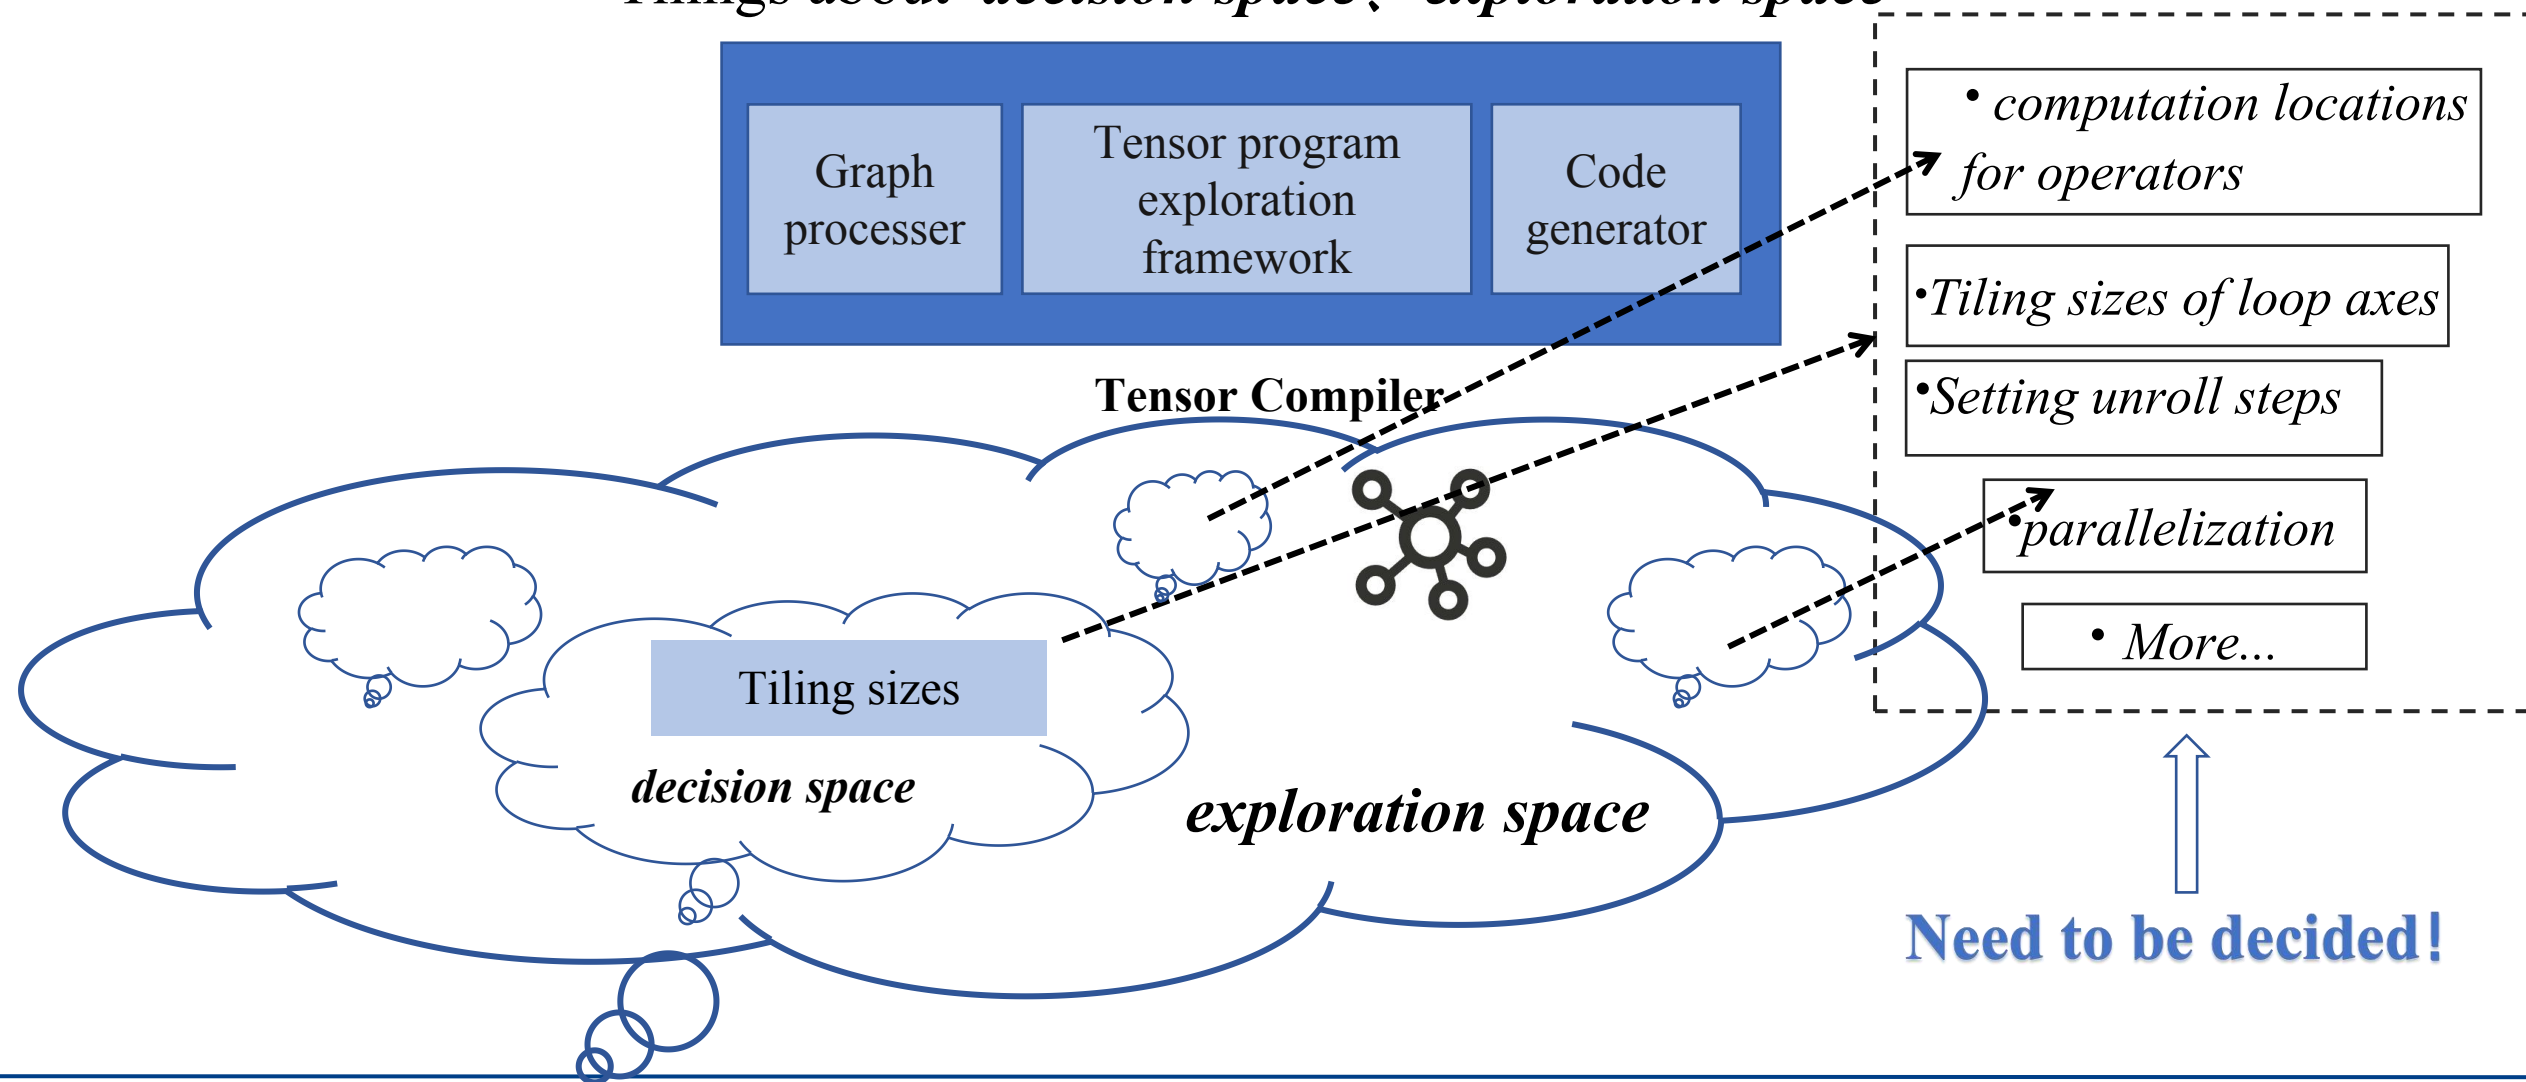
Background

In a clear, structured way, to *document* computational processes and *optimization* decisions



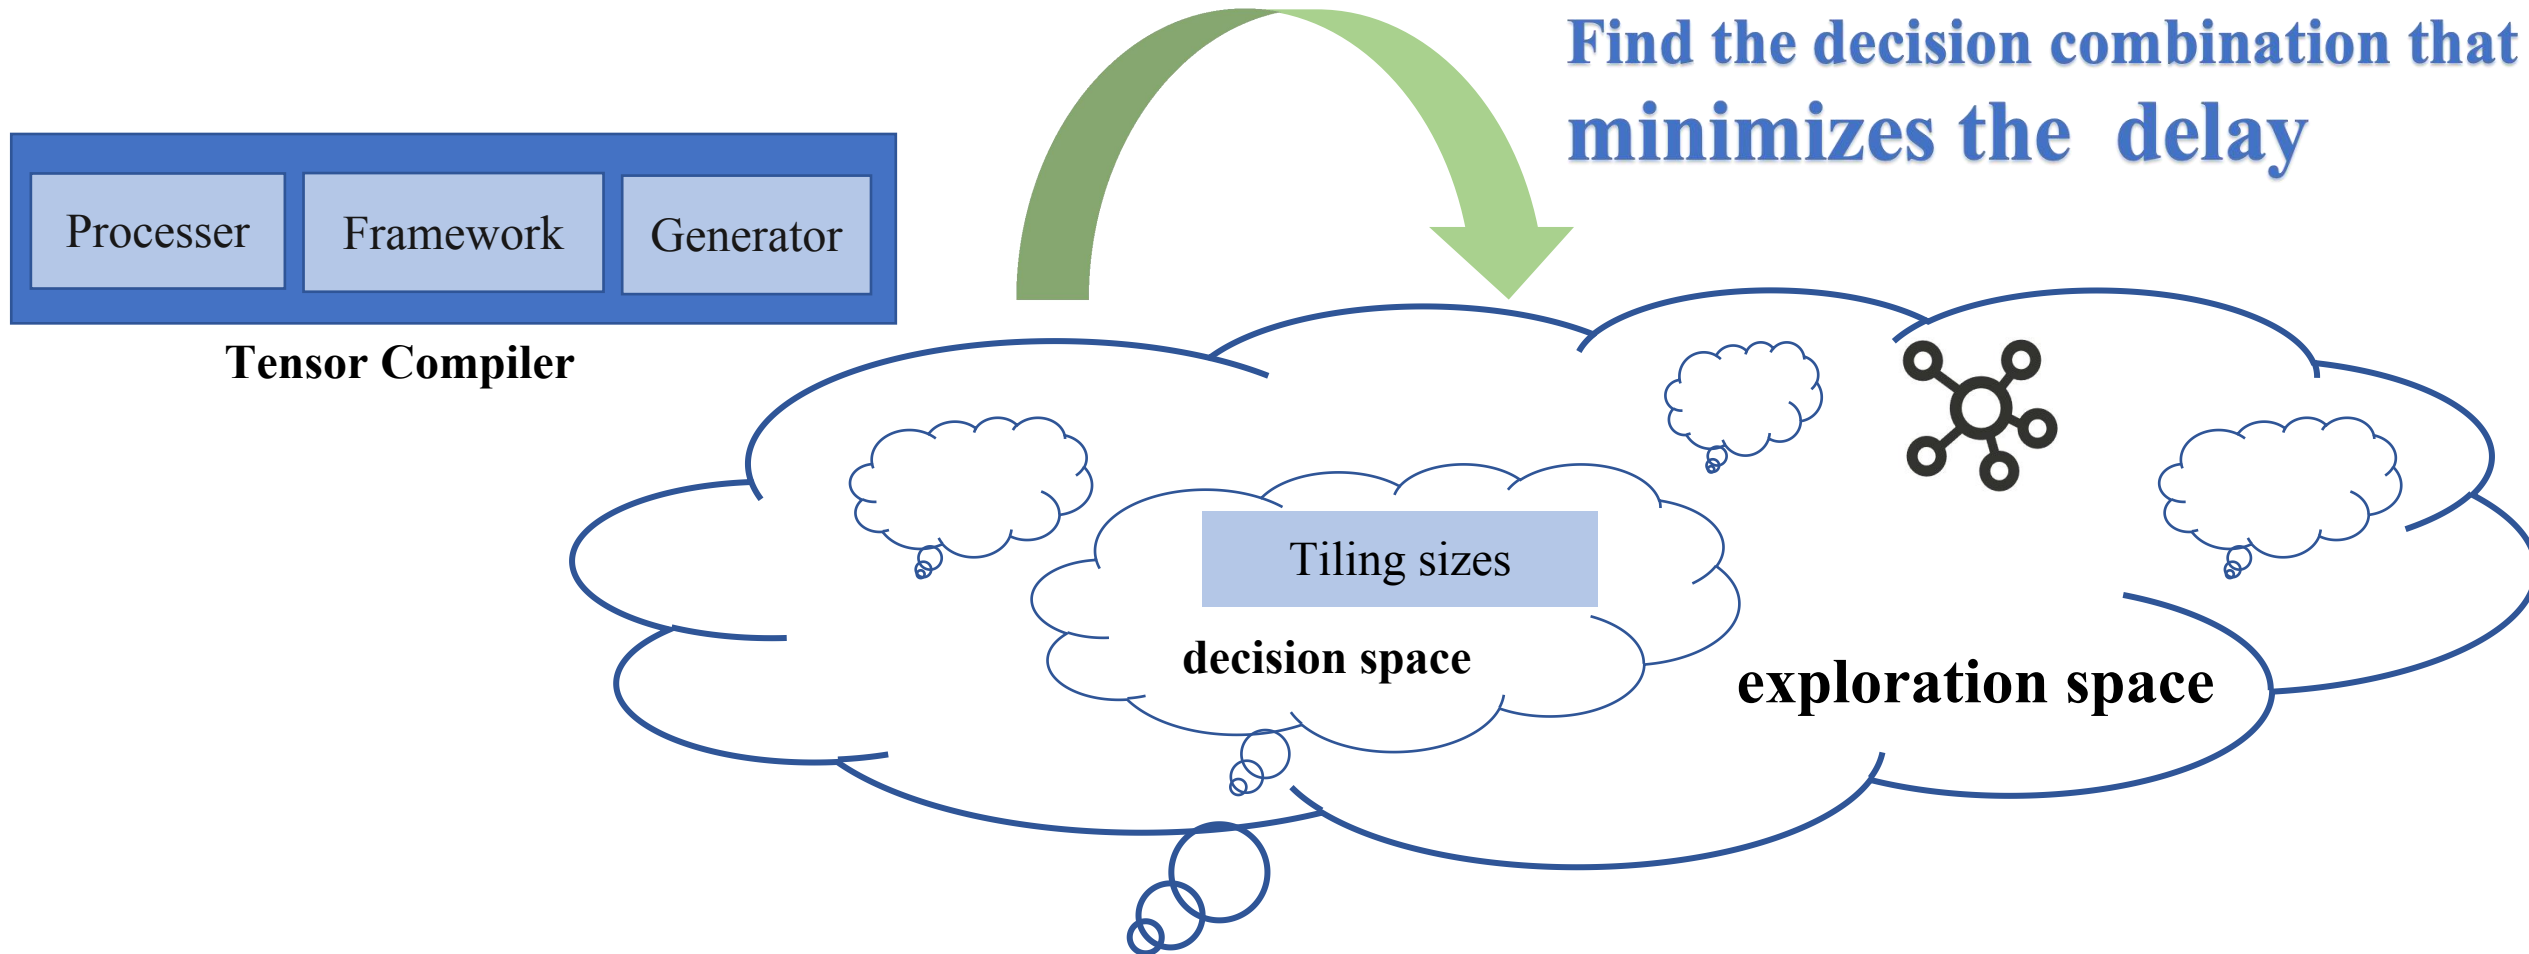
Background

Things about *decision space*, *exploration space*



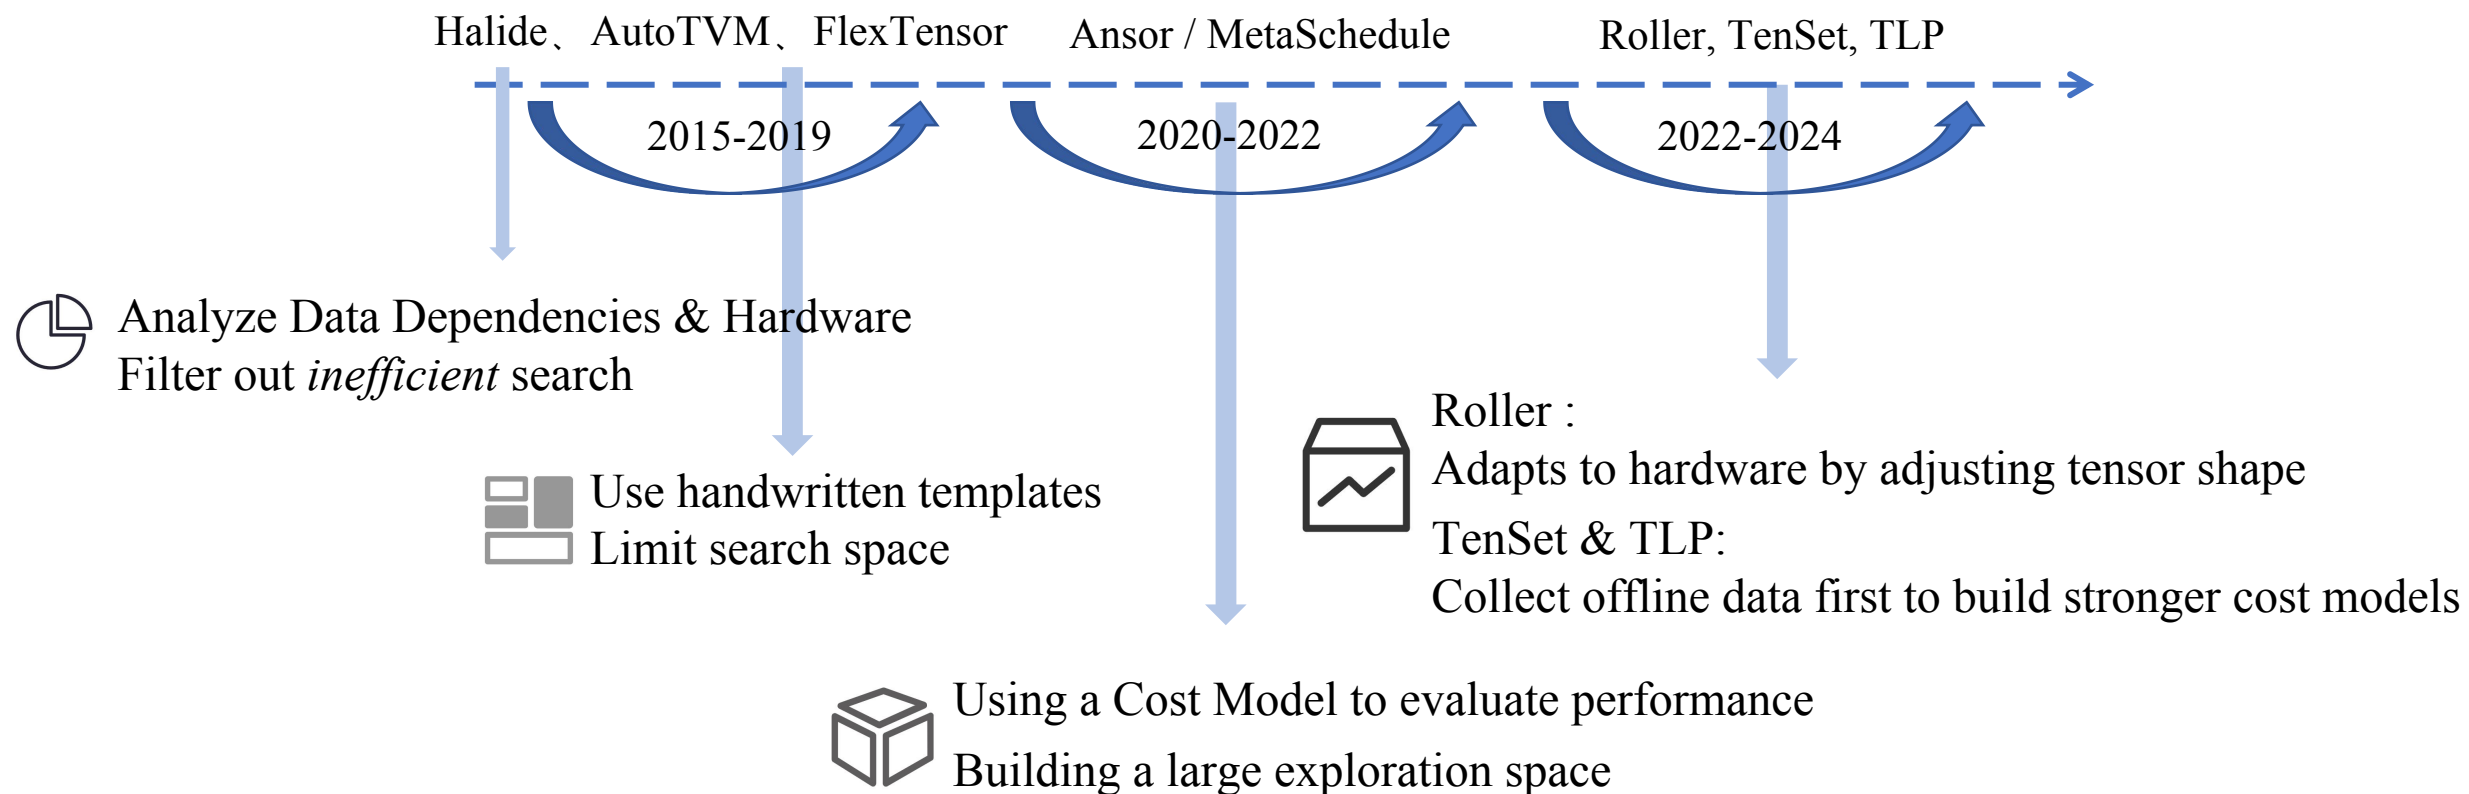
Background

The *objective* of tensor program exploration framework



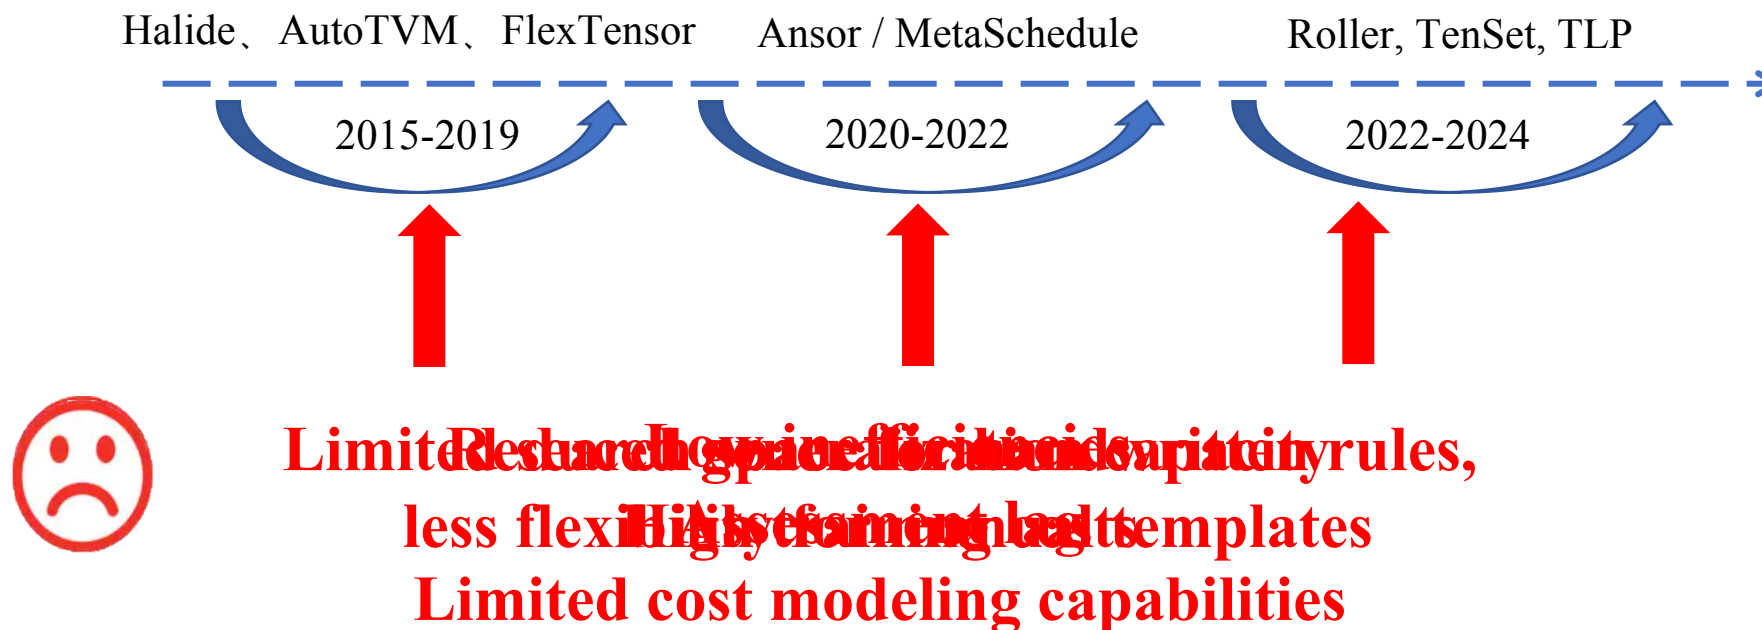
Background

Related SoA work



Background

Related SoA work



Background



Related SoA work



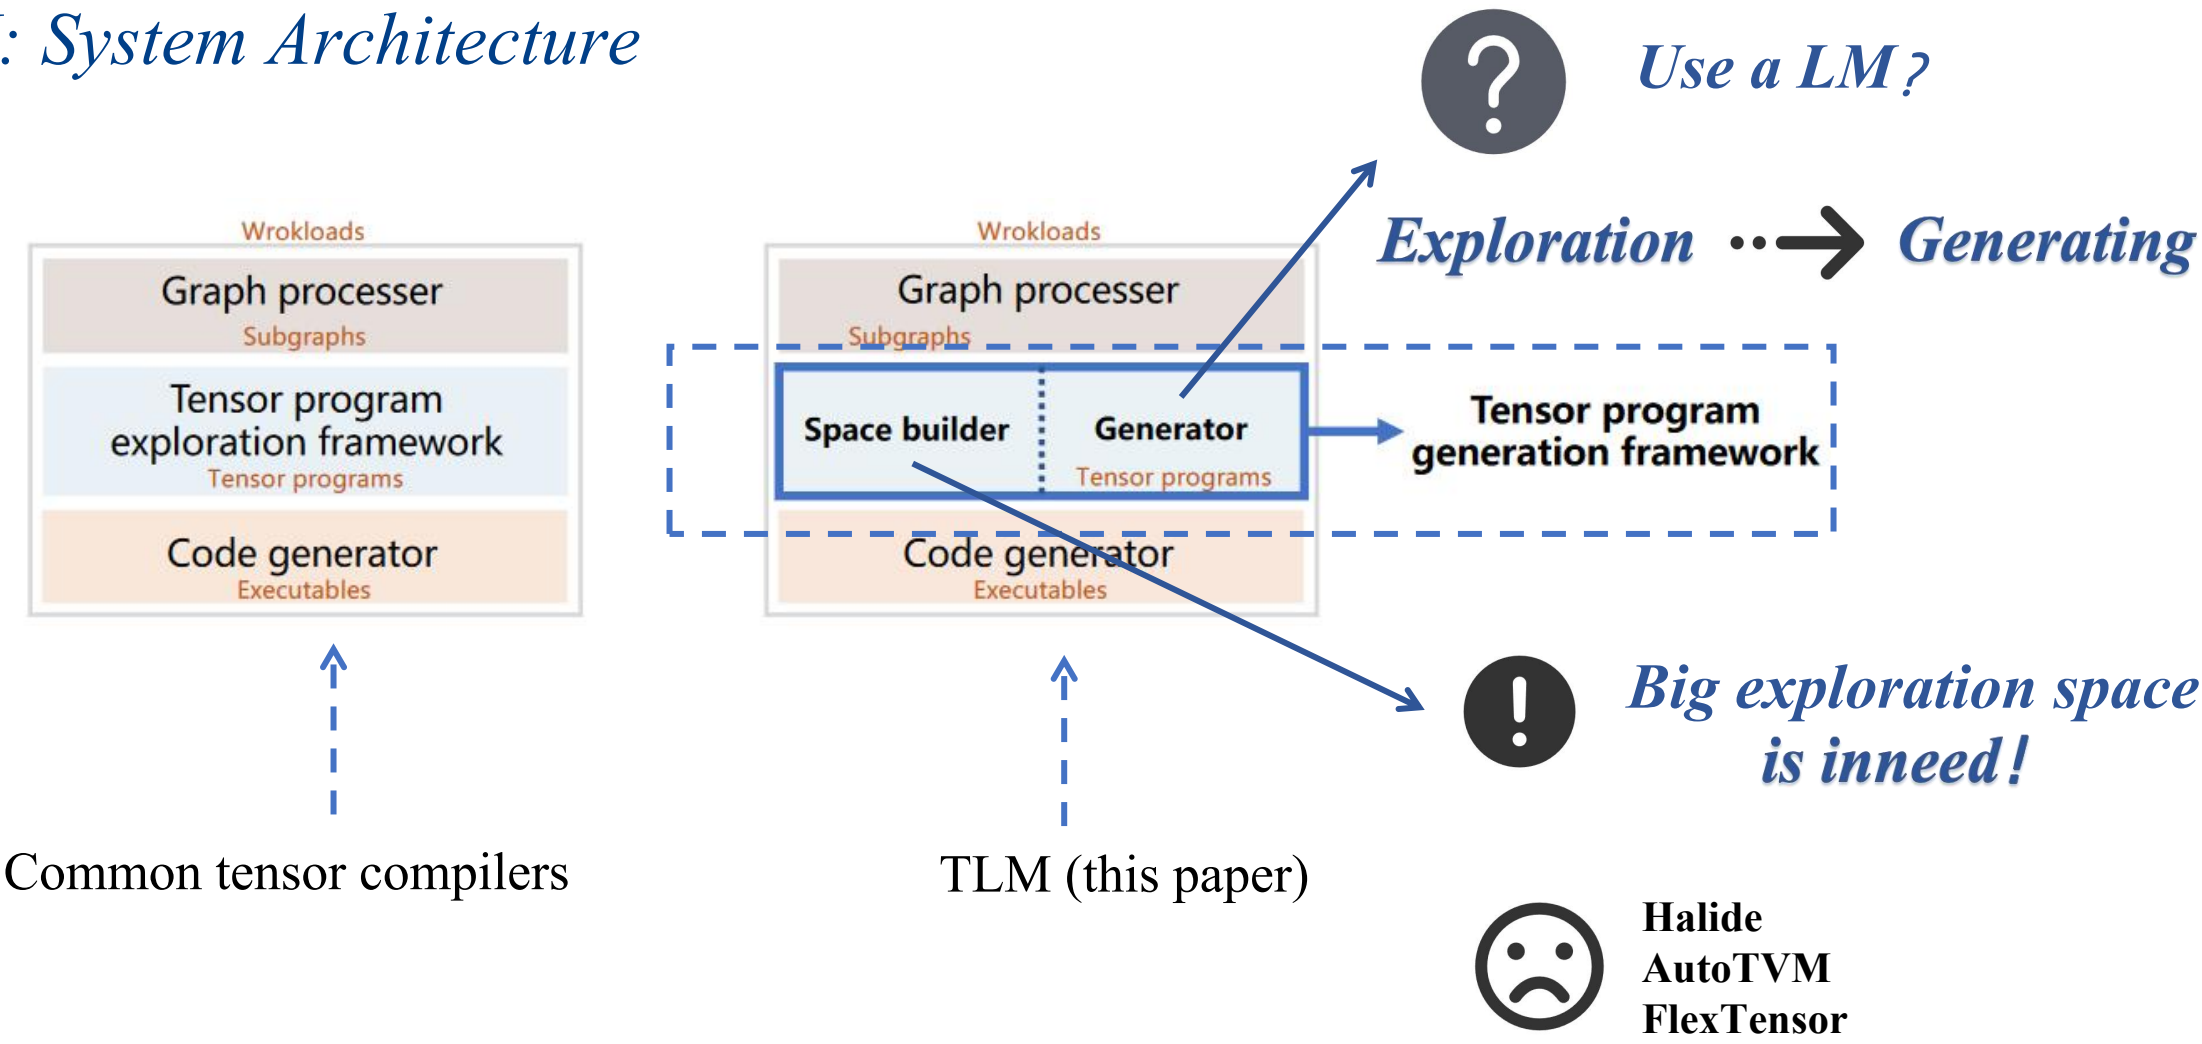
Design Concept	Exploration Orientation	Generation Orientation
Search Methods	Halide, AutoTVM (Templates、Heuristics) Roller (hardware alignment) , Ansor, MetaSchedule (random sampling)	-
Data-driven	TenSet, TLP (Offline Data + Search)	TLM (this paper)



- *Background*
- ***System Design***
- *Evaluation*
- *Conclusion*

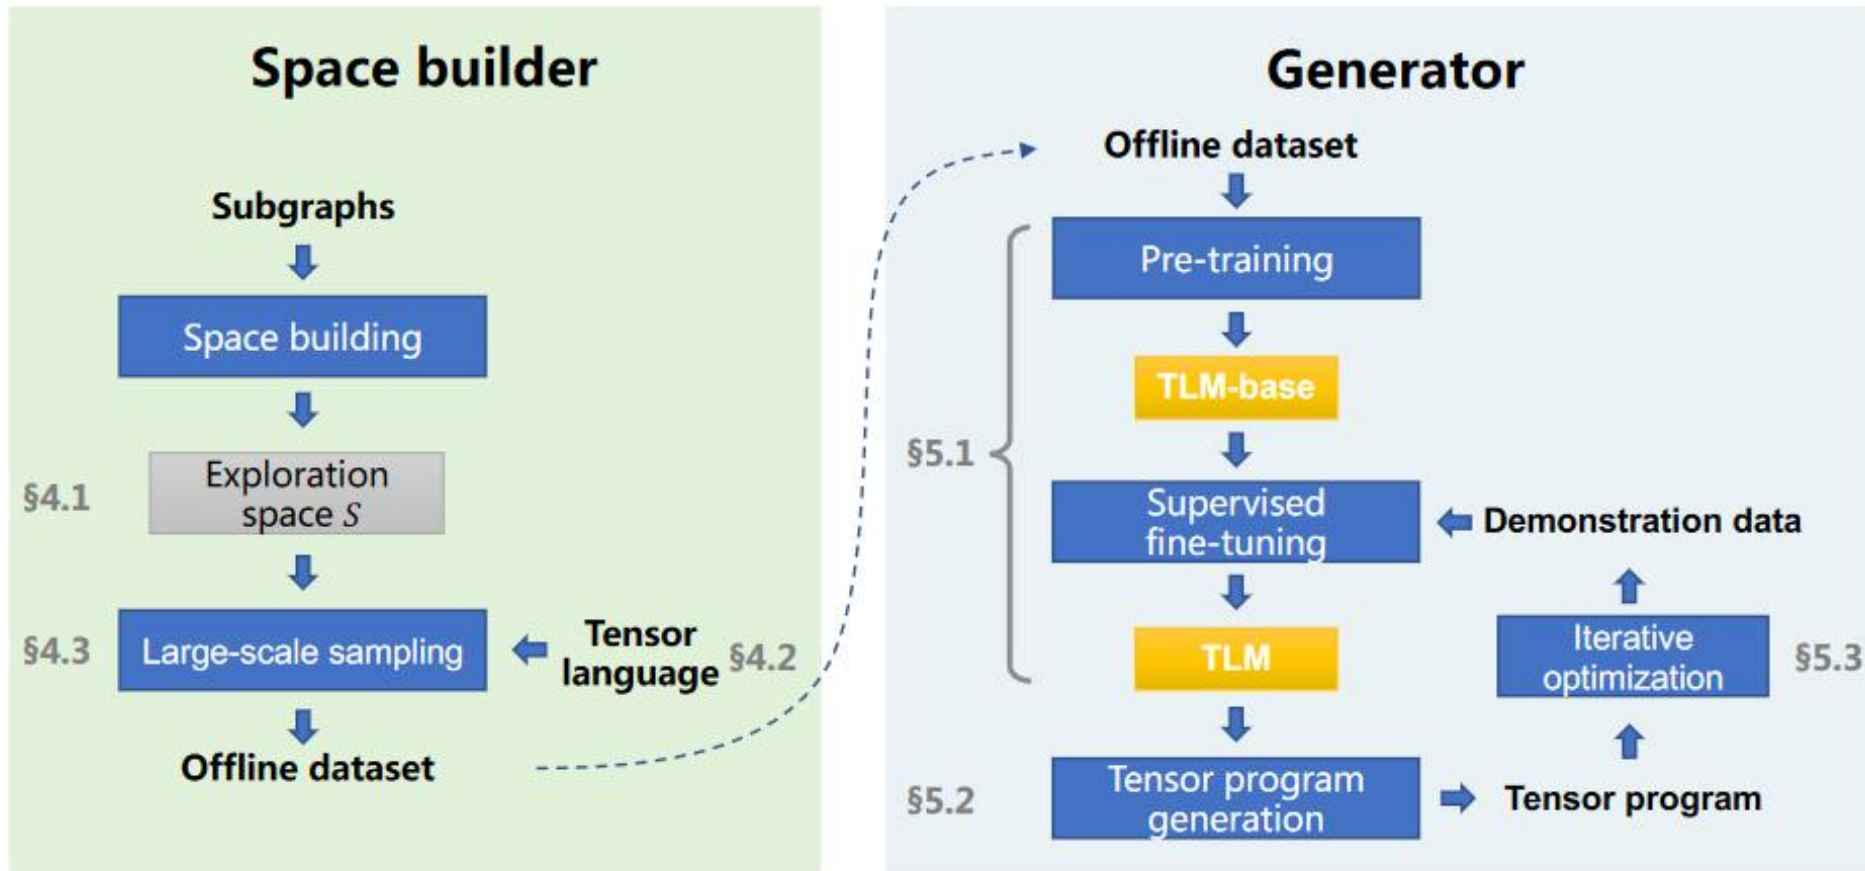
System Design

TLM: System Architecture

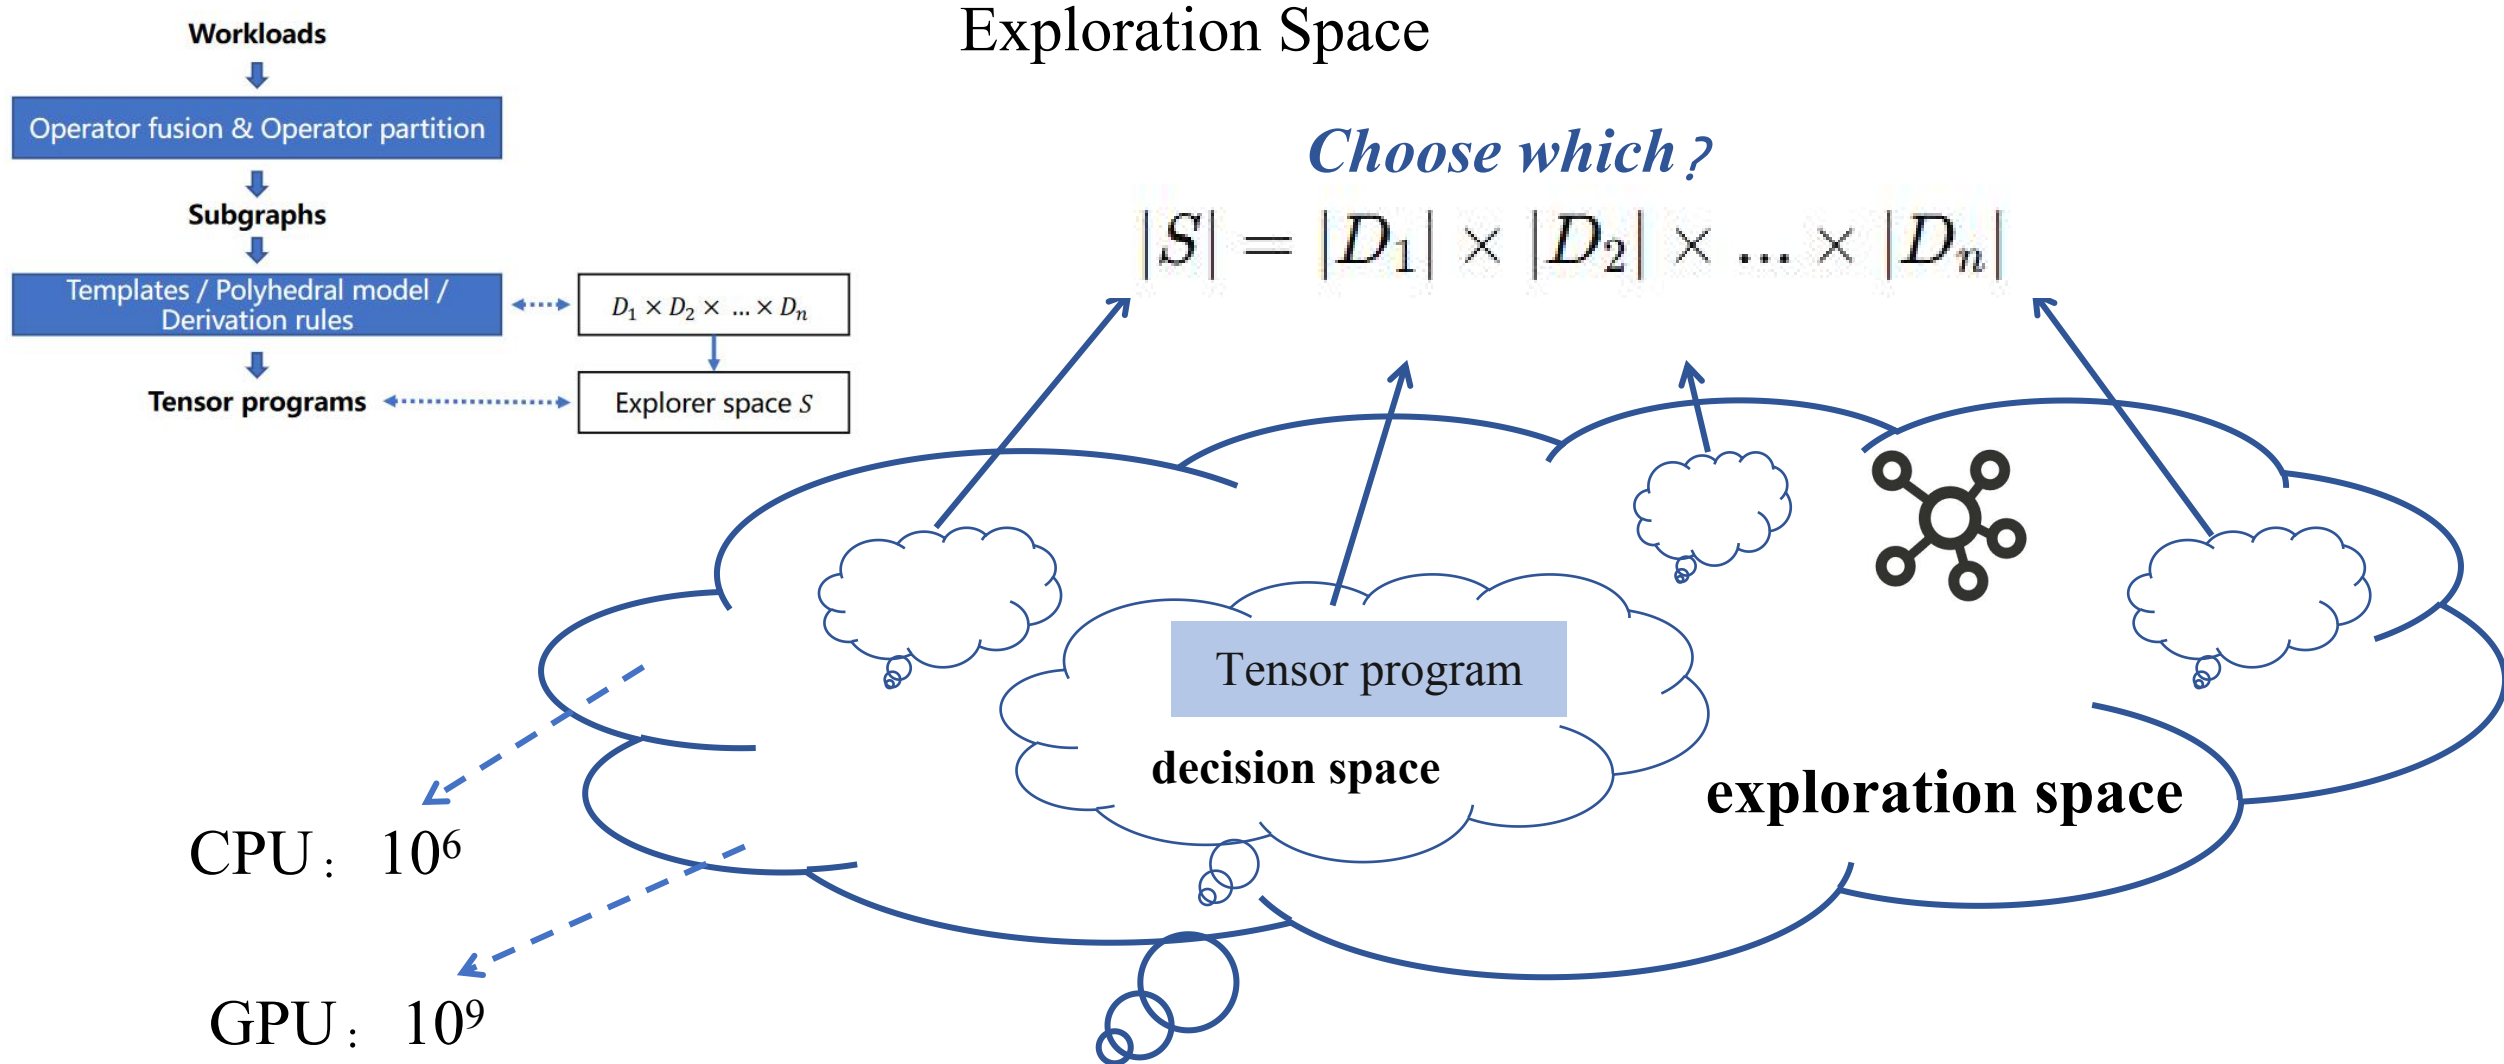


System Design

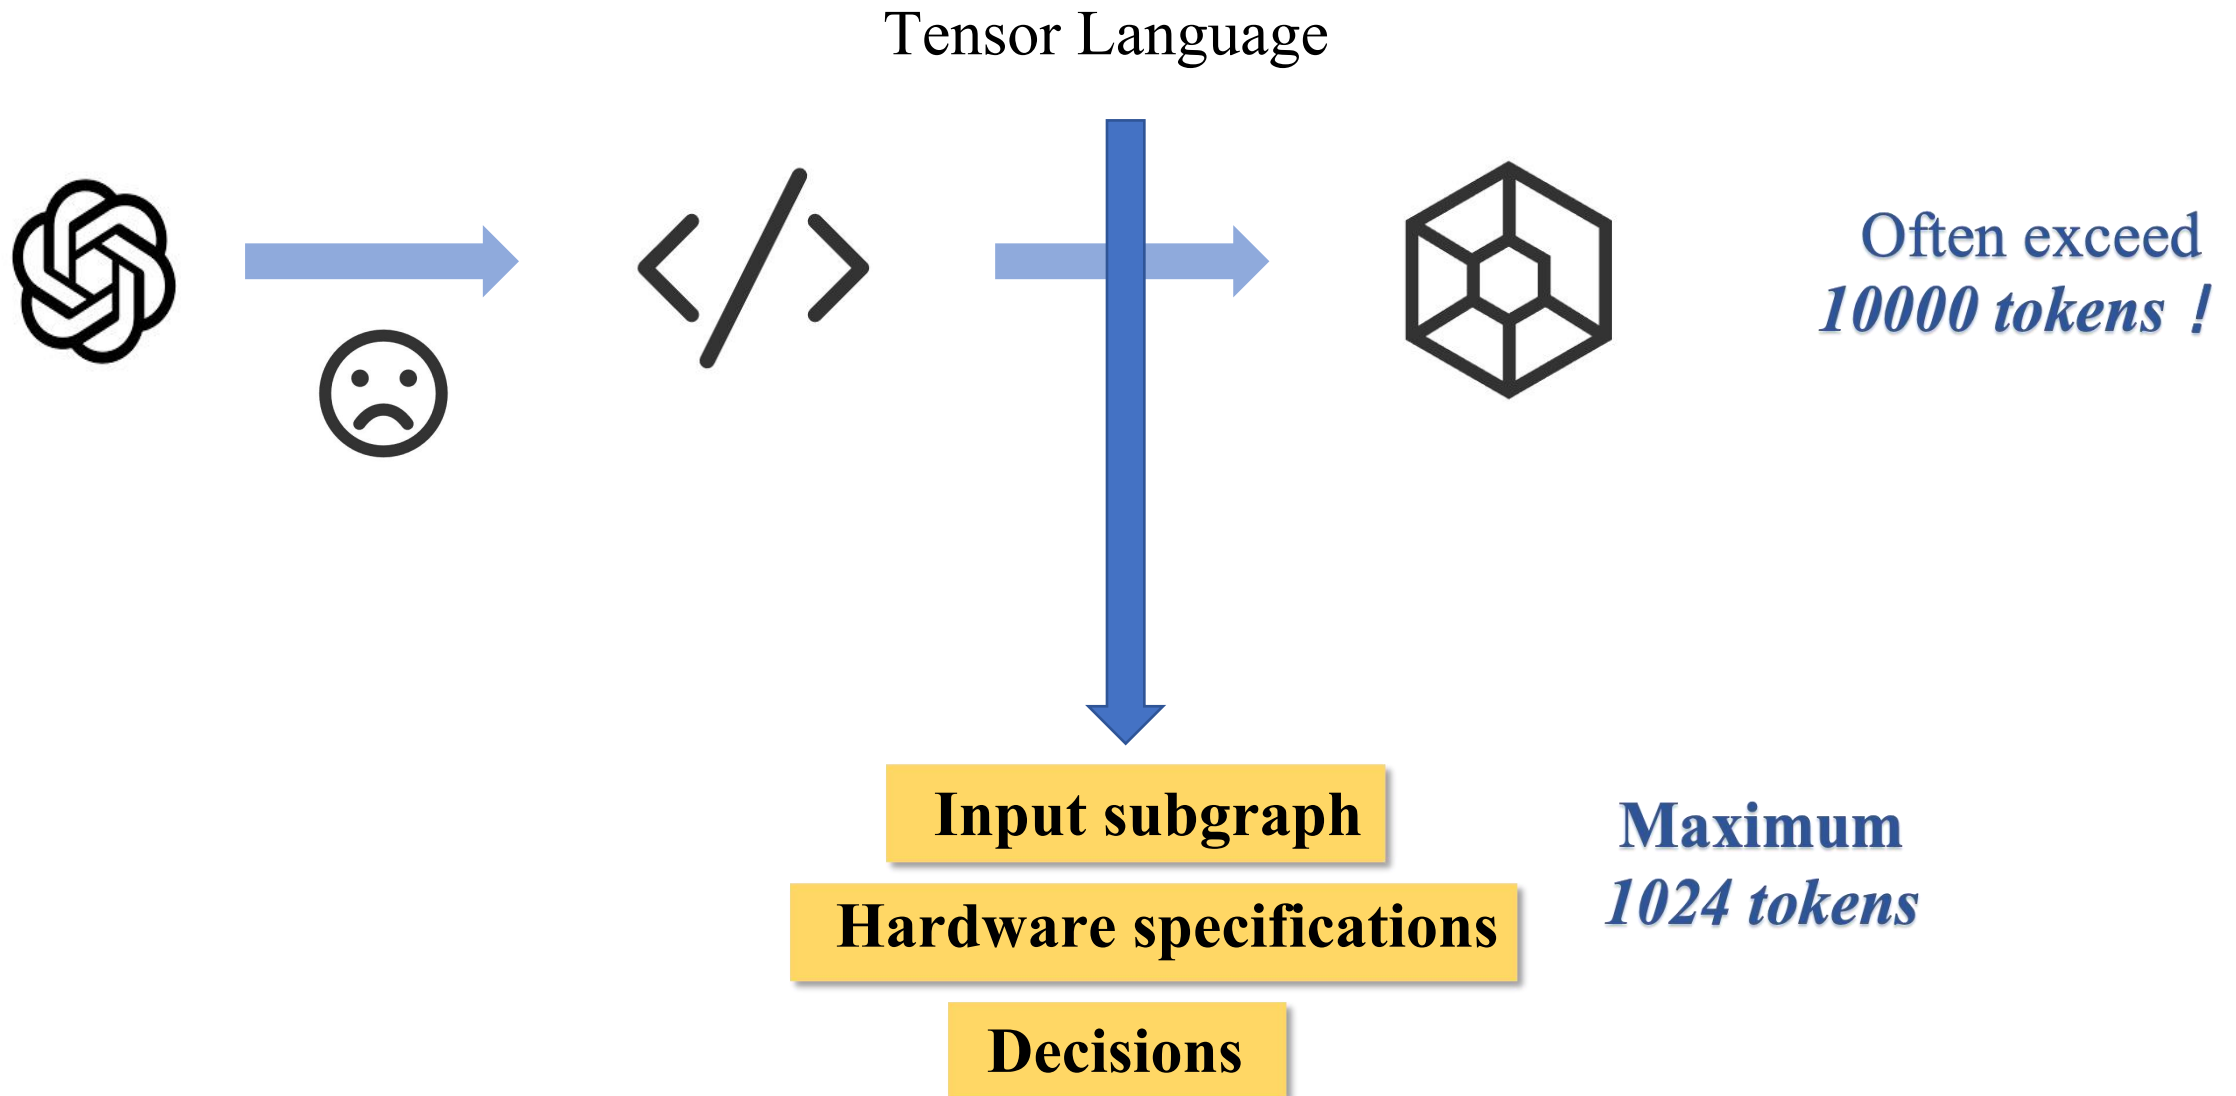
TLM: System Architecture



System Design



System Design





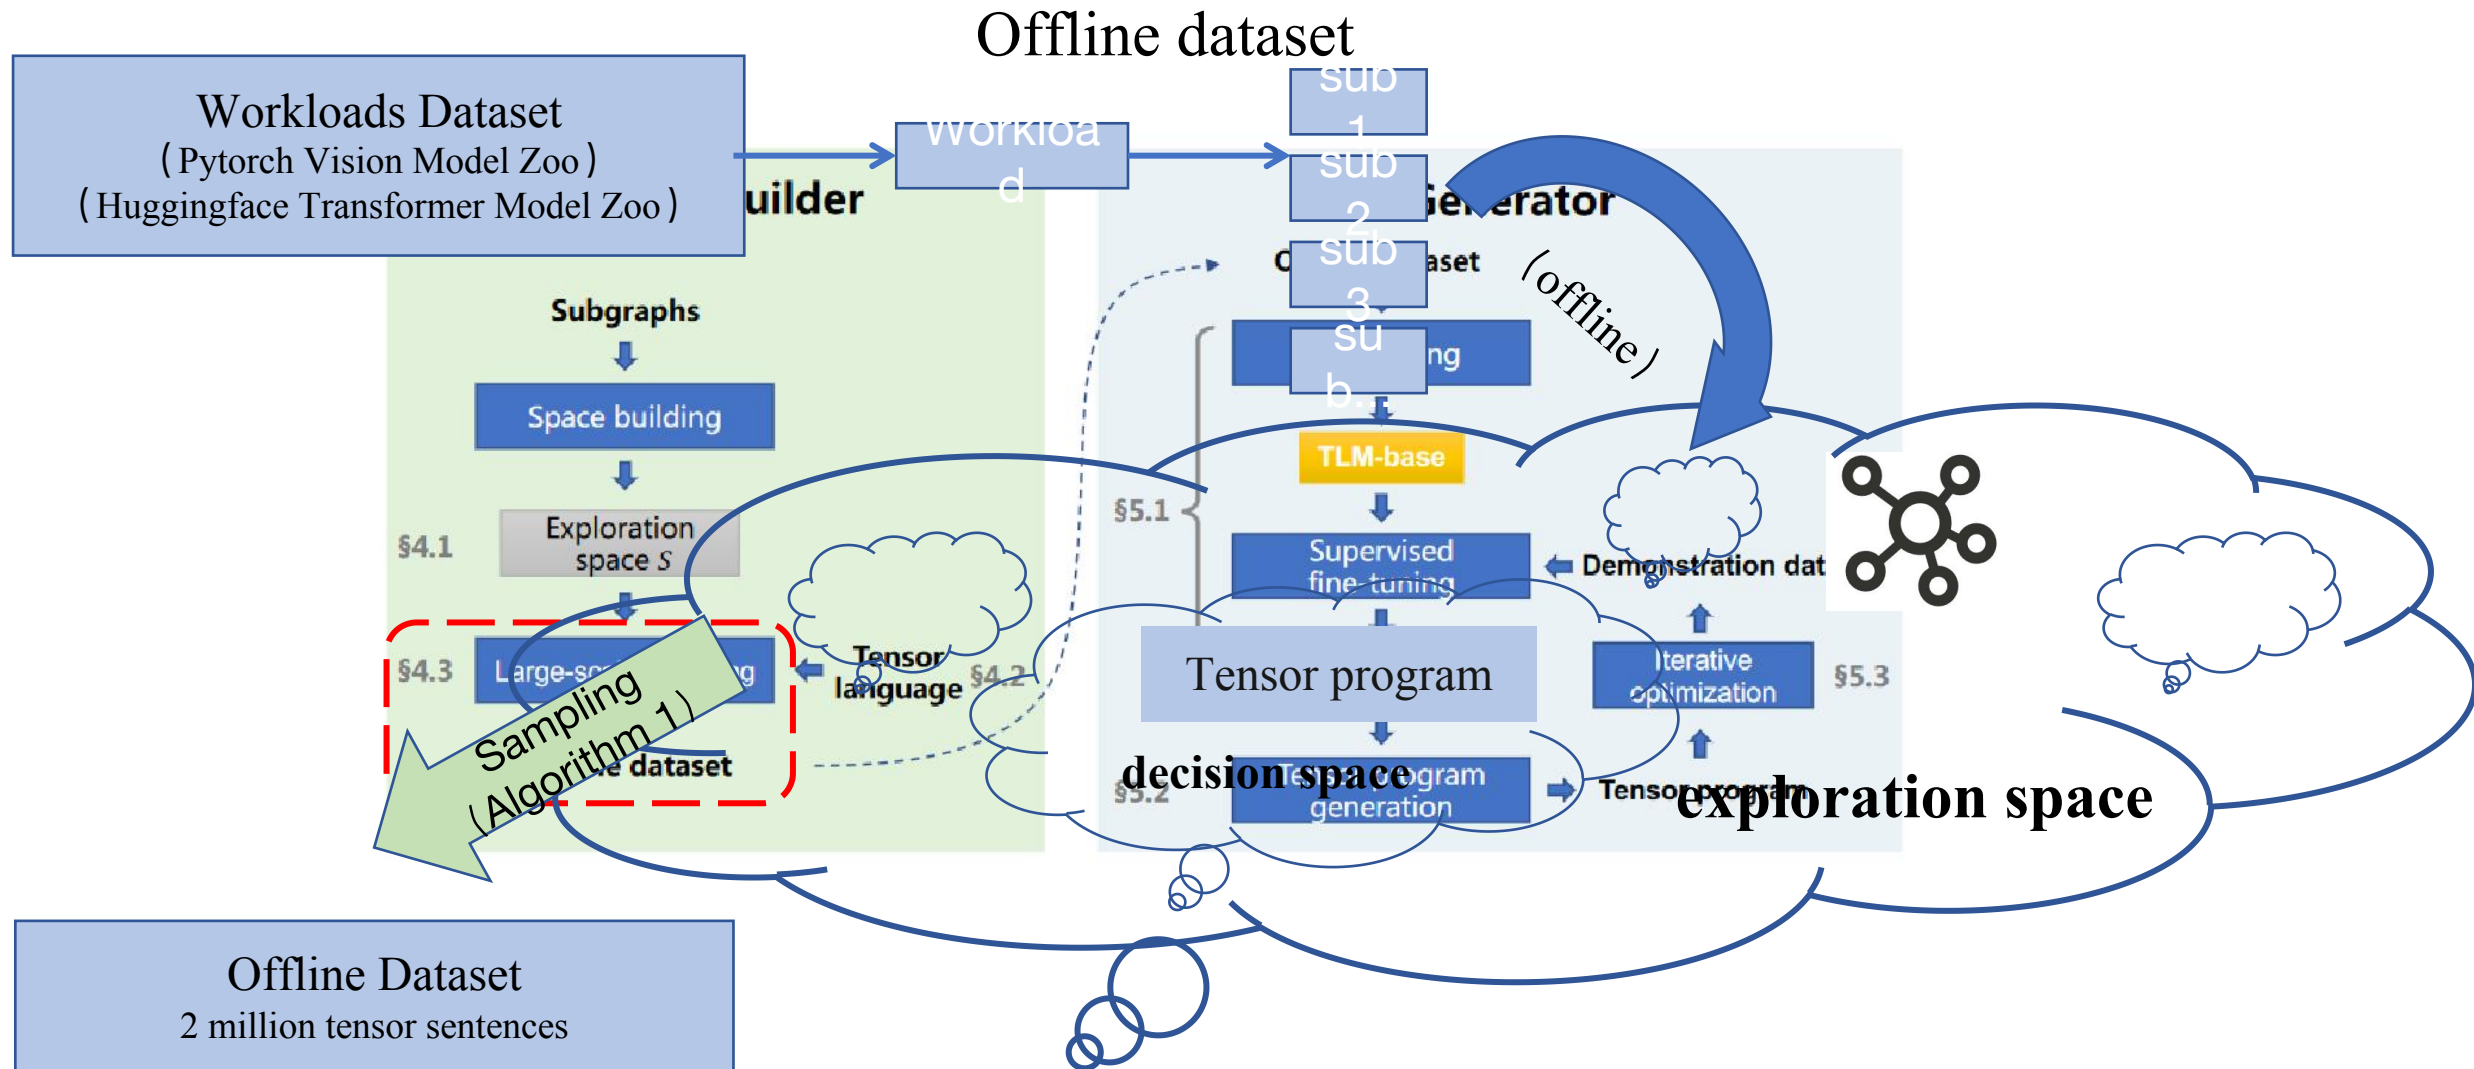
Tensor Language

Algorithm 1: Sampling tensor sentences from decision spaces.

```
1 Func GenerateSampleData(subgraph, hardware):
2   tokens = []
3   ExtractTokensFromSubgraph(subgraph, tokens)
4   ExtractTokensFromHardware(hardware, tokens)
5   decision_spaces = DetermineDecisionSpaces(subgraph,
6     hardware)
7   foreach space in decision_spaces do
8     switch space.type do
9       case "tile_size" do
10        | HandleTileSizeSpace(space, tokens)
11       case "unroll" do
12        | HandleParallelSpace(space, tokens)
13       // Additional space types
14       case ... do
15        | ...
16   return tokens
17
18 Func HandleTileSizeSpace(space, tokens):
19   tokens.append("split")
20   tokens.extend(Serialize(space.operator))
21   tokens.extend(Serialize(space.axis))
22   tiles = RandomSample(space)
23   tokens.extend(Serialize(tiles))
24   // Other properties
```

p0 p1 T_matmul_NT p2 T_add 00a059b856ac30ac172b6252254479a6 1024 1024 512 1024 1024 512
1024 512 llvm -keys=cpu -mcpu=core-avx2 -model=i7 4 64 64 0 0 0 0 2 SP 2 0 1024 32 1 4 1 SP 2 4
512 8 1 4 1 SP 2 8 1024 1024 1 RE 2 0 4 1 5 8 2 6 9 3 7 FSP 4 0 0 2 FSP 4 3 1 2 RE 4 0 3 1 4 2 5 CA 2 4 3 PPT
SPC 2 0 1024 32 1 4 1 SPC 2 4 512 8 1 4 1 SPC 2 8 1024 1024 1 CLS FU 4 0 1 2 3 AN 4 0 3 PRS 2 PR 2 0
auto_unroll_max_step\$0 VEC\$

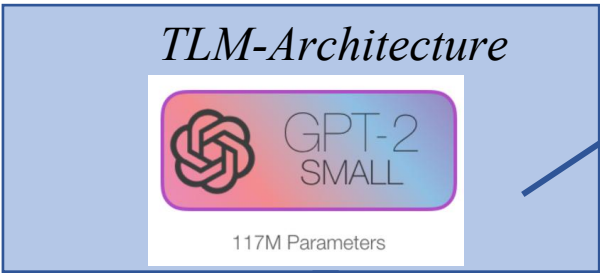
System Design



System Design

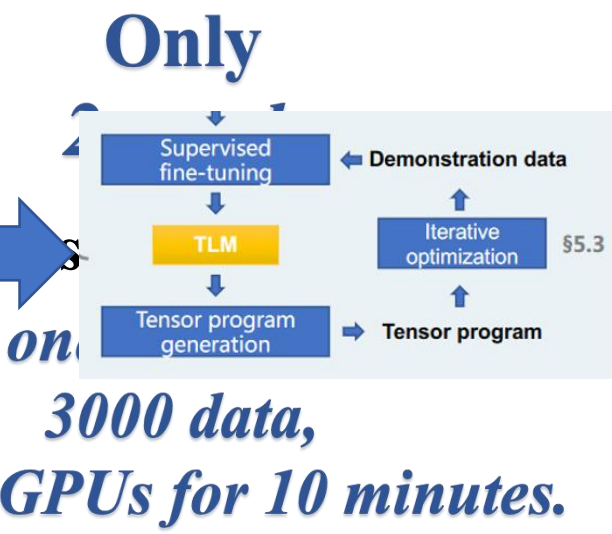
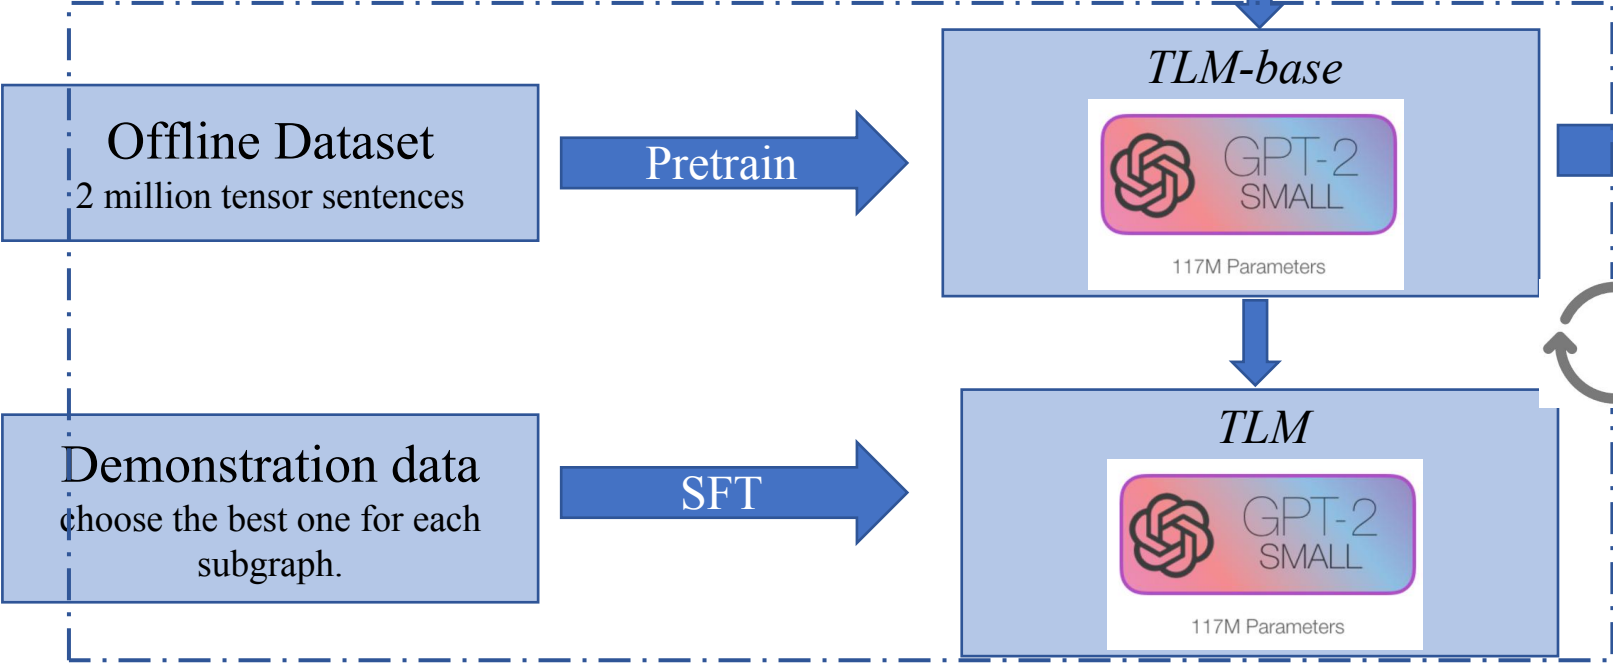
Inspired by [ChatGPT training process](#)

Model Details



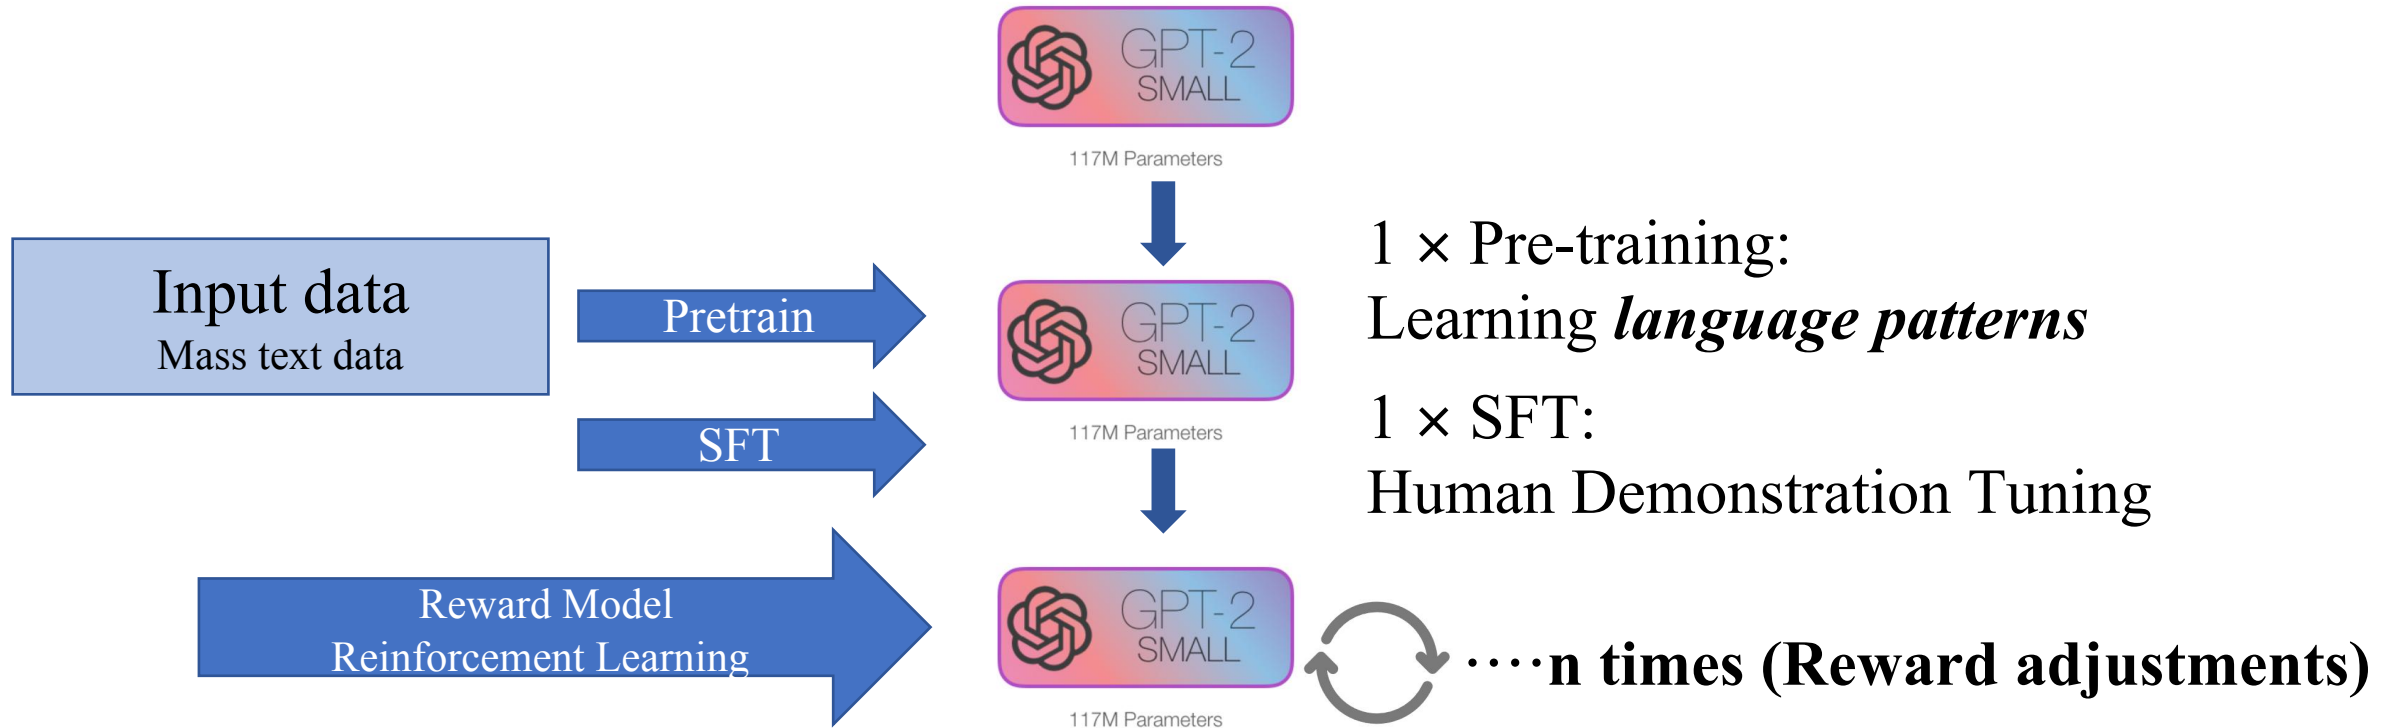
12 Transformer layers

12 attention heads
768 hidden cells



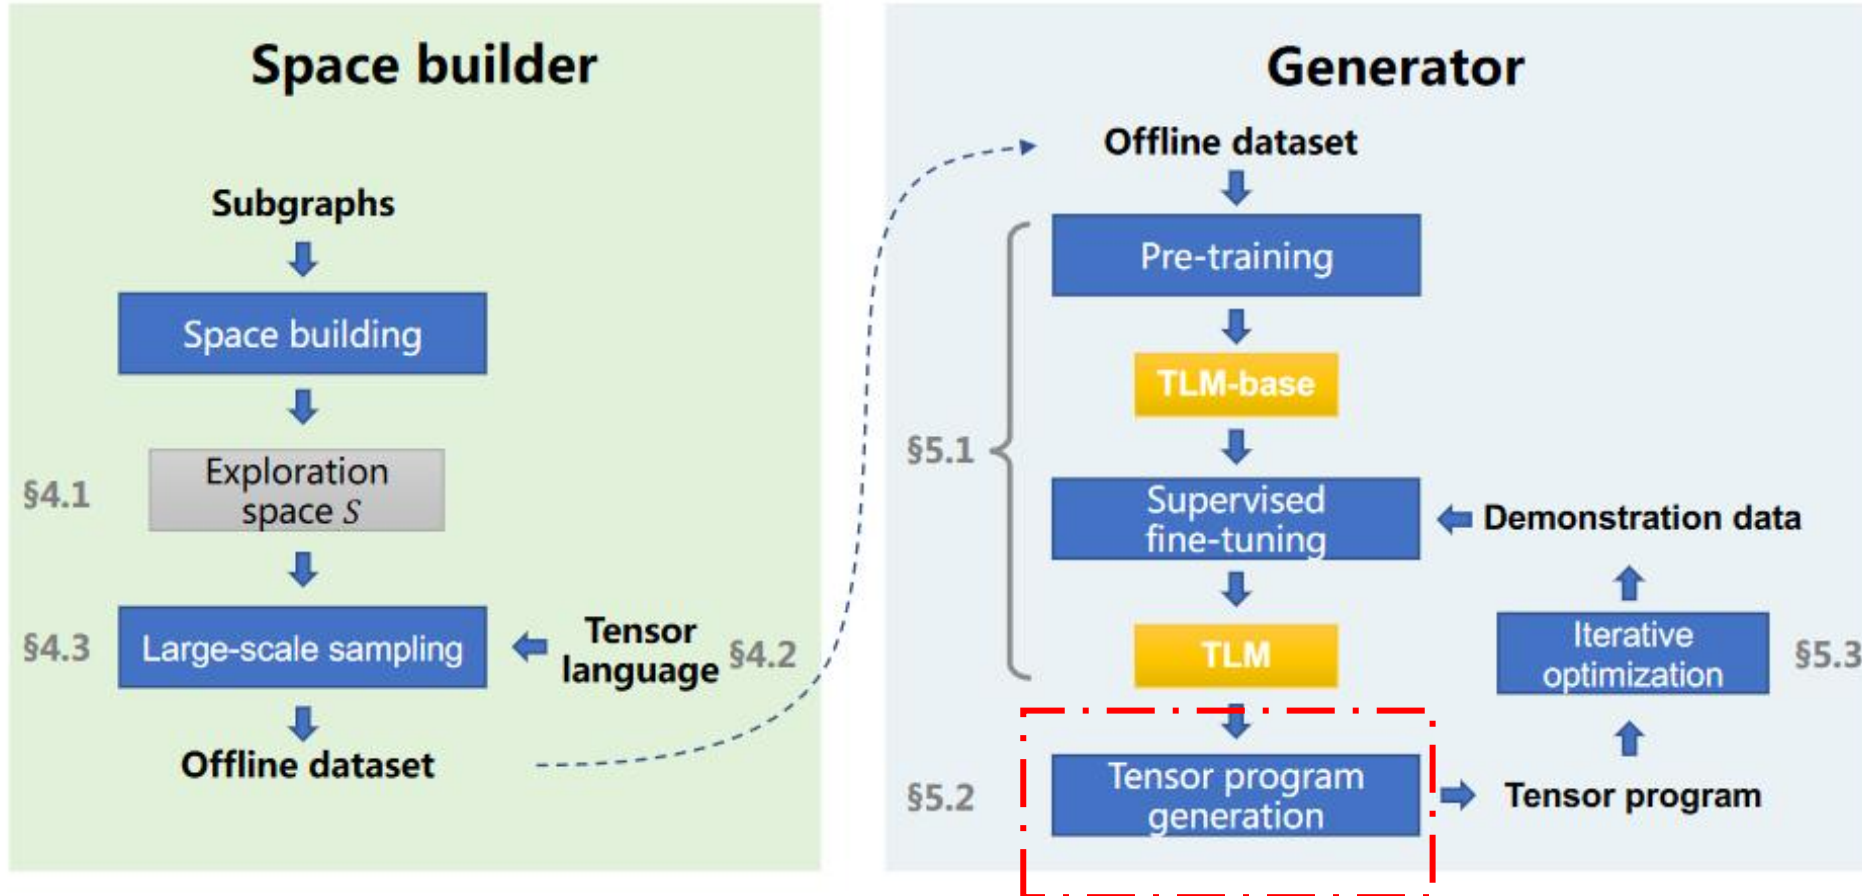
System Design

ChatGPT training process



System Design

Tensor Program Generation



Tensor Program Generation

Extract token from subgraph and hardware
initialize program



Traversing the decision space



Each time TLM generates a decision, such as
“tile size = 32”, the framework checks if it is
valid!



Algorithm 2: Generating tensor programs aided by TLM in decision-making.

```
1 Func GenerateTensorProgram(subgraph, hardware):  
2   tokens = []  
3   ExtractTokensFromSubgraph(subgraph, tokens)  
4   ExtractTokensFromHardware(hardware, tokens)  
5   program = GetInitProgram(subgraph, hardware)  
6   decision_spaces = DetermineDecisionSpaces(subgraph,  
7     hardware)  
8   foreach space in decision_spaces do  
9     switch space.type do  
10      case "tile_size" do  
11        ApplyTileSize(space, tokens, program)  
12      case "unroll" do  
13        ApplyParallel(space, tokens, program)  
14      // Additional space types  
15      case ... do  
16        ...  
17   return program  
18 Func ApplyTileSize(space, tokens, program):  
19   tokens.append("split")  
20   tokens.extend(Serialize(space.operator))  
21   tokens.extend(Serialize(space.axis))  
22   response_tokens = TLM(tokens)  
23   tiles = ConvertTokensToTiles(response_tokens)  
24   if not CheckValidTiles(space, tiles) then  
25     raise Exception("Invalid Tensor Program")  
26   tokens.extend(Serialize(tiles))  
27   program.apply(space.operator, space.axis, tiles)  
28   // Other properties
```

Tensor Program Generation

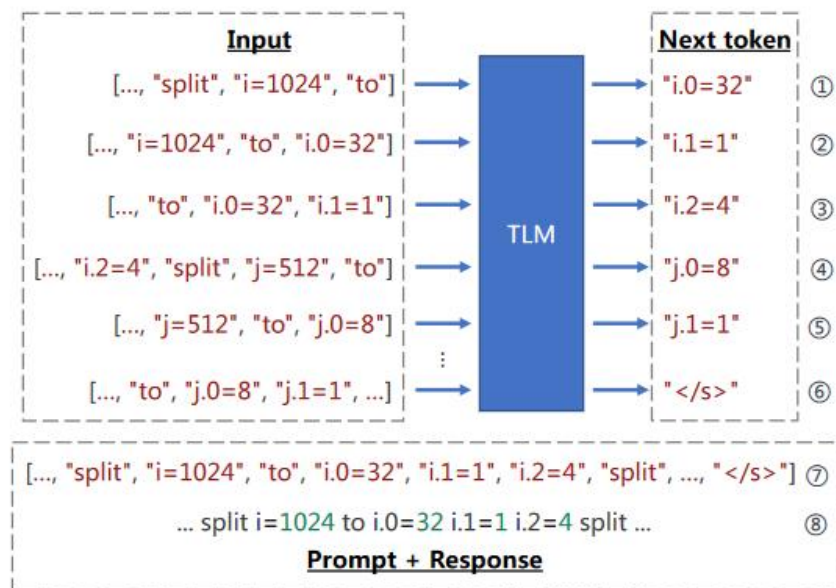


Figure 7: Generating a tensor sentence for a matrix multiplication operator with dimensions $m = 1024$, $n = 512$, and $k = 1024$, with the tiling size component depicted therein.

p0 p1 T.matmul NT p2 T.add 00a059b856ac30ac172b6252254479a6 1024 1024 512 1024 1024 512
1024 512 llvm -keys=cpu -mcpu=core-avx2 -model=i7 4 64 64 0 0 0 0 2 SP 2 0 1024 32 1 4 1 SP 2 4
512 8 1 4 1 SP 2 8 1024 1024 1 RE 2 0 4 1 5 8 2 6 9 3 7 FSP 4 0 0 2 FSP 4 3 1 2 RE 4 0 3 1 4 2 5 CA 2 4 3 PPT
SPC 2 0 1024 32 1 4 1 SPC 2 4 512 8 1 4 1 SPC 2 8 1024 1024 1 CLS FU 4 0 1 2 3 AN 4 0 3 PRS 2 PR 2 0
auto_unroll_max_step\$0 VEC\$

Algorithm 2: Generating tensor programs aided by TLM in decision-making.

```

1 Func GenerateTensorProgram(subgraph, hardware):
2     tokens = []
3     ExtractTokensFromSubgraph(subgraph, tokens)
4     ExtractTokensFromHardware(hardware, tokens)
5     program = GetInitProgram(subgraph, hardware)
6     decision_spaces = DetermineDecisionSpaces(subgraph,
7         hardware)
8     foreach space in decision_spaces do
9         switch space.type do
10             case "tile_size" do
11                 ApplyTileSize(space, tokens, program)
12             case "unroll" do
13                 ApplyParallel(space, tokens, program)
14             // Additional space types
15             case ... do
16                 ...
17     return program
18
19 Func ApplyTileSize(space, tokens, program):
20     tokens.append("split")
21     tokens.extend(Serialize(space.operator))
22     tokens.extend(Serialize(space.axis))
23     response_tokens = TLM(tokens)
24     tiles = ConvertTokensToTiles(response_tokens)
25     if not CheckValidTiles(space, tiles) then
26         raise Exception("Invalid Tensor Program")
27     tokens.extend(Serialize(tiles))
28     program.apply(space.operator, space.axis, tiles)
29     // Other properties

```



Iterative optimization

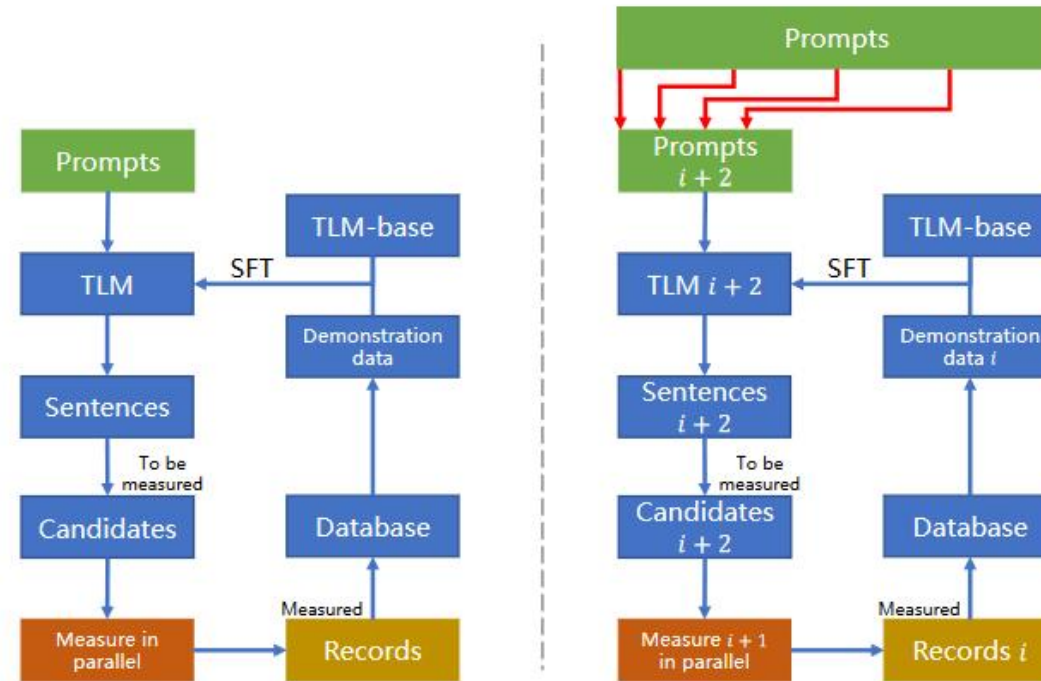
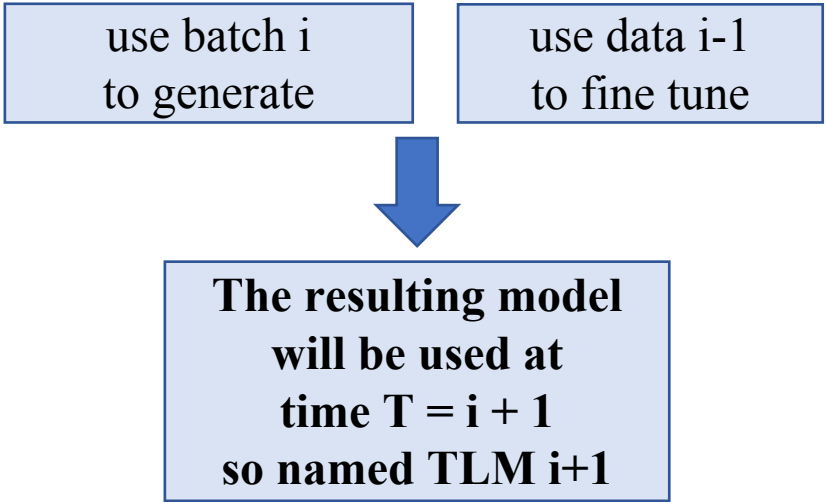
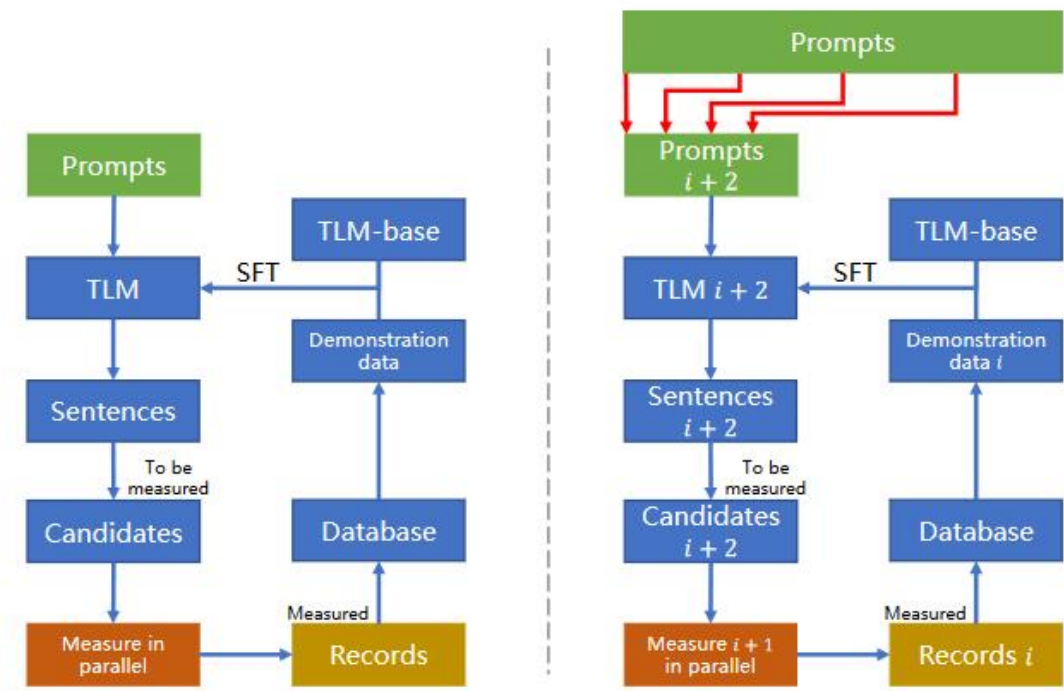


Figure 8: Flowchart of the iterative optimization process.

Why “ $i + 2$ ”
instead of “ $i + 1$ ” or “ i ”?



$T = i$



时间	测量进程 (左)	训练进程 (右)	TLM
T=0	测量 Batch 0, 获得示例数据 0	训练还没开始	TLM-base
T=1	测量 Batch 1, 获得示例数据 1	用示例数据 0 训练, 得到 TLM 2	TLM 2
T=2	测量 Batch 2, 获得示例数据 2	用示例数据 1 训练, 得到 TLM 3	TLM 3
T=3	测量 Batch 3, 获得示例数据 3	用示例数据 2 训练, 得到 TLM 4	TLM 4



Iterative optimization

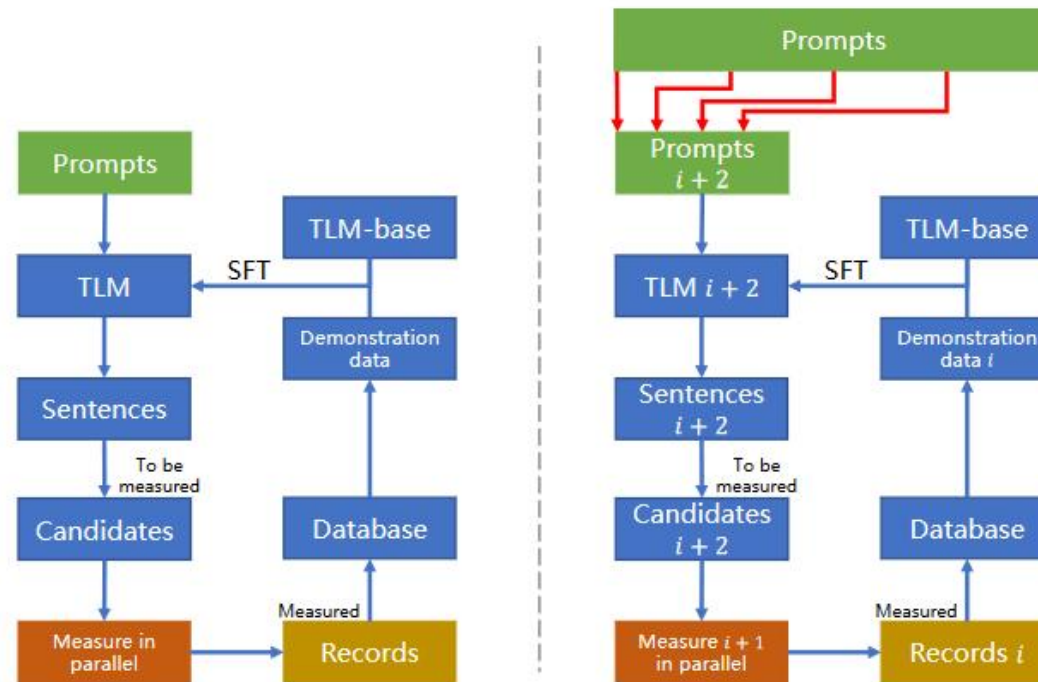


Figure 8: Flowchart of the iterative optimization process.



- *Background*
- *System Design*
- ***Evaluation***
- *Conclusion*

Evaluation

Experimental Settings

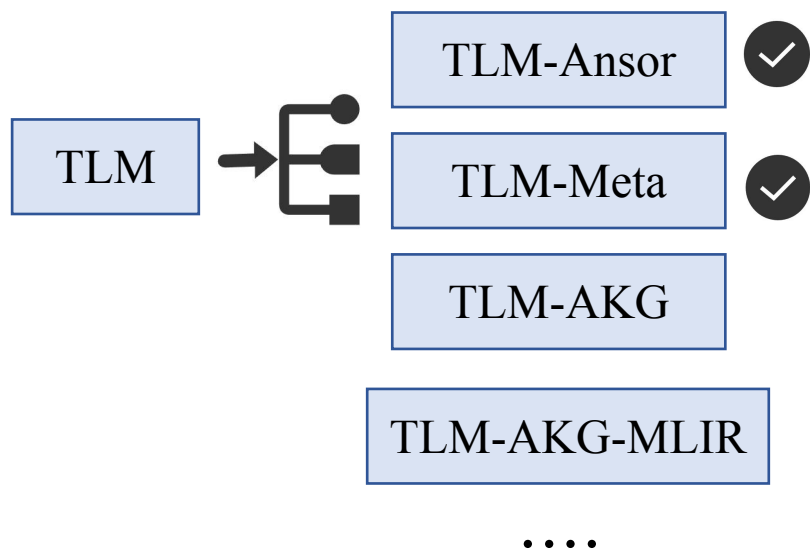


Table 1: Workloads in the TLM test set.

Model	Input shape	Model	Input shape
ResNet-50 [32]	[1, 3, 224, 224]	DenseNet-121 [33]	[8, 3, 256, 256]
MobileNet-V2 [34]	[1, 3, 224, 224]	BERT-large	[4, 256]
ResNeXt-50 [35]	[1, 3, 224, 224]	Wide-ResNet-50 [36]	[8, 3, 256, 256]
BERT-base [37]	[1, 128]	ResNet3D-18 [38]	[4, 3, 144, 144, 16]
BERT-tiny	[1, 128]	DCGAN [39]	[8, 3, 64, 64]
GPT-2	[1, 128]	LLAMA [40]	[4, 256]

CPU: 4-core Intel i7-10510U, 16GB RAM, AVX2 support, Ubuntu 20.04.

GPU: 48-core Intel Xeon Gold 6226, 376GB RAM, 4 x 32GB NVIDIA V100, Ubuntu 20.04, CUDA 11.6, cuDNN 8.4.0.



Convergence Behavior of Demonstration Data

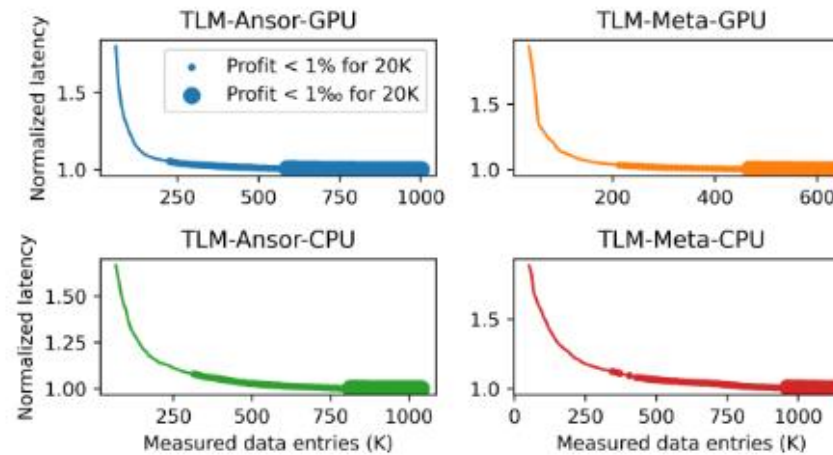


Figure 9: Demonstration data convergence curves of TLM-Ansor and TLM-Meta on the GPU and the CPU.

Convergence : as a performance improvement of less than 1% after measuring 20,000 pieces of data (all subgraphs).



Subgraph Benchmark

Table 2: The overall speedup for the 23 TLM-Ansor subgraphs. The higher the overall speedup, the better. In the table, "Times" represents the measurement times for each subgraph.

	Times	Ansor			TLM-Ansor
		64	1K	10K	10K
TLM-Ansor	1	1.26	0.98	0.92	0.85
	10	1.40	1.08	1.03	0.95
	16	1.43	1.10	1.04	0.96
	32	1.45	1.12	1.06	0.98
	64	1.45	1.12	1.06	0.98
	1K	1.46	1.13	1.07	0.99
	10K	1.48	1.14	1.08	1.00
Ansor	10K	1.37	1.06	1.00	0.92

Table 3: The overall speedup for the 40 TLM-Meta subgraphs.

	Times	MetaSchedule			TLM-Meta
		64	1K	10K	1K
TLM-Meta	1	1.00	0.69	0.68	0.67
	10	1.41	0.96	0.95	0.94
	16	1.45	1.00	0.99	0.97
	32	1.46	1.01	1.00	0.97
	64	1.49	1.02	1.01	0.99
	1K	1.50	1.02	1.01	1.00
MetaSchedule	10K	1.48	1.01	1.00	0.99



End-to-End Workload Benchmark

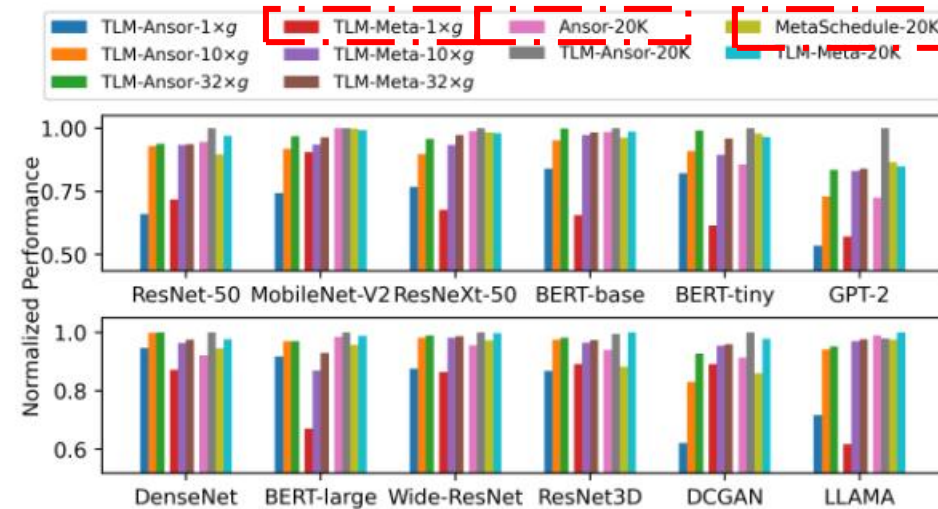


Figure 10: Workload inference performance comparison with Ansor/MetaSchedule on V100.



Comparison with TensorRT/PyTorch on V100

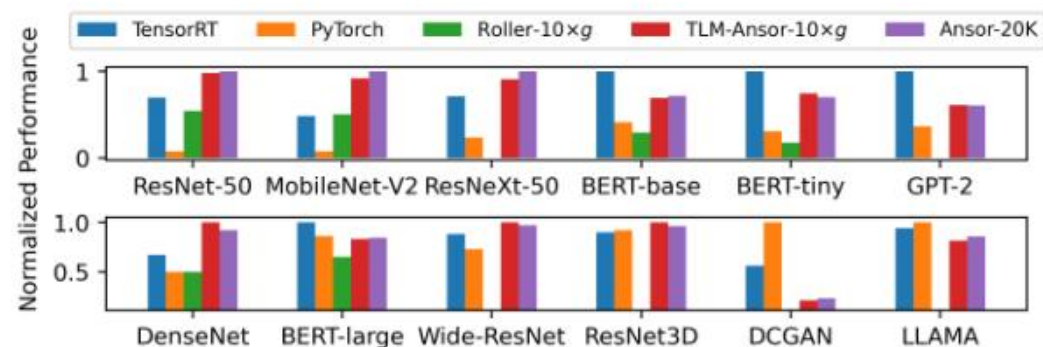


Figure 13: Workload inference performance comparison with TensorRT/PyTorch on V100.

Convergence : as a performance improvement of less than 1% after measuring 20,000 pieces of data (all subgraphs).



- *Background*
- *System Design*
- *Evaluation*
- ***Conclusion***



Enabling Tensor Language Model to Assist in Generating High-Performance Tensor Programs for Deep Learning

- ▶ Finding new scenarios to introduce LM.
- ▶ An LM was trained to apply the new scenario.
- ▶ To accompany this, a tensor language was designed and very solid framework work was done.
- ▶ Innovative optimization for very time-consuming traditional scenarios
- ▶ In the experimental part, the validation was done mainly for the **efficiency**, **low demand**, and **low training time** mentioned.

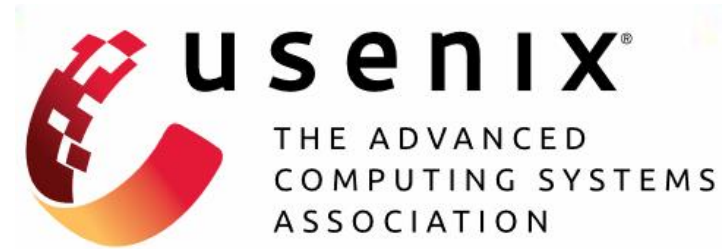


➤ 这个paper有什么问题，基于这个paper还能做什么？

- 硬件适配不足：
 - TLM 在 BERT、LLAMA、DCGAN 等任务上落后于 TensorRT/PyTorch，因其生成的张量程序未充分利用硬件特性（如 V100 的 TensorCore、HBM2 带宽）（在矩阵乘法和卷积主导的任务中，TLM 延迟高，竞争力不足。）
 - 【如果参考类似TLM的构造，可以将硬件特性融入张量句子生成，引入硬件感知损失函数？等方面作为一个创新点】
- 实验没有考虑子图间依赖，都是随机挑选种类。
 - 【如果参考TLM的设计的话，还可以，扩展 TLM 预测全局调度策略（？），优化子图间通信（GPU 内存拷贝、共享）。
 - 对于subgraph 引入子图依赖图（dependency graph），TLM 预测整体延迟最优配置。



- 这个paper提到的idea，能不能用在自己的方向/project上面？
 - 文章用到的都是 V100 GPU 和 Intel CPU，没有测试边缘设备，也没有测试过其他硬件（如 NPU、ARM 架构的 CPU）
 - 也许可以优化张量语言，适配 ARM 或 RISC-V 架构。
 - 在边缘场景，也许可以利用 $1\times g$ 测量的特性，减少边缘编译时间，支持低功耗推理。
 - 测评指标也许可以加入功耗指标。
- 这个paper能不能泛化？
 - 对于端侧设备来说，尤其是agent，可能优化空间不多。但如果对于具身智能来说，可能需要处理多任务，那么也许可以设计一个多任务张量调度生成框架，统一优化张量程序生成、任务调度和内存分配。
 - 针对 TLM 的离线数据依赖（L1），设计在线自适应机制，利用端侧运行时反馈微调模型。



Enabling Tensor Language Model to Assist in Generating High-Performance Tensor Programs for Deep Learning

Thank you !

Authors: Yi Zhai¹, Sijia Yang², Keyu Pan³, Renwei Zhang², Shuo Liu¹
Chao Liu and Zichun Ye², Jianmin Ji¹, Jie Zhao⁴,
Yu Zhang and Yanyong Zhang¹

Reporter: Hangshuai He

¹ University of Science and Technology of China,

² Huawei Technologies Co., Ltd.

³ ByteDance Ltd.

⁴ Hunan University

Enabling Tensor Language Model to Assist in Generating High-Performance Tensor Programs for Deep Learning

Authors: Yi Zhai¹, Sijia Yang², Keyu Pan³, Renwei Zhang², Shuo Liu¹
Chao Liu and Zichun Ye², Jianmin Ji¹, Jie Zhao⁴,
Yu Zhang and Yanyong Zhang¹

Reporter: Hangshuai He

¹ University of Science and Technology of China,

² Huawei Technologies Co., Ltd.

³ ByteDance Ltd.

⁴ Hunan University

