

Autothrottle: A Practical Bi-Level Approach to Resource Management for SLO-Targeted Microservices

**Zibo Wang, Pinghe Li, Chieh-Jan Mike Liang,
Feng Wu, Francis Y. Yan**



**Presenter: Chuanhua Fan
2025.05.09**

Dictionary

- **Background**
- **Design**
- **Evaluation**
- **Related work**
- **Conclusion**
- **Think**

Dictionary

- **Background**
- Design
- Evaluation
- Related work
- Conclusion
- Think

Background

Achieving resource efficiency while preserving end-user experience is important for cloud application operators.

To ensure a seamless end-user experience, many user-facing latency-sensitive applications impose an SLO.

Background

之前的做法:

cloud application operators resort to resource over-provisioning to avoid SLO violations.

现在的做法:

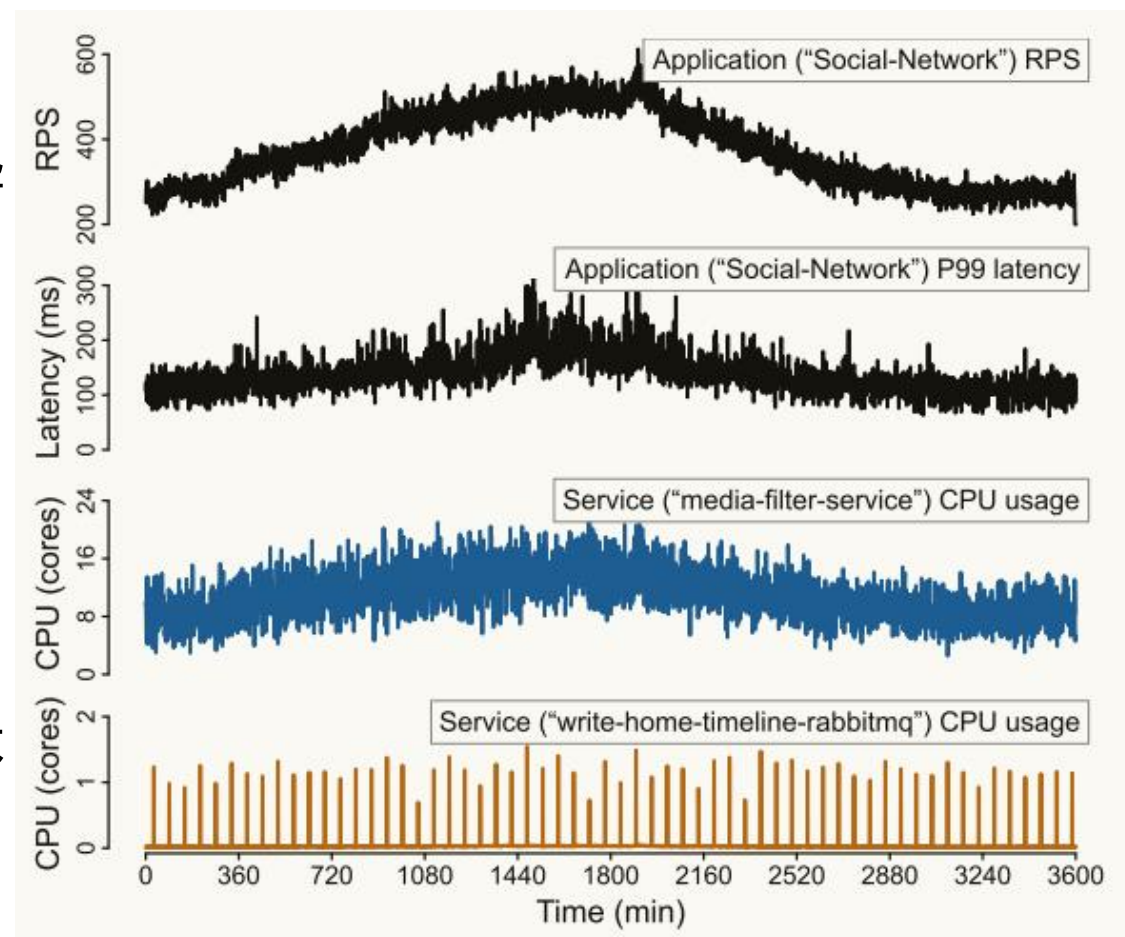
Recycling excess resources to save a significant amount of resources.

A key enabler for such resource saving is **SLO-targeted resource management**. Its **goal** is to continuously minimize the total resources allocated, while still satisfying the end-to-end latency SLO.

Background

The distributed nature of microservices has brought new difficulties to resource management:

- 1、由于不同的用户请求对每个服务的压力不同，异构服务可以表现出截然不同的资源使用模式。
- 2、应用程序性能和每个服务的资源使用情况是不同级别的度量，不一定表现出很强的相关性。
- 3、在观察分配变化对端到端性能的影响时会产生不希望的延迟。



Background

How does this article address these issues?

本文采用了分布式系统行为不同的级别，以及应用程序级SLO反馈和服务级资源控制的体系结构解耦机制，为慢速目标微服务设计了Autothrottle。

What is the goal?

最小化基于微服务的应用程序的总CPU分配，同时避免违反用户请求的延迟SLO。

Dictionary

- Background
- **Design**
- Evaluation
- Related work
- Conclusion
- Think

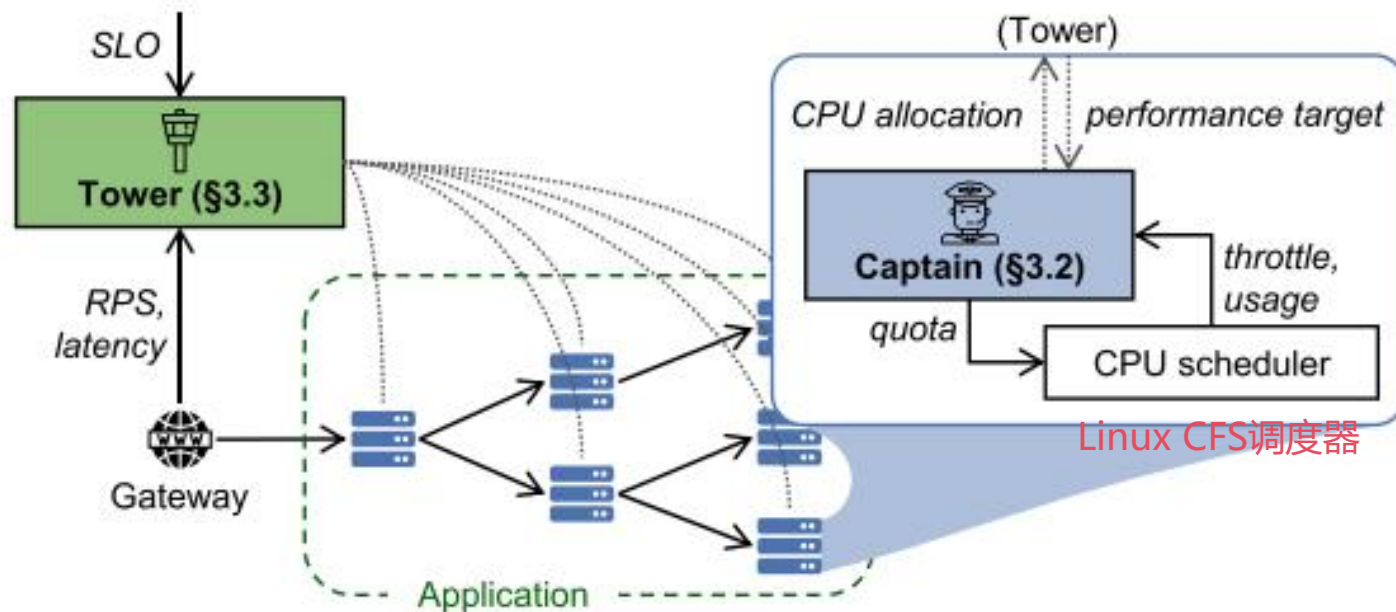
Design

Autothrottle框架

服务控制器——Captain

应用范围控制器——Tower

连接桥梁——CPU throttle ratios



Design

CPU Throttle Ratio (CPU节流比例) :

公式:

$$\text{CPU Throttle Ratio} = \frac{\Delta \text{nr_throttled}}{\text{Number of CFS Periods in Window}}$$

例:

假设在1秒（10个CFS周期）内，nr_throttled从50增加到55，则节流次数增量 $\Delta=5$ ，节流比例 = $5/10=0.55/10=0.5$ ，即50%的时间窗口内发生了节流。

Design

- Captain

1、Multiplicative scale-up

乘法放大，进一步使增量的大小与测量的CPU节流比和目标比率之间的差异成比例。

Algorithm 1: Captain: scaling up and down

```
1 /* executes every  $N$  periods */
2 throttleCount = throttle count during last  $N$  periods;
3 throttleRatio = throttleCount/ $N$ ;
4 margin = max(0, margin + throttleRatio - throttleTarget);
5 if throttleRatio >  $\alpha \times$  throttleTarget then
6   /* multiplicatively scale up */
7   | quota = quota  $\times$  (1 + throttleRatio -  $\alpha \times$  throttleTarget);
8 else
9   /* instantaneously scale down */
10  | history = CPU usage history in the last  $M$  periods;
11  | proposed = max(history) + margin  $\times$  stdev(history);
12  | if proposed  $\leq \beta_{\max} \times$  quota then
13  | | quota = max( $\beta_{\min} \times$  quota, proposed);
14  | end
15 end
```

Design

- **Captain**

1、Multiplicative scale-up

通过调整 α ，管理人员可以平衡以下两个目标：

- 快速响应真实高负载 ($\alpha < 1$)
- 避免错误操作扩大CPU配额 ($\alpha > 1$)

Algorithm 1: Captain: scaling up and down

```
1  /* executes every  $N$  periods */
2  throttleCount = throttle count during last  $N$  periods;
3  throttleRatio = throttleCount/ $N$ ;
4  margin = max(0, margin + throttleRatio - throttleTarget);
5  if throttleRatio >  $\alpha \times$  throttleTarget then
6    /* multiplicatively scale up */
7    quota = quota  $\times$  (1 + throttleRatio -  $\alpha \times$  throttleTarget);
8  else
9    /* instantaneously scale down */
10   history = CPU usage history in the last  $M$  periods;
11   proposed = max(history) + margin  $\times$  stdev(history);
12   if proposed  $\leq \beta_{\max} \times$  quota then
13     quota = max( $\beta_{\min} \times$  quota, proposed);
14   end
15 end
```

Design

- Captain

2、Instantaneous scale-down

瞬时缩放，防止CPU配额突然发生很大的变化，避免在工作负载高峰期间对短暂的平静反应过度;反之亦然。确保资源管理既敏捷又可控，增加系统的稳定性。

Algorithm 1: Captain: scaling up and down

```
1  /* executes every  $N$  periods */
2  throttleCount = throttle count during last  $N$  periods;
3  throttleRatio = throttleCount/ $N$ ;
4  → margin = max(0, margin + throttleRatio - throttleTarget);
5  if throttleRatio >  $\alpha \times$  throttleTarget then
6    /* multiplicatively scale up */
7    quota = quota  $\times$  (1 + throttleRatio -  $\alpha \times$  throttleTarget);
8  else
9    /* instantaneously scale down */
10   history = CPU usage history in the last  $M$  periods;
11   proposed = max(history) + margin  $\times$  stdev(history);
12   if proposed  $\leq \beta_{\max} \times$  quota then
13     | quota = max( $\beta_{\min} \times$  quota, proposed);
14   end
15 end
```

Design

- Captain

3、Rollback mechanism after scaling down

恢复“鲁莽”的缩减。最后加上等于两个配额之差的额外CPU配额，分配多一点的CPU，以考虑由于错误的缩小而可能发生的潜在处理延迟。

Algorithm 2: Captain: rollback mechanism

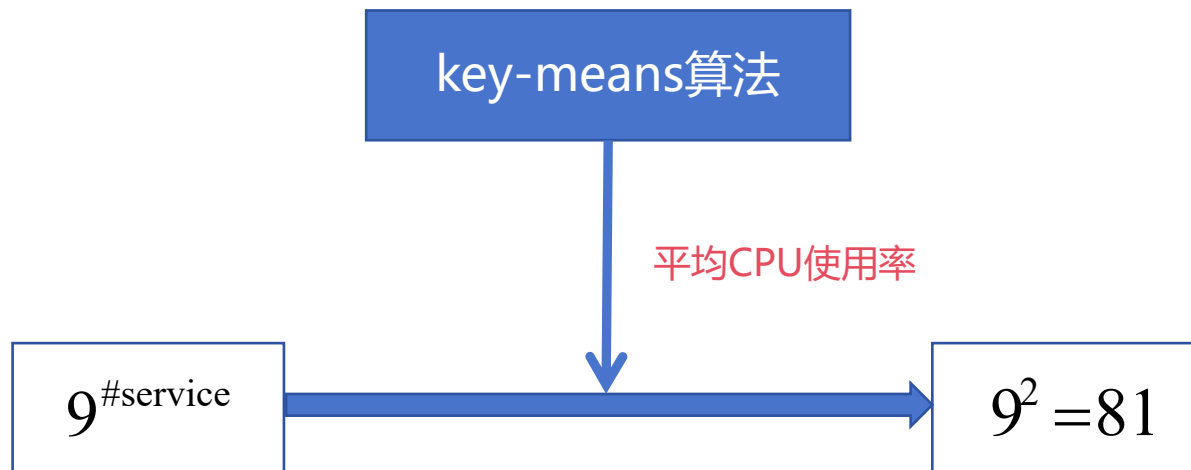
```
1  /* executes every period for  $N$  periods after each scale-down */
2  lastQuota = CPU quota before scale-down;
3  throttleCount = throttle count since scale-down;
4  throttleRatio = throttleCount/ $N$ ;
5  if throttleRatio  $> \alpha \times$  throttleTarget then
6    /* revert to the previous (higher) quota before scale-down
7     with an additional allocation equal to the quota difference */
8    quota = lastQuota + (lastQuota - quota);
9    margin = margin + throttleRatio - throttleTarget;
10 end
```

Design

- **Tower**

1、使用Contextual Bandits在线学习，根据RPS动态调整每个服务的CPU节流目标，在满足延迟SLO的同时尽量少用CPU资源。

2、减少行动空间



Dictionary

- Background
- Design
- **Evaluation**
- Related work
- Conclusion
- Think

Evaluation

1、基准应用程序（三个微服务应用测试对象）

- Train-Ticket
- Hotel-Reservation
- Social-Network

2、比较基线（对照组）

- Kubernetes (K8s-CPU and K8s-CPU-Fast)
- Sinan (ML方案)

3、四种流量模式：

- Diurnal
- Constant
- Noisy
- Bursty

4、长期测试：使用21天真实云厂商日志，验证策略的长期稳定性。

5、工具选择：使用Locust

Evaluation

Workload	Autothrottle	K8s-CPU	K8s-CPU-Fast	Sinan
Diurnal	30.4	58.0 (↓47.59%)	41.2 (↓26.21%)	278.4 (↓89.08%)
Constant	21.7	24.8 (↓12.50%)	27.3 (↓20.51%)	279.9 (↓92.25%)
Noisy	15.5	23.6 (↓34.32%)	17.7 (↓12.43%)	251.8 (↓93.84%)
Bursty	17.7	27.1 (↓34.69%)	21.9 (↓19.18%)	268.3 (↓93.40%)

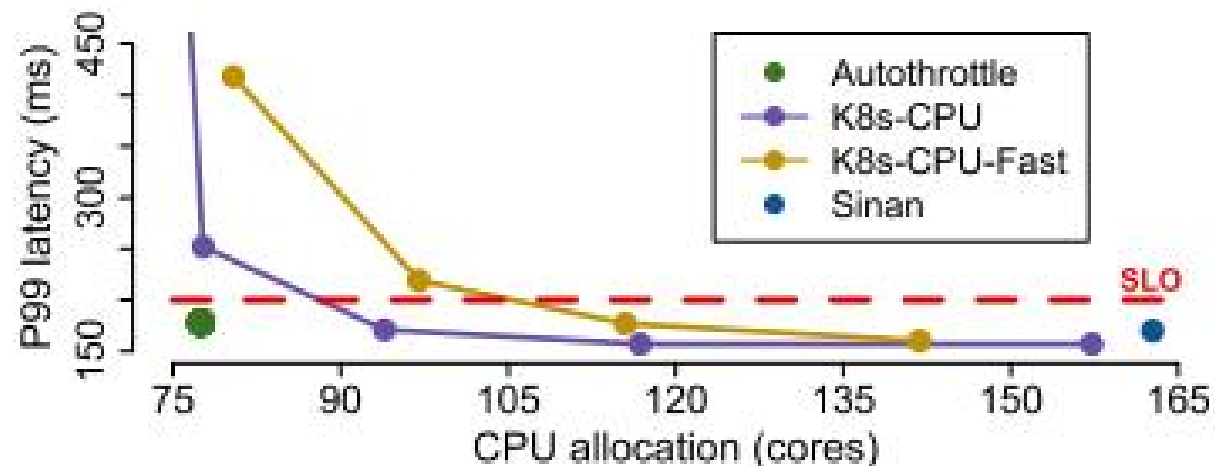
(a) Train-Ticket application (SLO: 1,000 ms P99 latency)

Workload	Autothrottle	K8s-CPU	K8s-CPU-Fast	Sinan
Diurnal	77.5	93.9 (↓17.47%)	115.5 (↓32.90%)	162.7 (↓52.37%)
Constant	88.7	115.6 (↓23.27%)	118.8 (↓25.34%)	149.7 (↓40.75%)
Noisy	57.5	66.5 (↓13.53%)	105.1 (↓45.29%)	105.2 (↓45.34%)
Bursty	50.0	67.5 (↓25.93%)	99.7 (↓49.85%)	111.9 (↓55.32%)

(b) Social-Network application (SLO: 200 ms P99 latency)

Workload	Autothrottle	K8s-CPU	K8s-CPU-Fast	Sinan
Diurnal	15.3	15.7 (↓2.55%)	16.5 (↓7.27%)	45.5 (↓66.37%)
Constant	11.2	11.5 (↓2.61%)	11.3 (↓0.88%)	21.2 (↓47.17%)
Noisy	10.8	12.1 (↓10.74%)	11.6 (↓6.90%)	65.9 (↓83.61%)
Bursty	10.1	15.7 (↓35.67%)	10.9 (↓7.34%)	63.1 (↓83.99%)

(c) Hotel-Reservation application (SLO: 100 ms P99 latency)



在160核集群上的实验结果，Autothrottle
在所有应用程序中都优于基线。

将相当数量的CPU分配给Autothrottle时，
K8s-CPU和K8s-CPU- fast将违反SLO

Evaluation

计算:

- P99延迟与CPU节流的Pearson相关系数
- P99延迟与CPU利用率的Pearson相关系数

现象:

CPU节流比CPU利用率表现出和P99延迟更高的相关性，存在更强的线性关系。

结论:

CPU节流与应用程序延迟的高相关性，所以我们使用CPU节流作为中间代理指标。

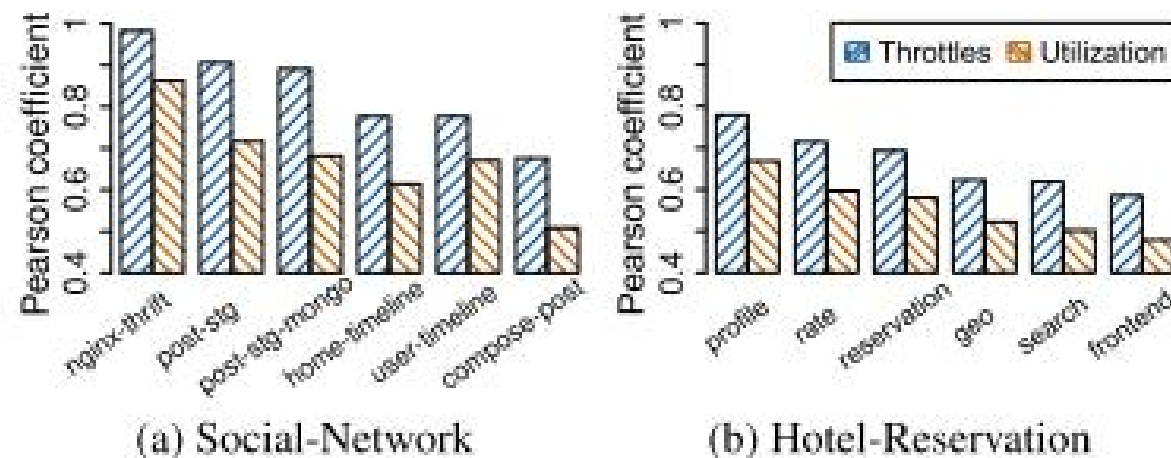



Figure 7: As a proxy metric, CPU throttles exhibit a higher correlation with application latencies than CPU utilization. The figure shows top microservices with highest CPU usage.

Evaluation

1、Long-term evaluation




(a) Autothrottle分配的cpu和K8s-CPU基线。红框表示K8s-CPU违反SLO规定的小时数。

- 将Autothrottle与性能最好的基准K8s-CPU进行比较

Evaluation

1、Long-term evaluation



(b) Social-Network's P99 latency, as achieved by Autothrottle and the K8s-CPU baseline. Dashed red line illustrates the 200 ms SLO.

- 该图显示了社交网络每小时的P99延迟
- 观察结果是：Autothrottle能够连续地保持接近200毫秒SLO的P99延迟
- 结论：应用程序使用Autothrottle性能更稳定

Evaluation

2、Large-scale evaluation

现象：Autothrottle需要分配的CPU内核更少

与性能最好的基准K8s-CPU和K8s-CPU-fast相比，Autothrottle最多可节省28.24%CPU内核和至少5.92%CPU内核。




图10: Autothrottle和基线分配的CPU核数，以满足Social-Network的P99 SLO。图中显示了Autothrottle在512核集群上的可伸缩性。

Dictionary

- Background
- Design
- Evaluation
- **Related work**
- Conclusion
- Think

Related work

1、垂直扩展

方法	机制	创新/局限
Kubernetes VPA	基于利用率阈值启发式调整	简单易用，依赖静态阈值
Autopilot	历史数据 + multi-armed bandit选择	动态适应，需长期数据积累
Sinan	ML预测SLO违规概率	精准但模型训练成本高
FIRM	RL定位根本原因并垂直扩展	针对SLO根本原因，策略收敛慢
Autothrottle	双层设计和CPU节流比例指标	实时响应，直接保障网络SLO

Related work

2、代理指标

传统指标	问题	Autothrottle方案
CPU利用率	高利用率 \neq SLO违规	CPU节流指标：直接反映资源竞争
队列长度	忽略单个请求复杂度，队列分散使测量难	闭环控制动态调整配额，保障吞吐量
排队延迟	取决于线程模型，需手动测试	无侵入监控，自适应阈值
总结	传统指标滞后/误导	创新：节流信号 + 实时反馈

Related work

3、横向与混合扩展

方法	机制	适用场景
Kubernetes HPA	基于利用率/QPS调整副本数	通用场景，响应延迟高
GRAF	图神经网络建模服务依赖	复杂依赖系统，计算开销大
COLA	多服务协同调整	避免局部优化，需全局协调
混合扩展	垂直优先（短期）+水平（长期）	平衡速度与弹性，策略设计复杂

Dictionary

- Background
- Design
- Evaluation
- Related work
- **Conclusion**
- Think

Conclusion

Autothrottle是一个两级学习辅助资源管理框架，用于慢速目标微服务。

微服务架构中，端到端延迟由多个服务共同决定，但直接通过延迟调整每个服务的资源分配面临以下挑战：

- 1、延迟反馈滞后（例如请求需经过多个服务处理）。
- 2、不同服务对延迟的贡献差异大（某些服务可能是瓶颈）。

Autothrottle通过CPU节流比作为中间代理指标，将全局的延迟SLO转化为每个服务的本地资源控制目标。

Dictionary

- Background
- Design
- Evaluation
- Related work
- Conclusion
- **Think**

Think

- 目前仅针对CPU资源进行优化，而实际系统中内存、网络带宽或磁盘I/O可能同样关键，未来可扩展框架以支持多种资源（例如内存）联合优化。
- 论文假设请求类型分布恒定，但在实际场景中，请求类型的动态变化可能导致性能目标失效。可探索基于请求类型的细粒度目标调整，或引入请求分类器作为上下文输入。

Thank you!