



D²MoE: Dual Routing and Dynamic Scheduling for Efficient On-Device MoE-based LLM Serving

MobiCom' 25

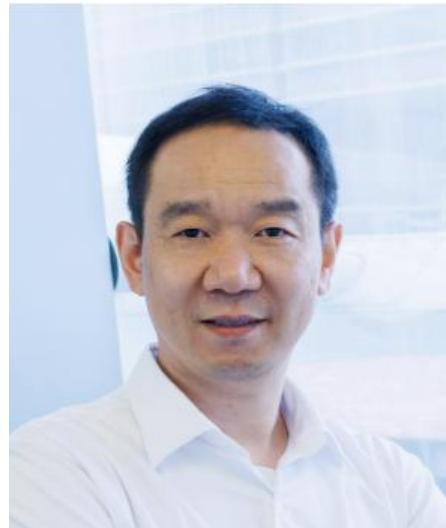
Haodong Wang, Qihua Zhou*, Zicong Hong*, Song Guo



深圳大学计算机与软件学院
College of Computer Science & Software Engineering,
Shenzhen University

Presenter: Yunpeng Xu
2025.9.18

Author



Song Guo



[Pervasive intelligence Lab](#)

研究方向：普适智能（Pervasive Intelligence）

研究核心聚焦于**大模型（LLMs）**、**边缘智能（Edge AI）**、**云边端协同**和**AI安全**的交叉优化，探索如何通过系统与算法的协同设计，实现高效、可靠、可扩展的智能服务。

- **Agentic AI Theory**

- 高效推理 (Efficient Reasoning)
- 测试时对齐 (Test-time Alignment)
- 多智能体协作 (Multi-Agent Collaboration)

- **AI & Web3.0 Security**

- AI 系统安全漏洞 (AI Vulnerabilities)
- AI 治理 (AI Governance)
- 区块链与分布式安全 (Blockchain Security)

- **Edge-Native AI System**

- AI 个人计算机 (Artificial Intelligence PC)
- 云-边-端协作 (Cloud-Edge-End Collaboration)
- 国产 GPU 与 AI 芯片生态 (Chinese GPU)

- **Pervasive AI Applications**

- 具身智能 (Embodied AI)
- 科研赋能 (AI for Sciencee)
- AI 内容生成 (AI Generated Content)

Contents

- **Background & Motivation**
- **Challenges**
- **Design**
- **Evaluation**
- **Thinking**

Background & Motivation

MoE-based LLM inference

The Real Bottleneck?
Computation **vs** Memory



- **Hardware.** The growth of memory capacity in edge devices **significantly lags behind** that of high-performance data centers in the cloud.
- **Model.** The **memory demands** of MoE-based LLMs have **grown rapidly** due to the expansion of model size and the need for higher performance.

Table 1: Trade-offs among different bit-width.

Bit-width	mem. (GB)	lat. (token/s)	acc. (ppl ↓)
2	3.04	50.47	20.95
3	3.80	45.91	15.10
4	4.48	43.82	14.72
5	5.10	40.15	14.63
6	5.60	37.72	14.62
8	7.24	35.34	14.55
16	13.60	23.45	14.55

How to bridge the gap?

Background & Motivation

Observation:

1. **Different bit-width** in MoE-based LLM quantization **bring different benefits** in terms of **accuracy-memory-latency**.
2. **The importance of experts changes dynamically** according to **different input samples**.
3. **Large bubble** between **I/O and computation** of quantized experts led to substantial **inference delays**.

Dynamically adjust the expert bit-width to align with hardware characteristics

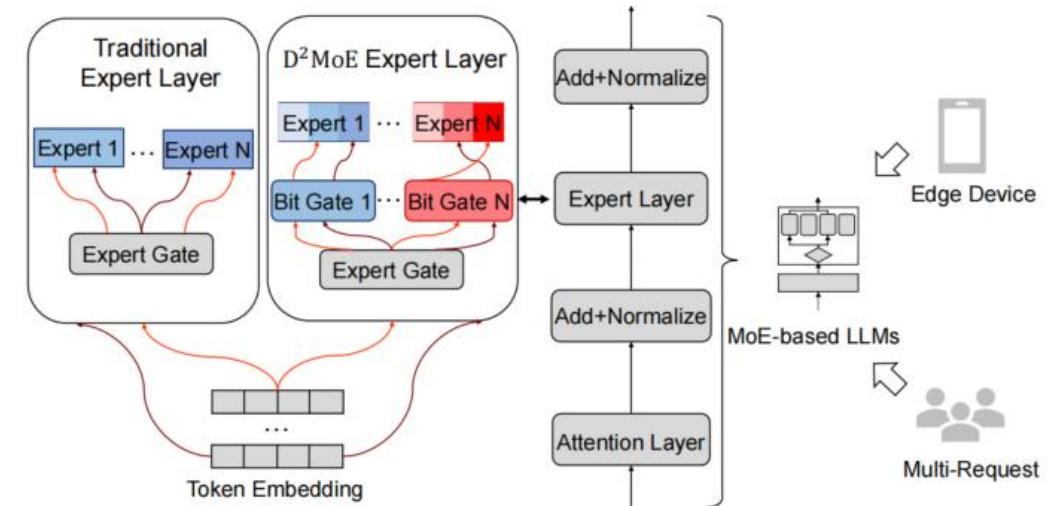


Figure 1: Traditional MoE single routing (expert ID only) vs. our D²MoE dual routing (ID and bit-width).

Challenges

◆ Challenge 1: **Unbalanced and inefficient** bit-width selection load

The **training** of lightweight gating network



- Imbalanced Selection
- Irrational Allocation

◆ Challenge 2: **High memory** overhead to store multiple versions of quantized experts

The basic LLM quantization



Edge devices with
limited memory

multiple quantization
models + **different** bit-width
versions

How To Handle?

Challenges

- ◆ Challenge 3: Significant runtime **overhead** on weight **dequantization** operations

During model inference:

Transfer the quantized weights into **the same data type** as the activation for matrix computation



Consume **20-70% of the entire inference latency**

- ◆ Challenge 4: **Inefficient I/O–Compute pipeline** with large parallelism **bubbles**

Existing methods to achieve I/O-compute parallelism:

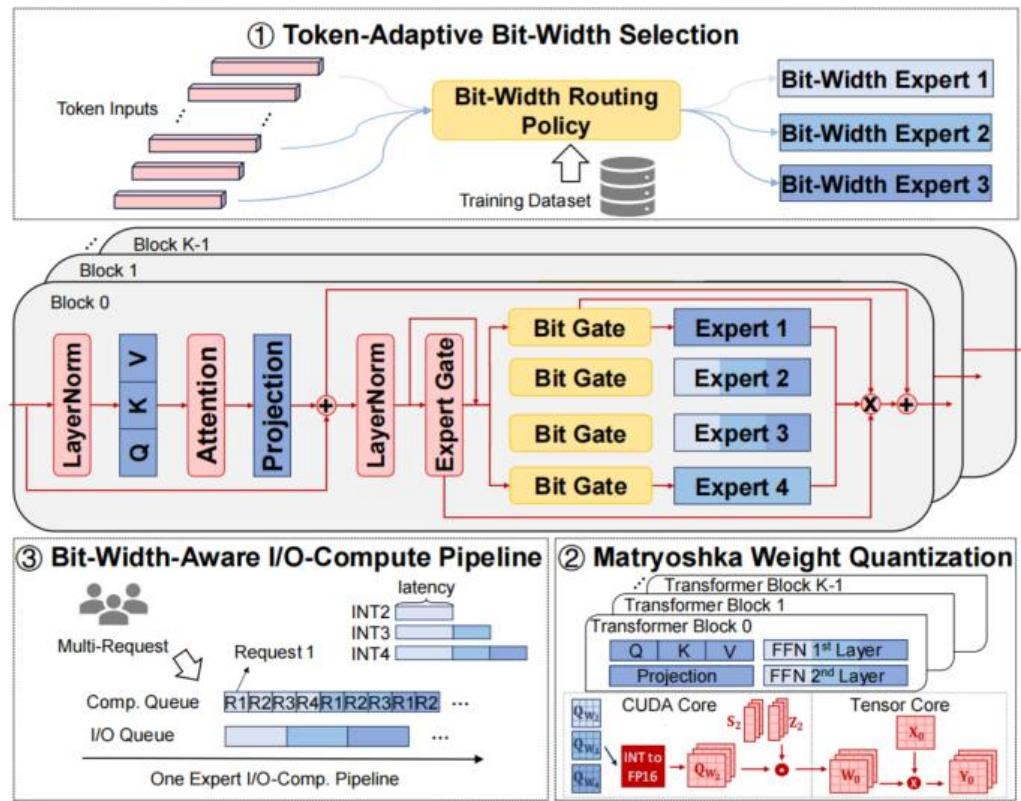
- model placement policies
- adaptive batch scheduling



Quantized LLMs serving must also consider **the varying bit-width selections** for different requests

Design

System architecture of D²MoE



● The offline preprocessing phase (On server)

- ① Token-adaptive bit-width selection
 - > **train a lightweight** gate to dynamically select the optimal bit-width expert for each token

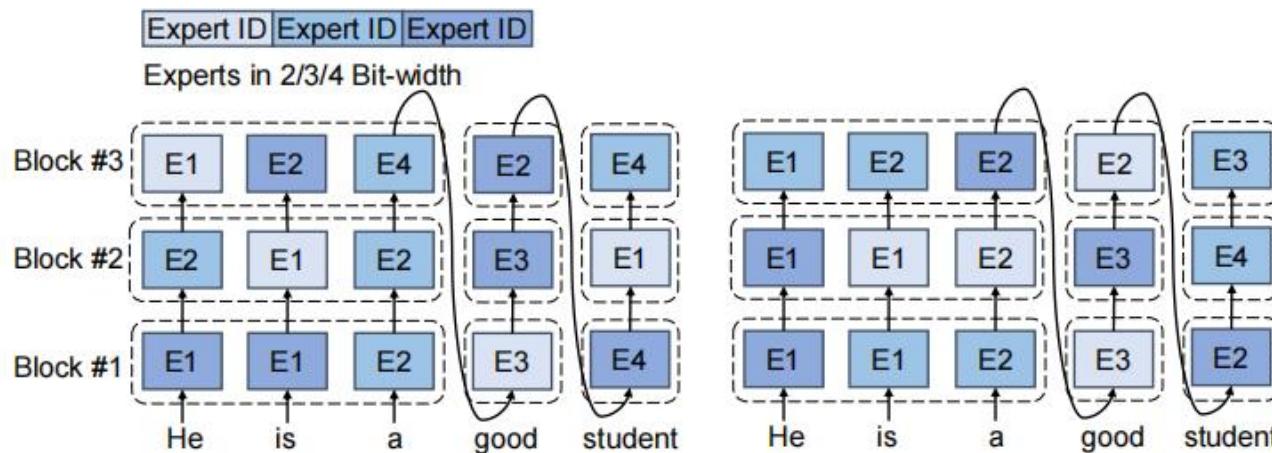
- ② Matryoshka weight quantization (MWQ)
 - > **compress** expert weights into a **nested, multi-bitwidth** format

● The online execution phase (On device)

- ③ The bit-width-aware I/O-compute pipeline
 - > employ the **HEBF** scheduler to overlap I/O and computation to **eliminate the idle bubbles**

Design

➤ Core modules 1: Token-Adaptive Bit-Width Selection



traditional method vs token-adaptve method

Two steps to deal with challenge 1 :

① Quantized Expert Capacity

Set a capacity limit c_k for each bit-width expert to prevent overfitting and force the router to **learn balance**.

② Dynamic Bit-Width Selection Loss

Propose a novel bit-width balancing loss to **balance the selection frequency** of different bit-width.

$$Loss = \frac{1}{T} \sum_{x \in S} \left(\underbrace{CE(p(x), q(x))}_{\text{任务精度损失}} + \frac{\alpha}{L} \sum_{l=1}^L \sum_{k=1}^K \underbrace{p_k^l(x)}_{\text{选择概率}} \underbrace{b_k}_{\text{比特成本}} \right).$$

Design

➤ Core modules 2: Matryoshka Weight Quantization

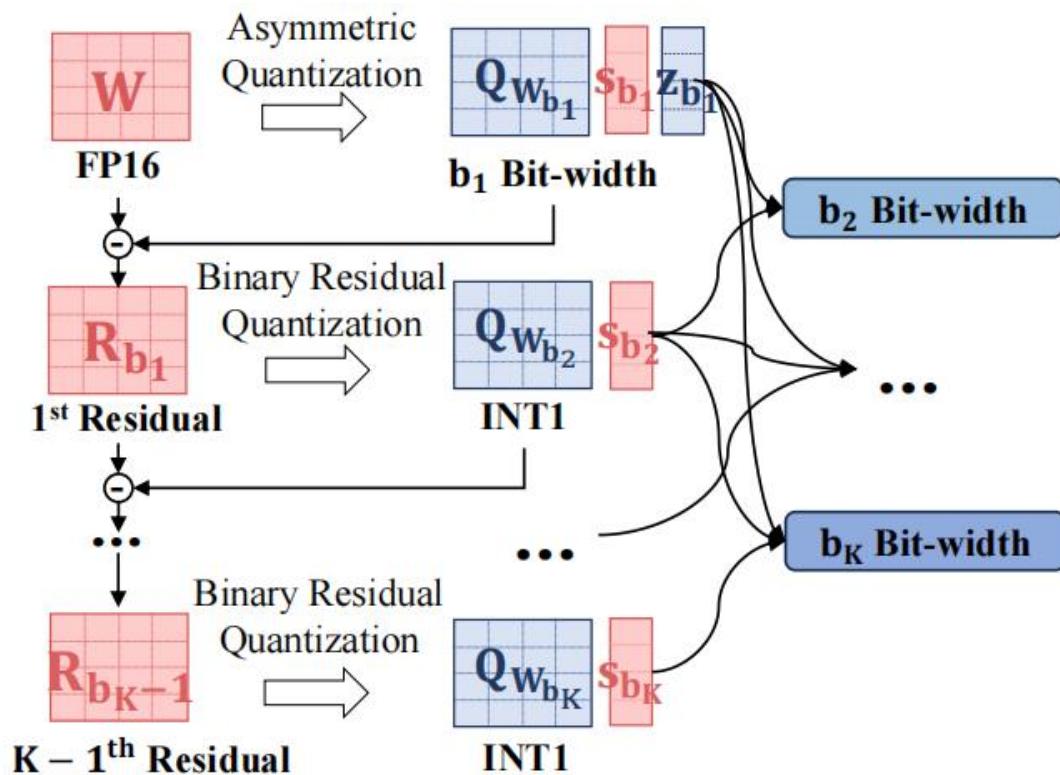


Figure 6: The workflow of MWQ.

Embed low bit-width weights within high bit-width weights to deal with challenge 2 :

- ① Asymmetric Quantization (to b_1)
Quantize the weights to the minimum supported bit-width (denoted as b_1 before).
- ② Binary Residual Quantization (b_2 to b_k)
Quantize the residual weights to increase the bit-width until the final b_k bit-width weight is obtained.

Design

Dequantization Kernel to deal with challenge 3:

tensor loading from various storage levels in the GPU

+

computation within the CUDA cores and Tensor cores

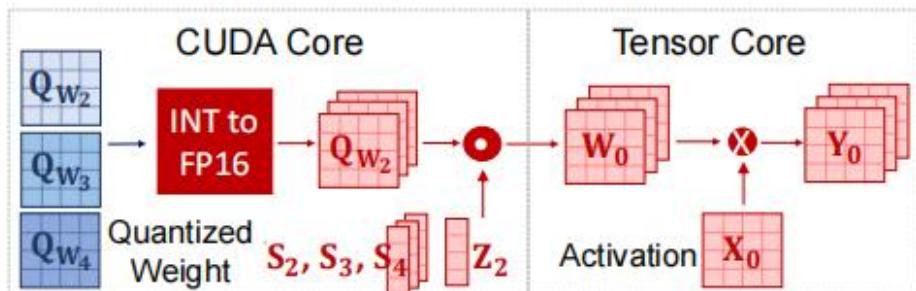


Figure 8: The dequantization overview of D²MoE.

The Limited Parallelism

constrain the dequantization efficiency **on edge device**

1. Loading Parallelism

- ① transfer quantized data from disk directly to GPU's global memory
- ② move activations from global memory to L2 cache

2. Computation Parallelism

- ① expert dequantization in the CUDA cores
- ② expert computation in the Tensor cores

Design

➤ Core modules 3: Bit-Width-Aware I/O-Compute Pipeline

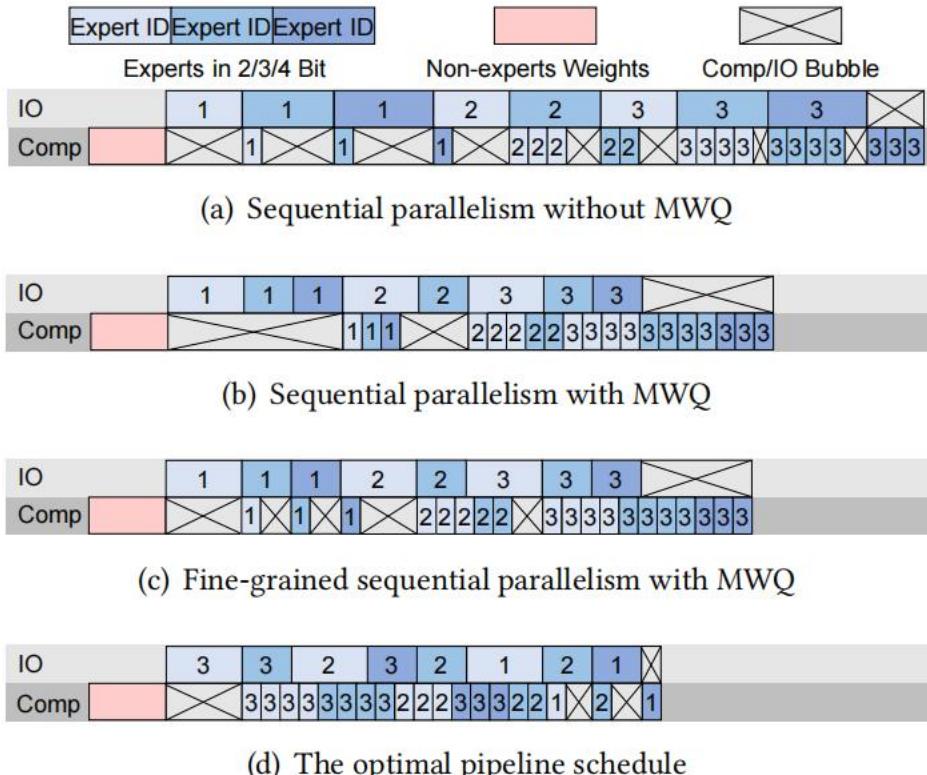


Figure 9: Comparison between different I/O-compute parallel strategies.

How to reduce the parallel bubbles?

- ① Low-bit-width experts can be **reused (MWQ)**
- ② Break the **ID-order** and dynamic scheduled by **bit-width + activation frequency (HEBF)**



Maximize I/O-Compute overlap



Minimize idle time and end-to-end latency

Evaluation

1. Experimental Setup

● Hardware

Hardware	Environment 1		Environment 2	
	Device	Memory	Device	Memory
GPU	NVIDIA RTX 3060	6GB	Jetson AGX Orin	64GB
CPU	Intel Core i7-11800H	32GB	ARM Cortex-A78AE (SoC share)	
Disk	Samsung 970 EVO	1T	Samsung 970 EVO	1T
Disk Read		3.5 GB/s		3.5 GB/s

● Baseline

- ① Hold-in-Memory (**all** INT8 and hold in GPU memory)
- ② Matryoshka-Free (INT2/3/4 and load on demand)
- ③ Hold-in-Memory-AWQ (**all** INT4 and hold in GPU memory)
- ④ EdgeMoE (**static** quantization and predict to load)
- ⑤ MoQE-DynalO (**uniform** bit-width and load dynamically)

● Metrics

① model accuracy

② throughput

language generation performance
zero-shot performance

● Configuration (compared to baseline)

① D2MoE-V1

- $b_1 = 2$ and $b_k = 4$
- expert capacity is set to $\{0.3, 0.4, 0.3\}$

② D2MoE-V2

- $b_1 = 5$ and $b_k = 8$
- the capacity of each expert is 0.25

Evaluation

2. End-to-End Results

Table 3: **Perplexity accuracy (lower is better)** and **zero-shot accuracy (higher is better)** of D²MoE and the baselines in LLaMA-MoE-3.5B and Mixtral 8×7B.

Model	Method	Perplexity ↓	PiQA	Arc.e	BoolQ	HellaSwag	Winogrande
LLaMA-MoE-3.5B	Hold-in-Memory	14.55	72.32	48.76	65.56	66.34	61.77
	Matryoshke-Free	14.58	72.29	48.71	65.48	66.28	61.76
	Hold-in-Memory-AWQ	15.89	69.32	46.52	62.54	61.55	57.44
	EdgeMoE	14.78	71.46	47.36	64.43	64.32	59.52
	MoQE-DynaIO-INT4	15.66	69.34	46.52	62.58	61.53	57.40
	MoQE-DynaIO-INT8	14.55	72.32	48.74	65.58	66.28	61.71
	D ² MoE-V1	15.68	69.32	46.52	62.50	64.28	59.52
	D ² MoE-V2	14.58	72.29	48.72	65.51	66.23	61.71
Mixtral 8×7B	Hold-in-Memory	4.04	82.4	82.6	80.56	84.1	76.5
	Matryoshke-Free	4.28	80.29	80.71	78.48	82.28	74.76
	Hold-in-Memory-AWQ	4.25	80.32	81.52	78.54	83.55	75.44
	EdgeMoE	4.38	78.46	80.36	77.43	82.32	75.52
	MoQE-DynaIO-INT4	4.25	80.34	81.52	78.58	83.53	75.40
	MoQE-DynaIO-INT8	4.08	82.32	81.74	79.58	83.28	74.71
	D ² MoE-V1	4.28	81.32	81.52	78.05	82.88	75.52
	D ² MoE-V2	4.09	82.29	81.72	79.51	83.23	74.71

Model Accuracy

① language generation performance

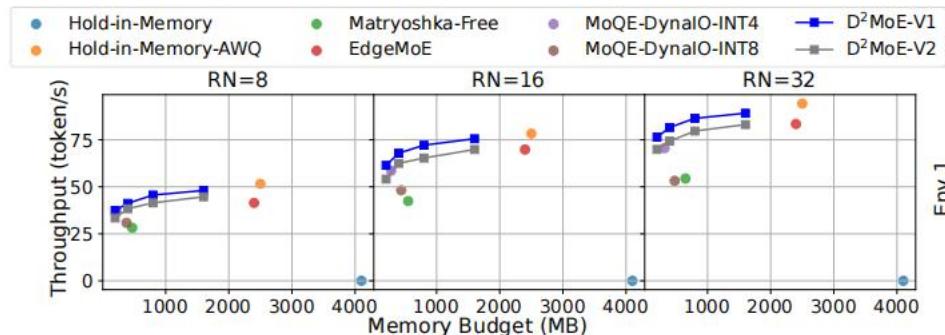
Calculate the **perplexity** to evaluate the model's **intrinsic language quality**.

② zero-shot performance

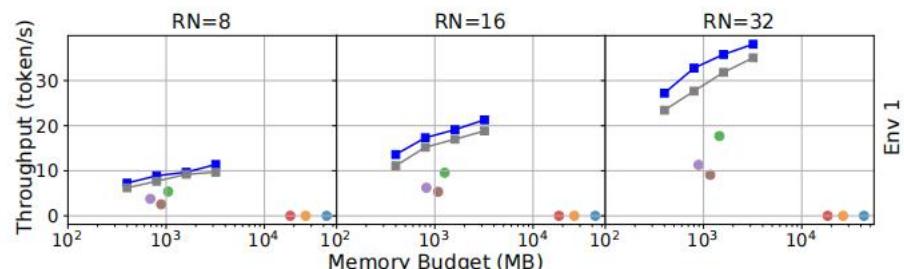
Test the model's **general capabilities** on tasks like commonsense reasoning and reading comprehension.

Evaluation

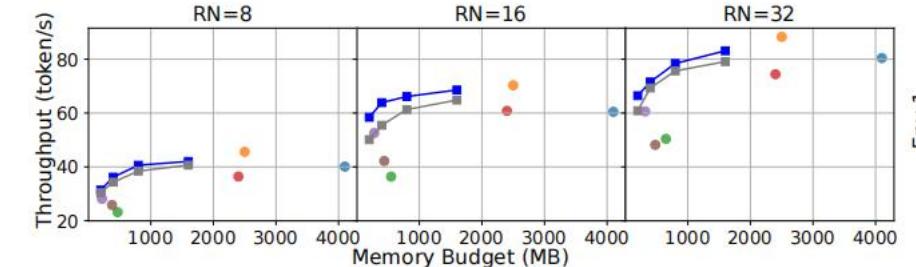
2. End-to-End Results



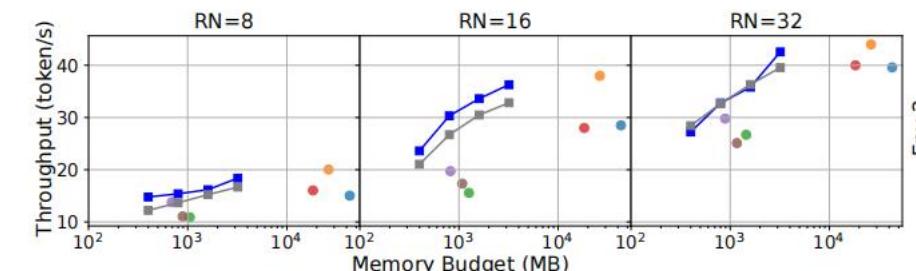
(a) Throughput of LLaMA-MoE in Environment 1.



(b) Throughput of Mixtral 8x7B in Environment 1.



(c) Throughput of LLaMA-MoE in Environment 2.

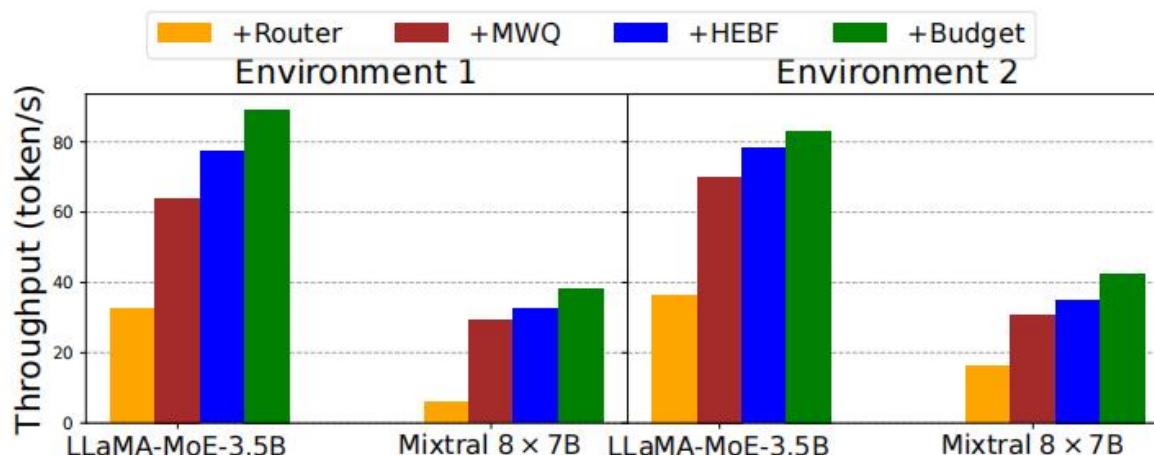


(d) Throughput of Mixtral 8x7B in Environment 2.

Throughput of D2MoE and baselines with different memory budgets

Evaluation

3. Ablation Study



- + Router: Enables **dynamic bit-width allocation** for each expert (**baseline**).
- + **MWQ**: **Embeds** lower bit-width weights into higher ones, reducing loading overhead.
 - Throughput improves **$1.91 \times - 4.95 \times$**
- + HEBF: **Parallelizes** expert I/O and computation
 - Additional **$1.11 \times - 1.21 \times$** improvement
- + Budget: Keeps **frequently** used experts in **GPU memory**
 - Further **$1.06 \times - 1.21 \times$** improvement

Thinking

- 文章有什么问题，基于这篇 Paper 还能做什么优化？
 1. D²MoE 主要针对 “同步请求” 设计，未来可以探索 “异步并发请求”的调度策略。
 2. 当前系统依赖 “按需加载”，未来可以引入 “预加载” 机制（热度预测），进一步减少 I/O 等待。
 3. 离线路由偏静态，专家容量设置固定，存在失配的可能性，未来可以引入在线自适应容量。
 4. 在移动设备上，NPU/CPU的反量化效率需要进一步优化，需面向NPU/CPU重新设定内核并行加载和并行计算机制。

Thinking

- **这篇 Paper 的 Idea 能不能应用在自己的工作上面?**
 1. D²MoE提出的“**动态精度**”和“**智能调度**”思想可以迁移到多模型推理、异构计算资源调度等领域。
 2. 另外，MWQ “**嵌套量化**”的存储思想也可以迁移至资源受限的边缘设备上以减少存储开销。
- **这篇 Paper 能不能泛化?**
 1. **模型架构泛化**，其他MoE变体等**稀疏模型**，甚至是**稠密模型中的FFN层**，这种**动态位宽、嵌套量化、智能调度**的思想都能有一定的性能提升。
 2. **核心概念泛化**，任何“**动态决策**”系统，任何“**分层存储**”或“**增量更新**”系统，任何存在**生产-消费**关系的流水线系统，上述核心思想都有性能提升的空间。



Q & A

**Presenter: Yunpeng Xu
2025.9.18**