



# White-Boxing RDMA with Packet-Granular Software Control

NSDI'25

Frank Chenxingyu Zhao, Jaehong Min, Ming Liu, Arvind Krishnamurthy  
University of Washington; University of Wisconsin-Madison

Presenter: Huang Kai  
2025.9.20

# Author

- **研究方向**

- 可编程网络、高性能网络、数据中心分布式系统、机器学习系统

- **近期论文**

- **可编程网络**: Gimbal: Enabling Multi-tenant Storage Disaggregation on SmartNIC JBOFs (SIGCOMM'21)
- **高性能网络**: CC-NIC: a Cache-Coherent Interface to the NIC(ASPLOS'24) Host Congestion Control (SIGCOMM'23)
- **数据中心分布式系统**: Xenic: SmartNIC-Accelerated Distributed Transactions (SOSP'21)
- **机器学习系统**: Efficient Direct-Connect Topologies for Collective Communications (NSDI'25) Punica: Multi-tenant LoRA serving (MLSys'24)



Arvind Krishnamurthy  
Short-Dooley Professor,  
University of Washington

# Content

- **Background**
- **Motivation**
- **Design**
- **Evaluation**
- **Thinking**

# Background

□ 数据中心网络对高吞吐低延迟有要求

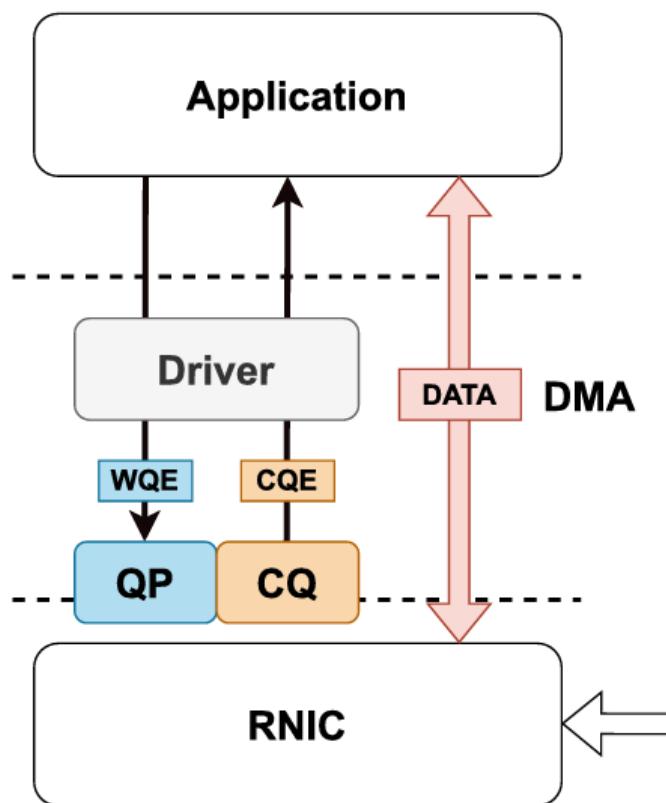
➤ RDMA NIC

特性	普通 NIC (传统网卡)	RNIC (RDMA 网卡)
核心技术	标准以太网/IP协议处理	支持 RDMA (远程直接内存访问)
CPU 开销	较高, CPU 深度参与数据传输与协议处理	极低, RDMA 绕过 CPU 进行内存直接访问
延迟	相对较高	非常低
吞吐量	受 CPU 限制	更高, 能更充分利用网络带宽
主要优势	成本较低, 通用性强	极低的延迟, 极高的吞吐量, CPU 卸载
适用场景	通用网络连接	高性能计算、数据中心、低延迟应用等

# Background

数据中心网络对低CPU开销、高吞吐低延迟有要求

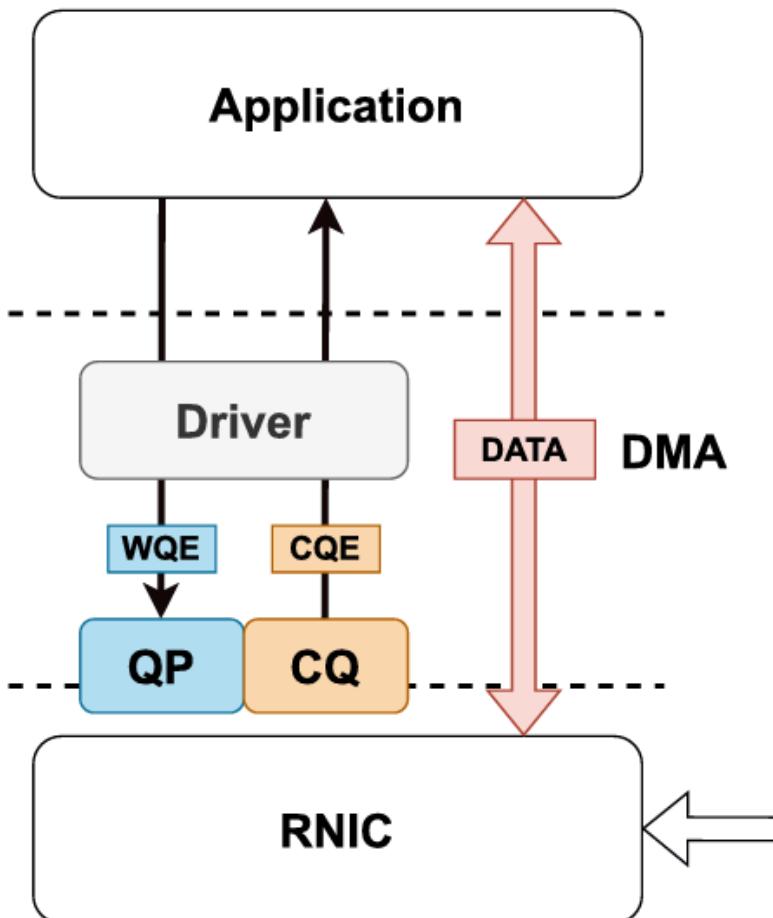
➤ RDMA 传输机制



- Work Queue Entry, WQE QP操作单元
- Queue Pair, QP 包括发送队列和接收队列
- Complete Queue, CQ 通知App任务完成
- Complete Queue Entry, CQE WQE完成后生成相应的CQE， 并传递给CQ

# Motivation

- RNIC(RDMA NIC) 算法定制 与 细粒度控制难度大



**FPGA**

High Dev. Cost

**ASIC**

Wait for Years

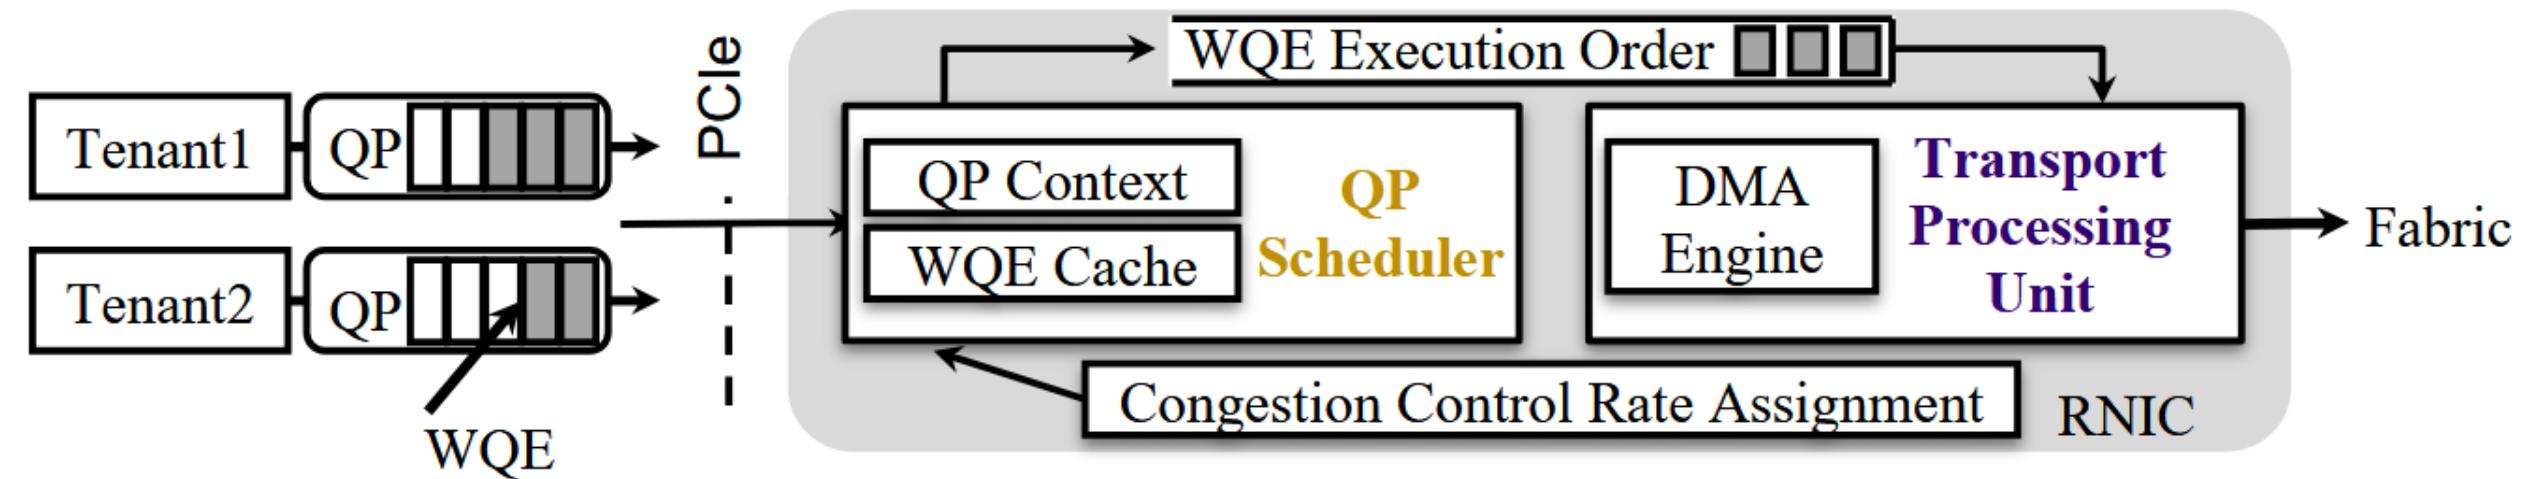
**SW Emulate**

Perf. Loss

**Blackboxing**

Coarse & Delay

# Motivation



**FPGA**

High Dev. Cost

**ASIC**

Wait for Years

**SW Emulate**

Perf. Loss

**Whiteboxing**

SW Ctrl  
HW Data

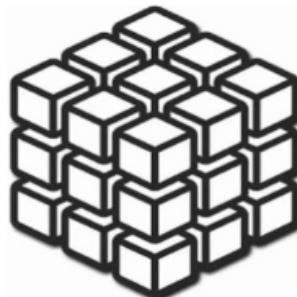
# Challenges

## □ SCR(Software Control RNIC) 存在的挑战



### Flexibility

Express Diverse  
Customizations



### Granularity

Packet-Granular Events  
from High Line Rates

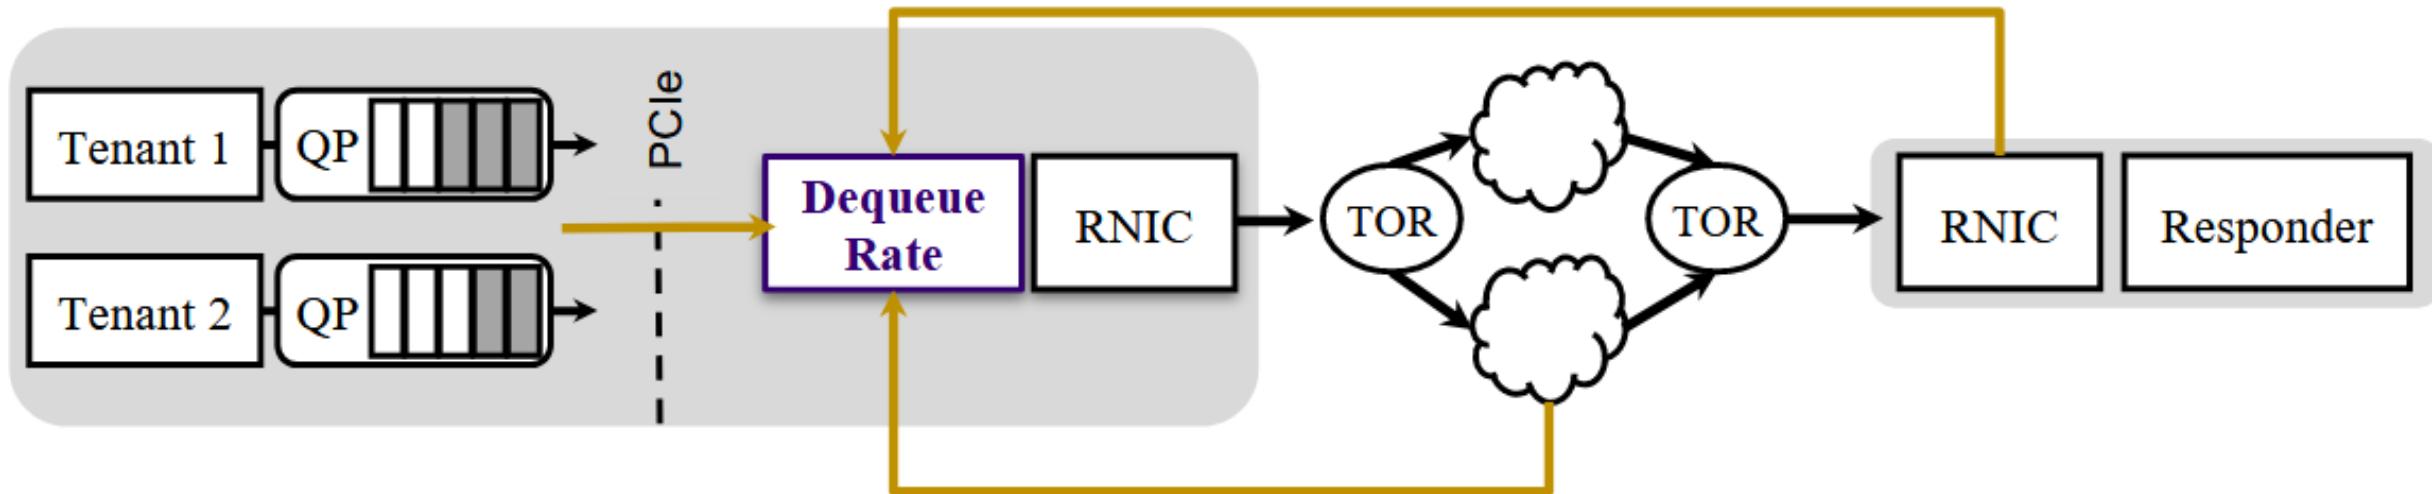


### Scalability

Handle Many  
Concurrent Flows

What to Control & How to Control

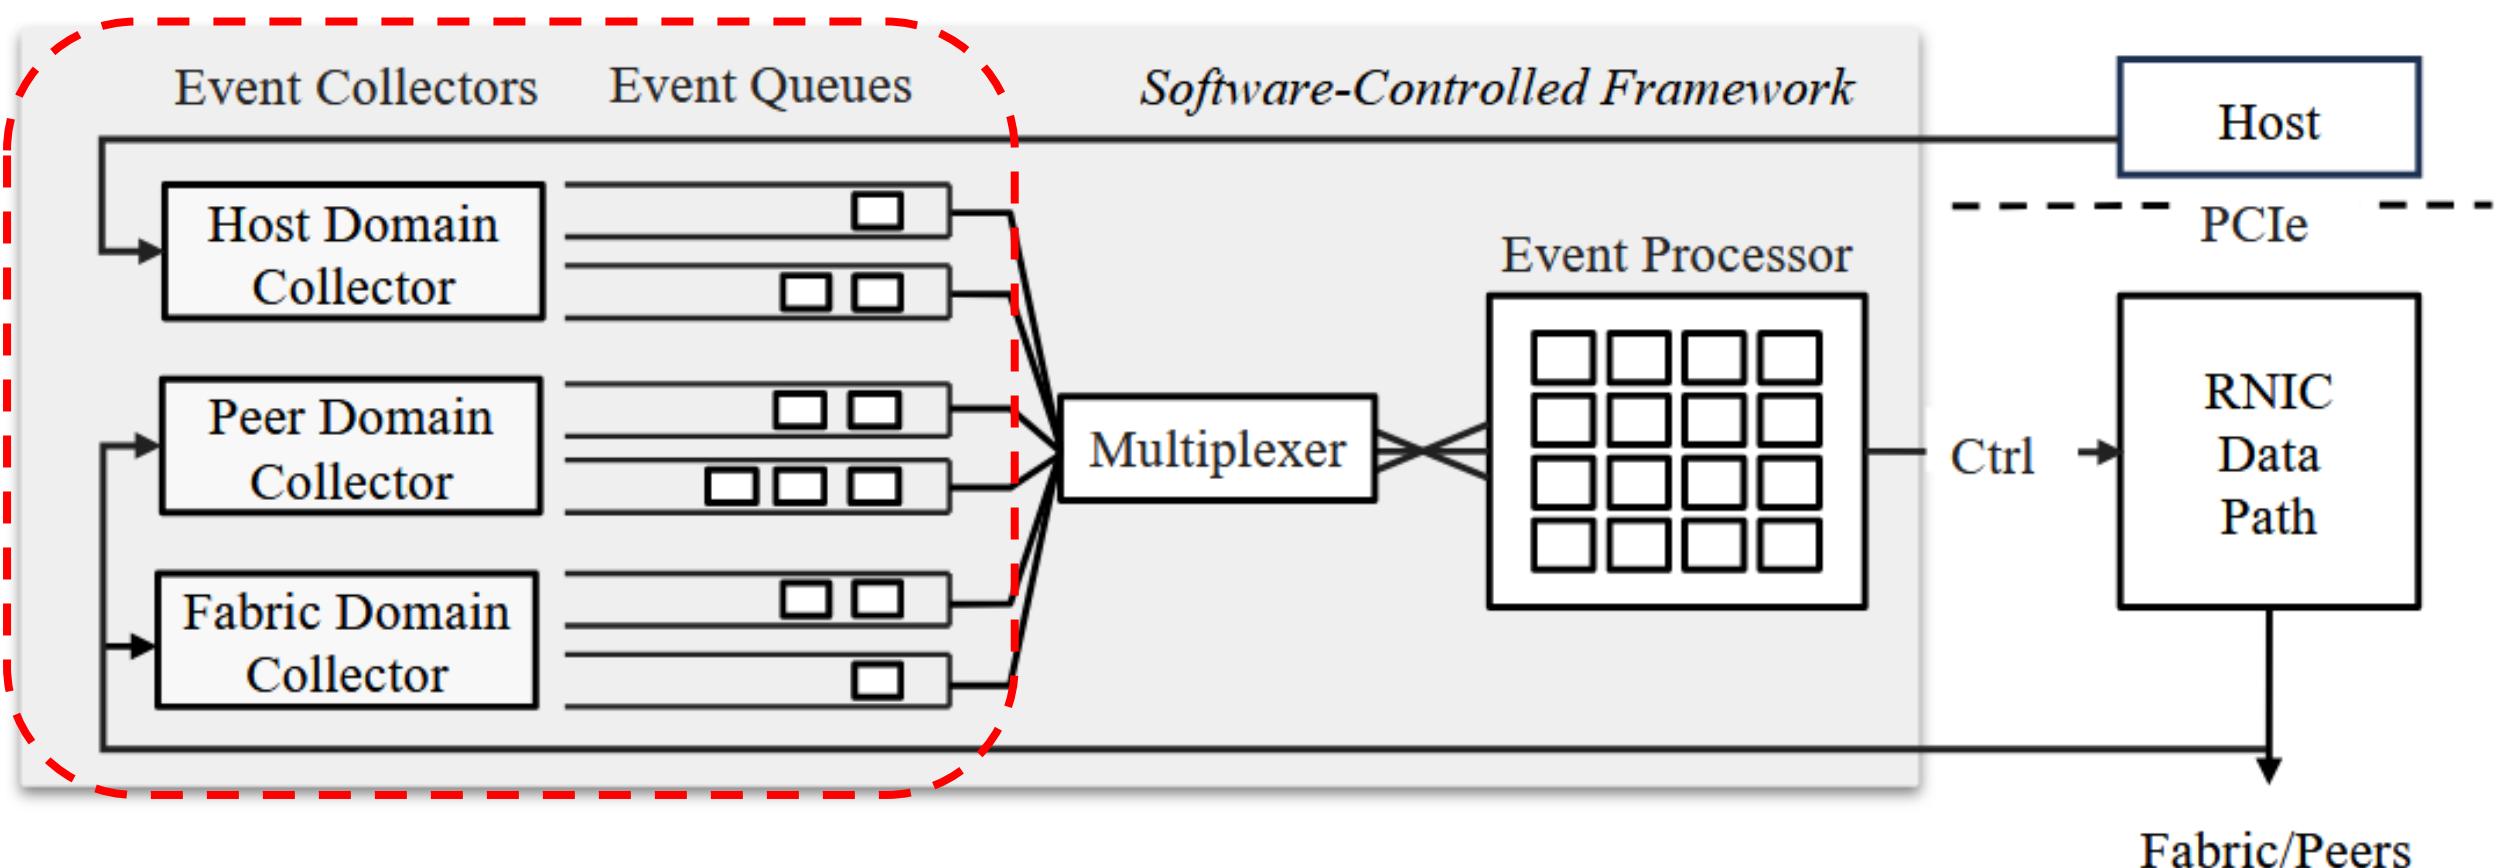
# What to Control -- Flexibility



□ 针对不同控制信号实现复杂场景下的细粒度控制

- 主机控制域：QP(流)出队速率
- 链路控制域：RNIC链路吞吐率、网络拥塞信号
- 端到端控制域：接收端反馈信号

# Design -- Overview

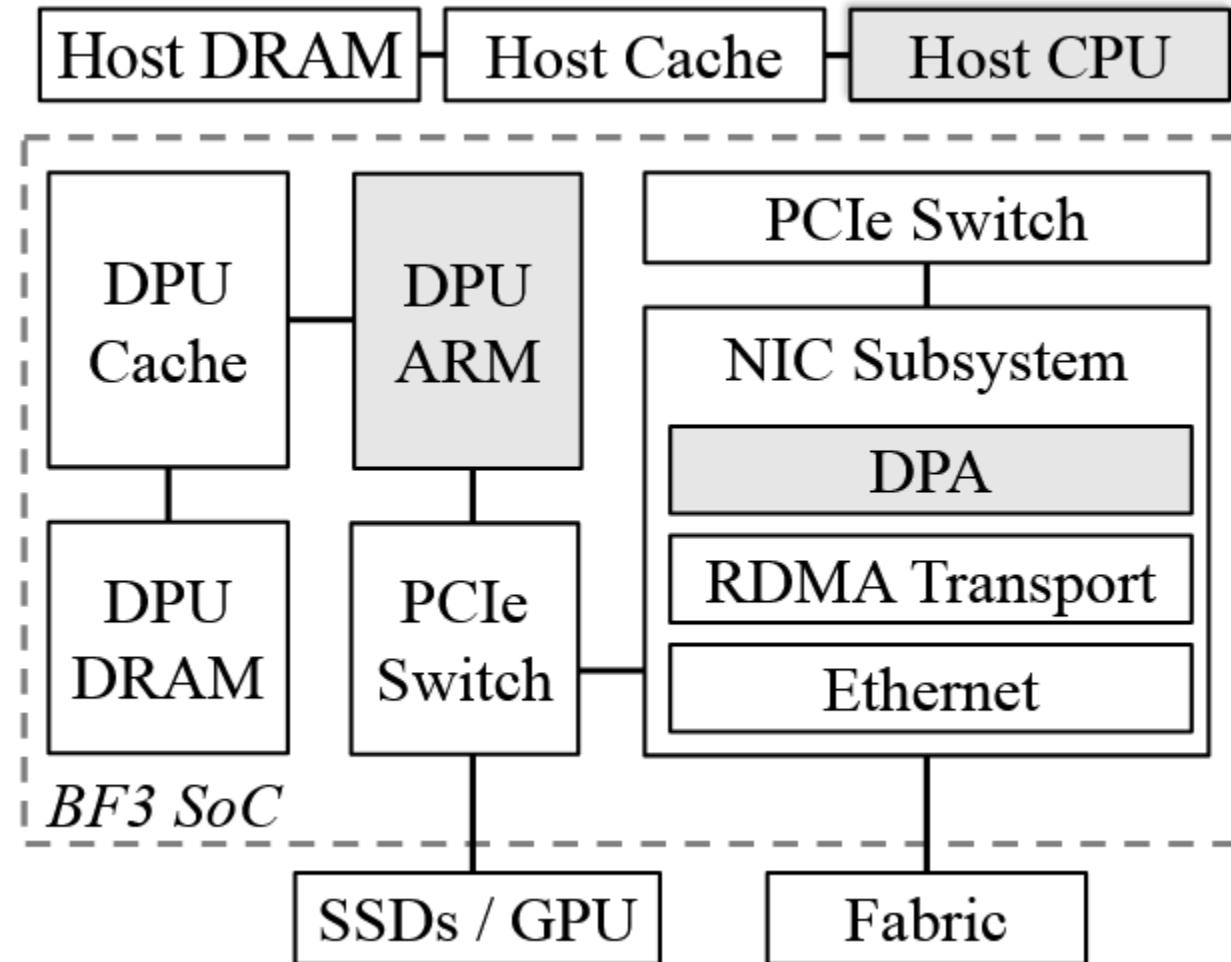


QP事件驱动的速率计算模型

# How to Control

## □ Nvidia BlueField-3架构处理器

- Host CPU
- DPU-embedded ARM CPU
- **DPA(Datapath Accelerators)**
  - 卸载控制平面 (CPU) 的数据转发任务，实现硬件级加速

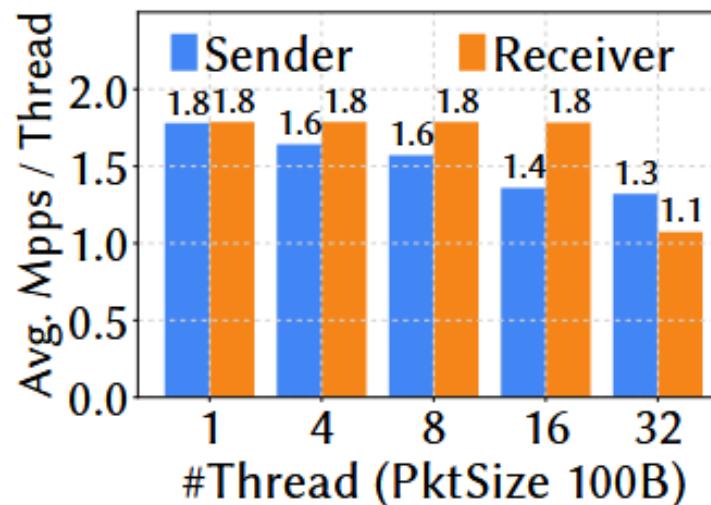
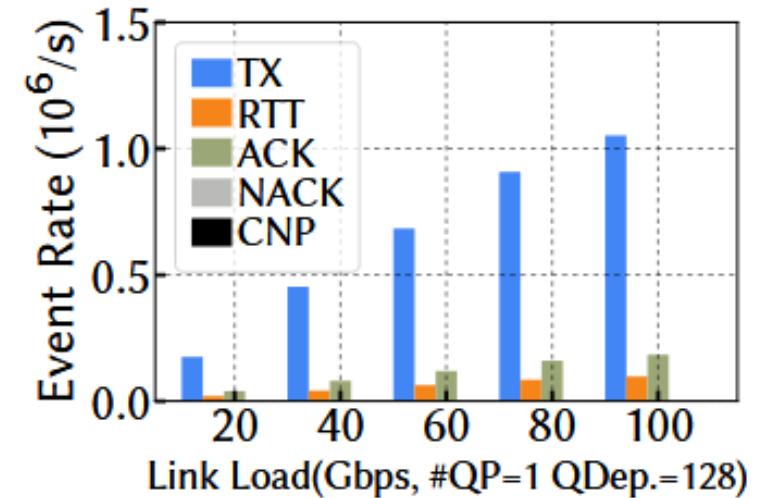
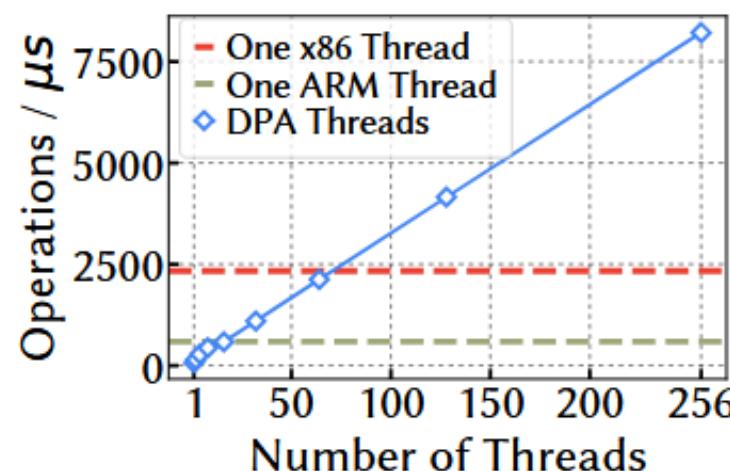


BlueField-3架构

# How to Control -- DPA

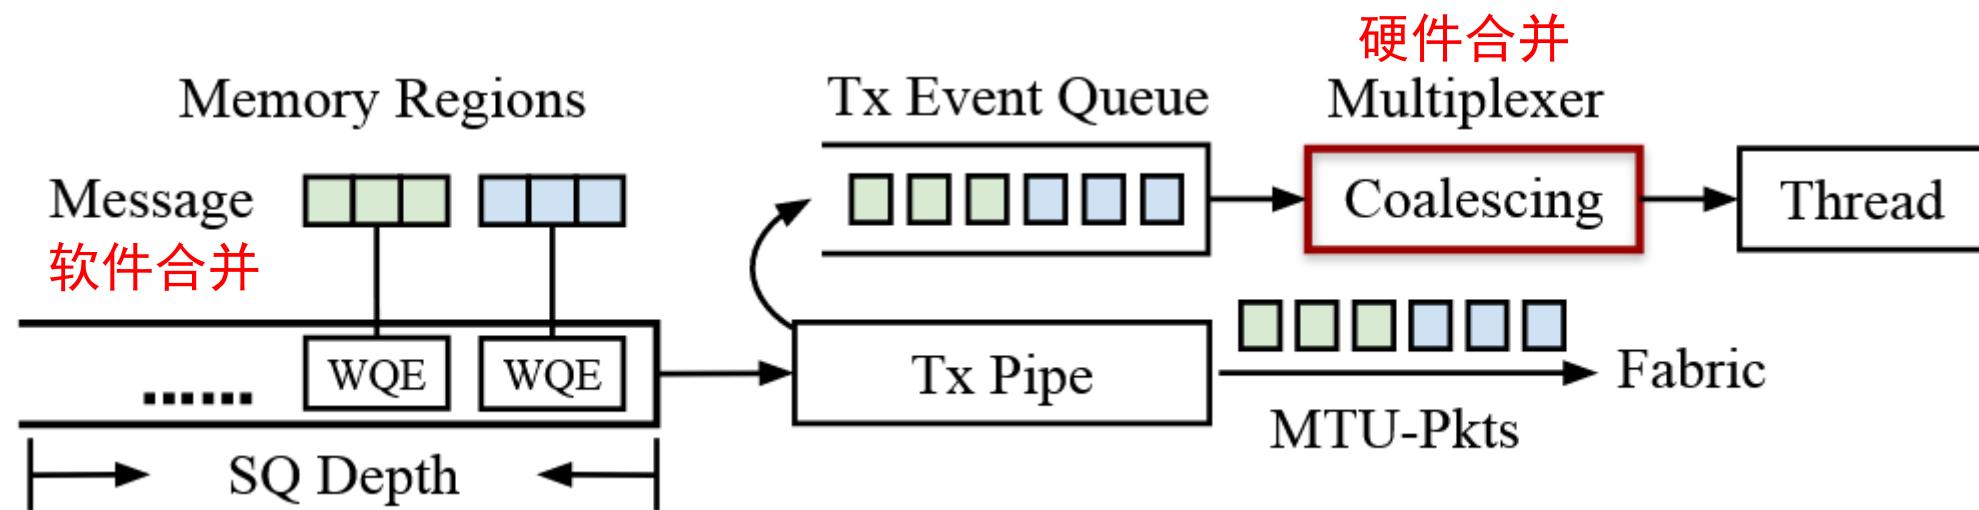
## DPA性能测试

- 多线程性能
- 端到端信号的采集
- 吞吐效率



# How to Control -- Granularity

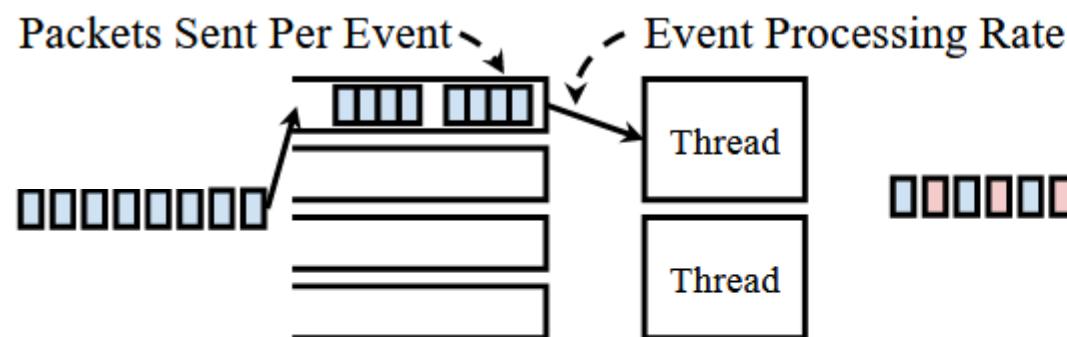
- 硬件合并 (ASIC -> RNIC TX/RX pipelines)
  - 只能合并链路信号
- 软件合并 (DPA -> QP SQ/RQ)
  - DPA进程间通信能力有限——CPU/DPU ARM 代理进程间通信



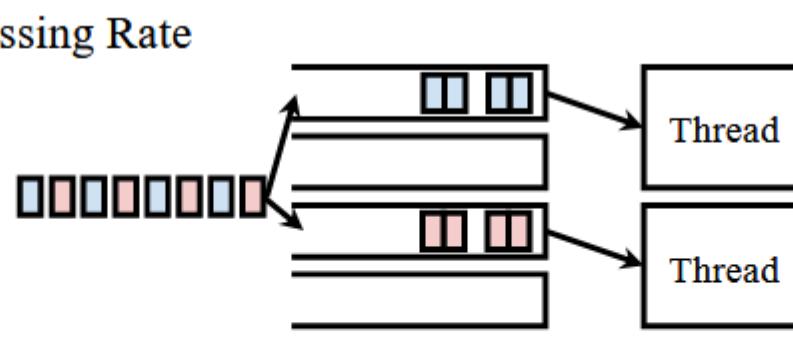
# How to Control Many Flows -- Scalability

□ 处理线程的数量不需要随着QP(流)的数量进行扩大

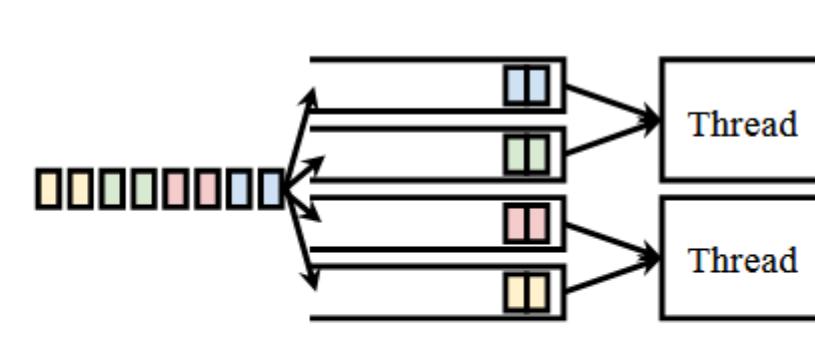
$$\begin{aligned} LineRate &= \sum_{i=1}^N BytesSentPerEvent_i \times EventRate_i \\ &= Bspe \times \sum_{i=1}^N EventRate_i \end{aligned}$$



1 QP : 2 Threads  
Granularity = 4



2 QPs : 2 Threads  
Granularity = 2



4 QPs : 2 Threads  
Granularity = 2

# Evaluation -- Setting

## □ 实验设置

- 每个节点的QP (流)的数量最多为1024个
- 数据包大小为64KB, MTU为1500B
- 吞吐速率为100Gbps
- DPA线程数为64

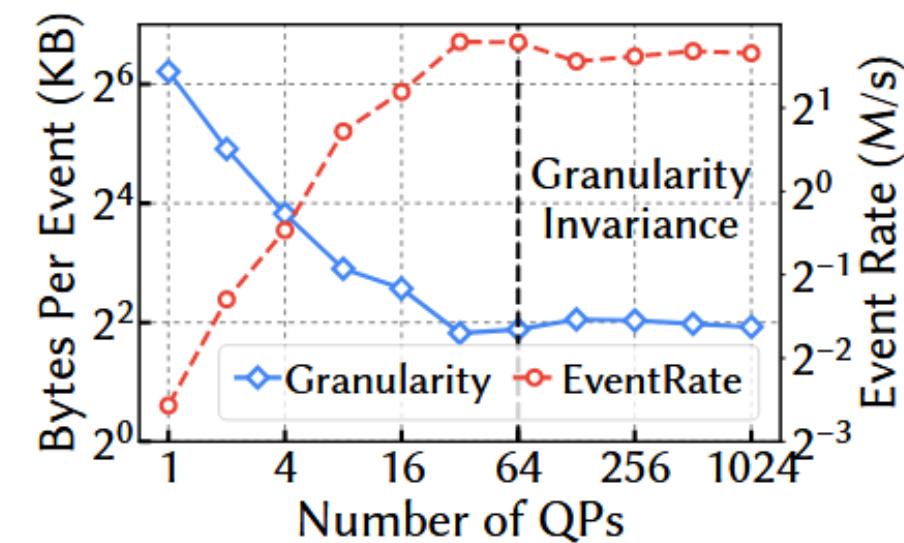
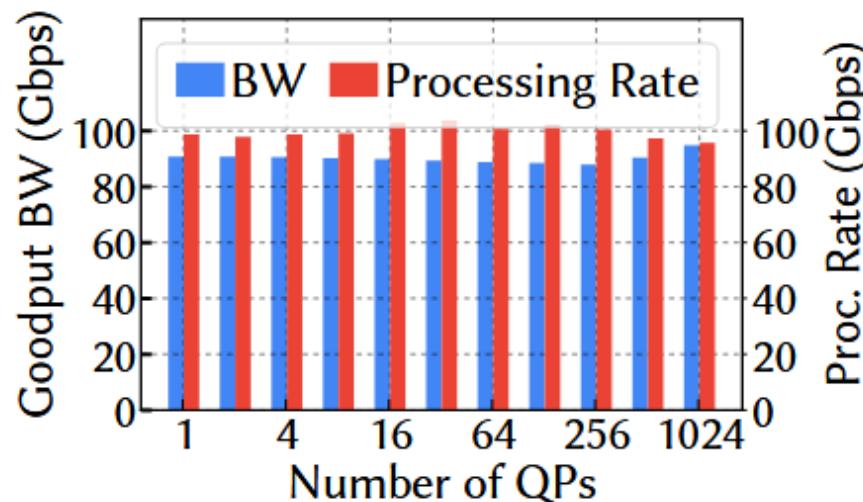
## □ 评价指标

- Event Rate: 事件 (信号) 处理速率 (million/s)
- Coalescing Granularity: 数据包处理粒度
- Processing Rate: Event Rate × Coalescing Granularity

# Evaluation -- Efficiency and Granularity

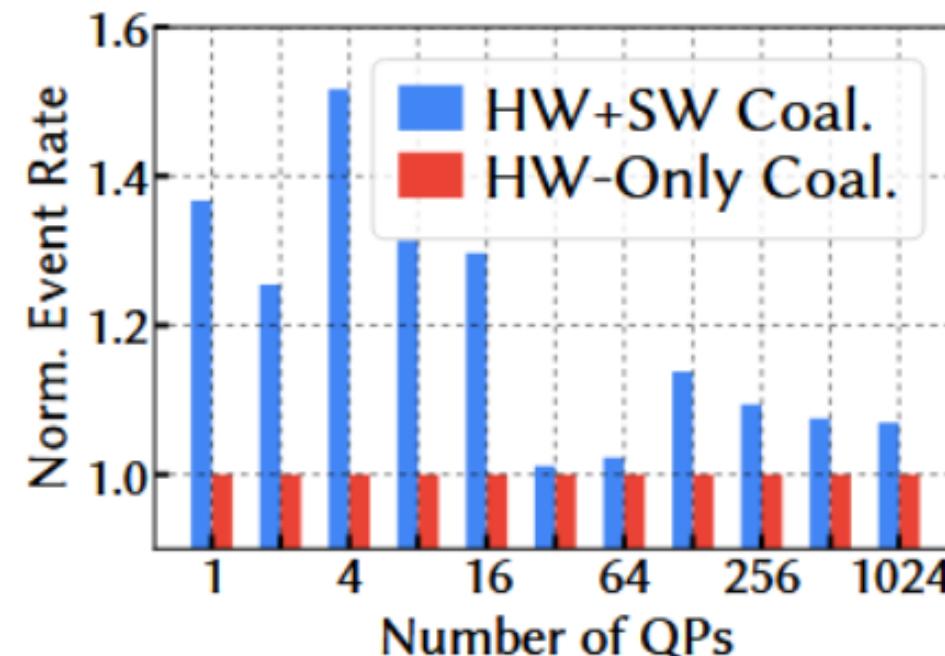
## □ 处理速率与粒度

- 处理速率与数据包带宽(BW)一致
- QP数大于64后，粒度保持不变
  - 粒度稳定时为4，符合包级粒度控制



# Evaluation -- Software Coalesce

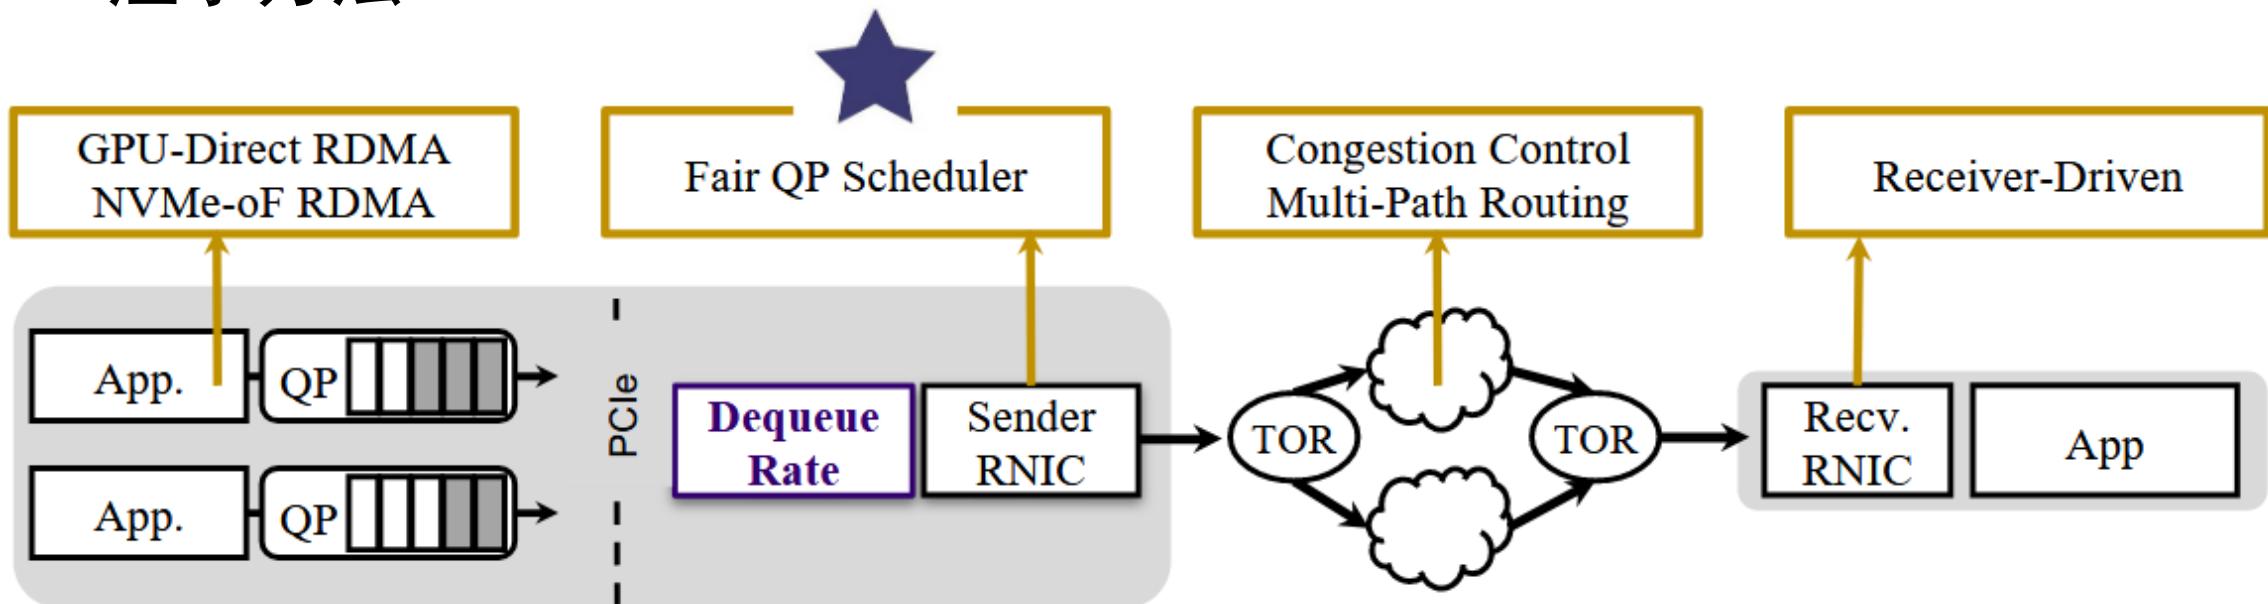
- 软硬件合并，事件（信号）处理速率
  - 软件(SW)合并后的事件处理速率相对于硬件(HW)合并有所提升
    - QP数较小时提升倍率更为明显



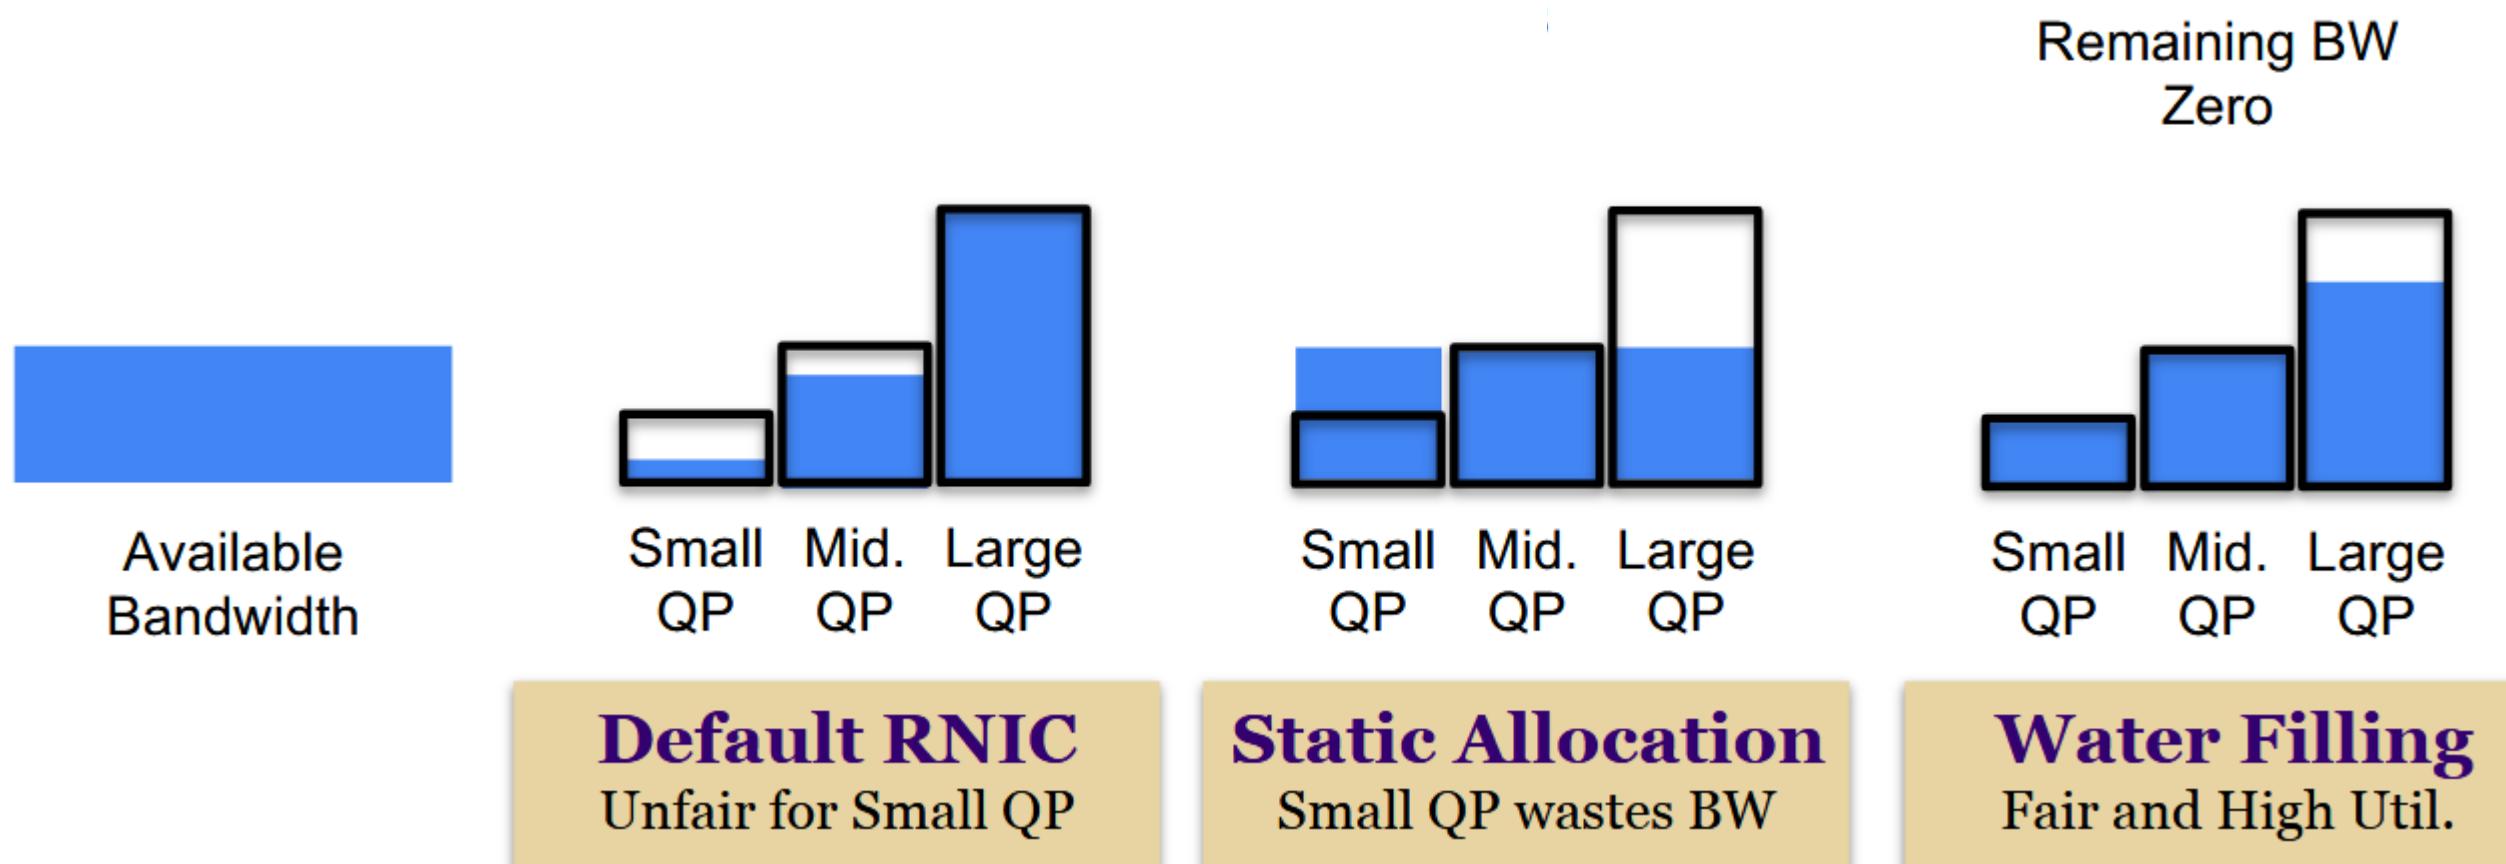
# Evaluation -- New Control Laws

□ 以公平队列调度为例

- 长流优先（默认方法）
- 静态分配
- 注水方法



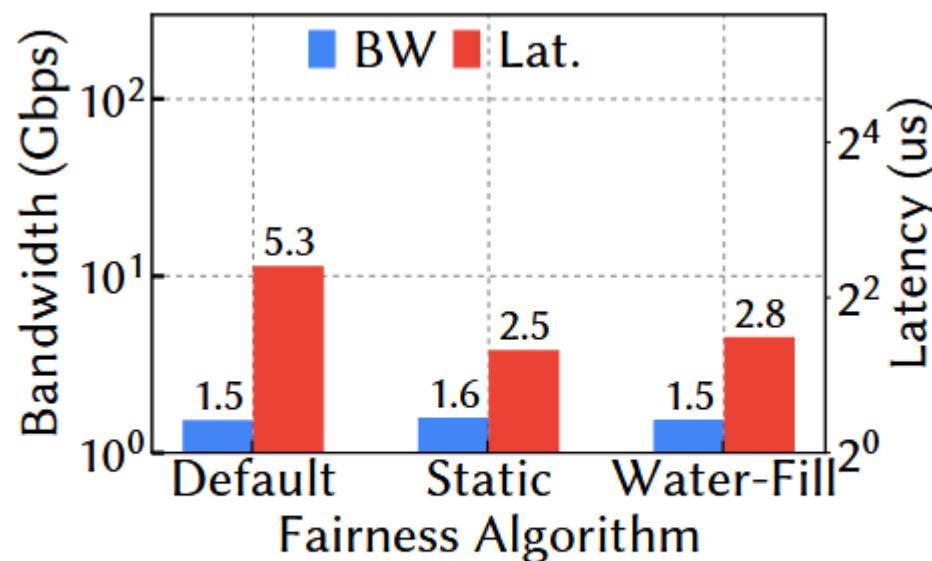
# Evaluation -- Fair QP Scheduler



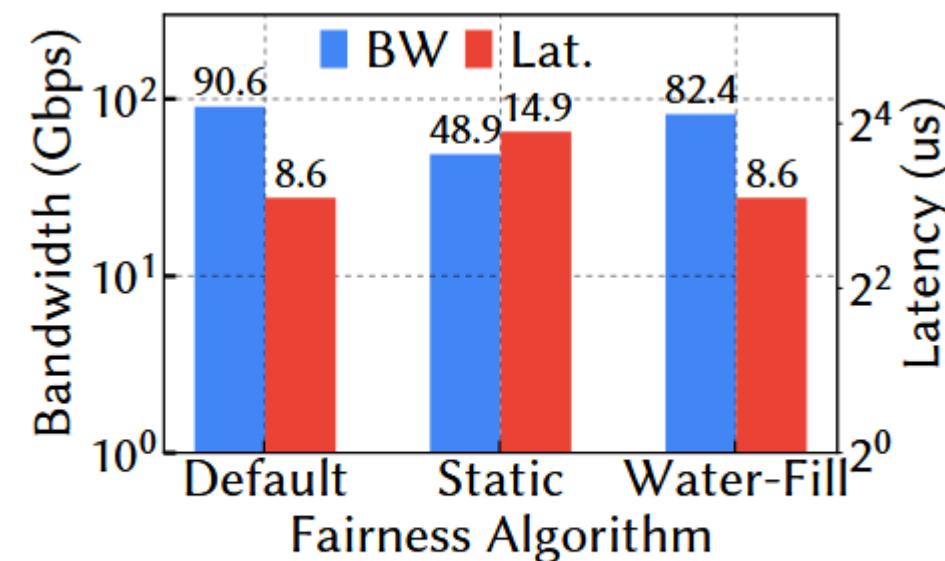
# Evaluation -- Fair QP Scheduler

## □ 公平队列调度——注水方法

- 延迟敏感QP——延迟优于默认方法
- 带宽敏感QP——带宽吞吐优于静态分配方法



(a) Latency-Sensitive QP



(b) Bandwidth-Hungry QP

## □ 能否提高

- 收集链路信号为局部信号，可以在RNIC内存中留存一块存储全局链路信号，并通过链路传输或者数据包携带全局信息定期更新，实现类似背压树的全局控制图
- CPU在运行过程中只用于辅助进程间通信，可以利用CPU处理低频率单线程的全局信号，从而避免大规模网络下，全局信息读写对处理粒度产生影响

## □ 与此类似

- 绝大多数流控机制的数据面与控制面强关联，需要接收对端信号后才会实现触发式流控。可以采用论文中提到的隔离控制面的方式，实现更灵活的流控方法

## □ 问题泛化

- 处理粒度取决于物理硬件的吞吐速率和线程处理速率，硬件依赖性强，泛化难度高



# Q & A

**Presenter: Huang Kai**