

WLB-LLM: Workload-Balanced 4D Parallelism for Large Language Model Training

**Zheng Wang, Anna Cai, Xinfeng Xie, Zaifeng Pan, Yue Guan, Weiwei Chu,
Jie Wang, Shikai Li, Jianyu Huang, Chris Cai, Yuchen Hao, Yufei Ding**

Presenter: Yi Liu



Author

Research Direction:

- high-performance system design
- end-to-end optimization for deep learning

Selected Publications:

- GMI-DRL(ATC'25)
- WLB-LLM(OSDI'25)
- FastTree(MLSys'25)



Zheng Wang
<https://zhengwang.info/>

Outline

- Background
- Design
- Evaluation
- Thinking

Background: Training LLM is costly

The scale of LLMs grow larger → Training is costly

Estimated training cost and compute of select AI models

Source: Epoch AI, 2024 | Chart: 2025 AI Index report

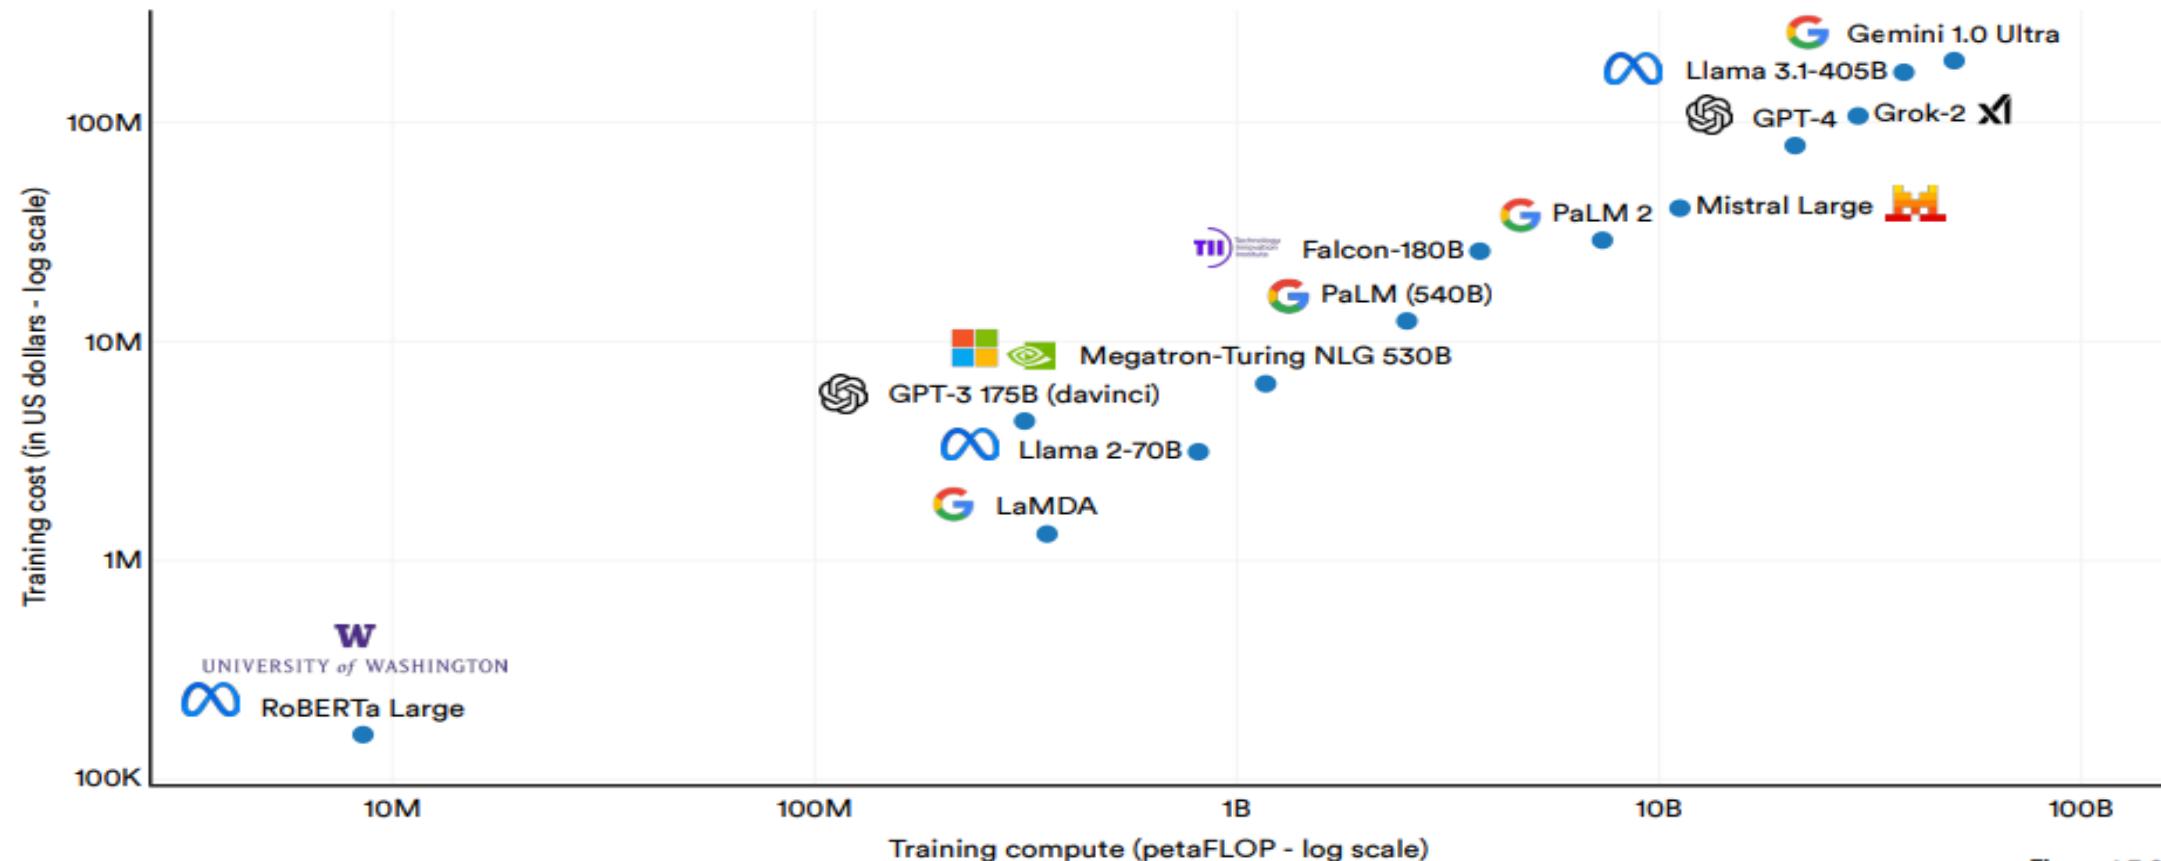
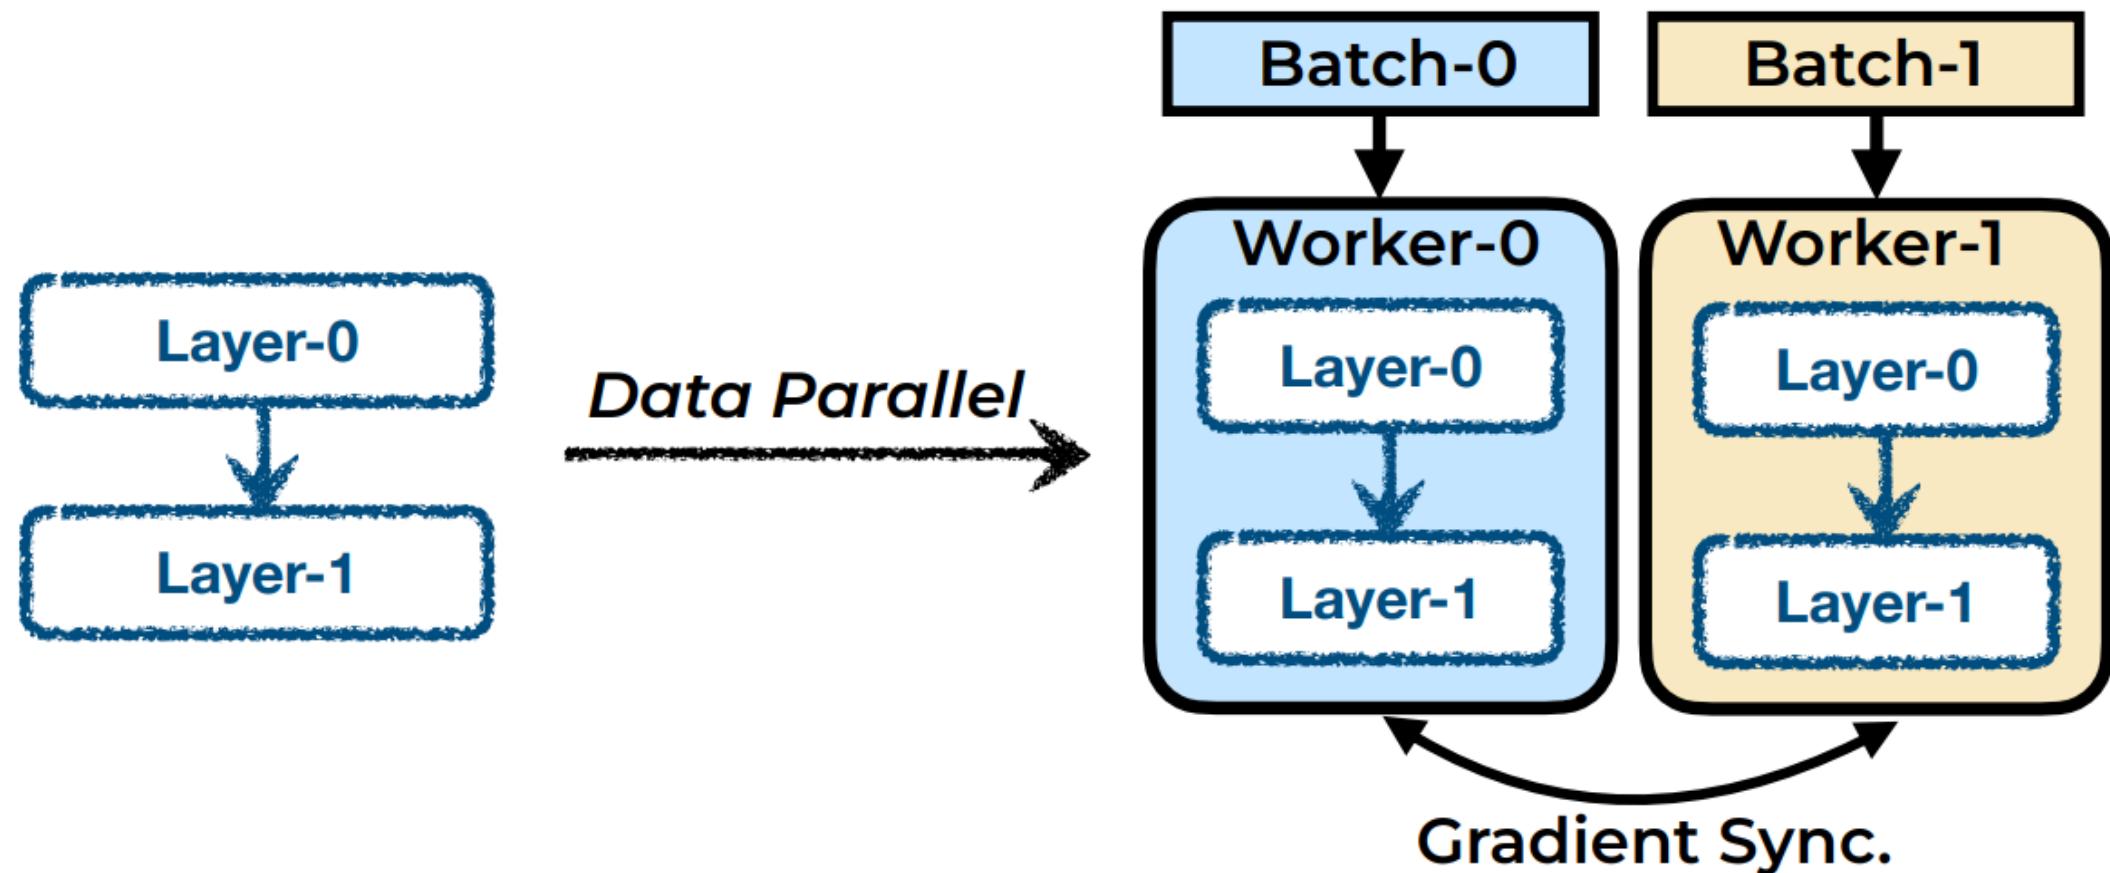


Figure 1.3.26

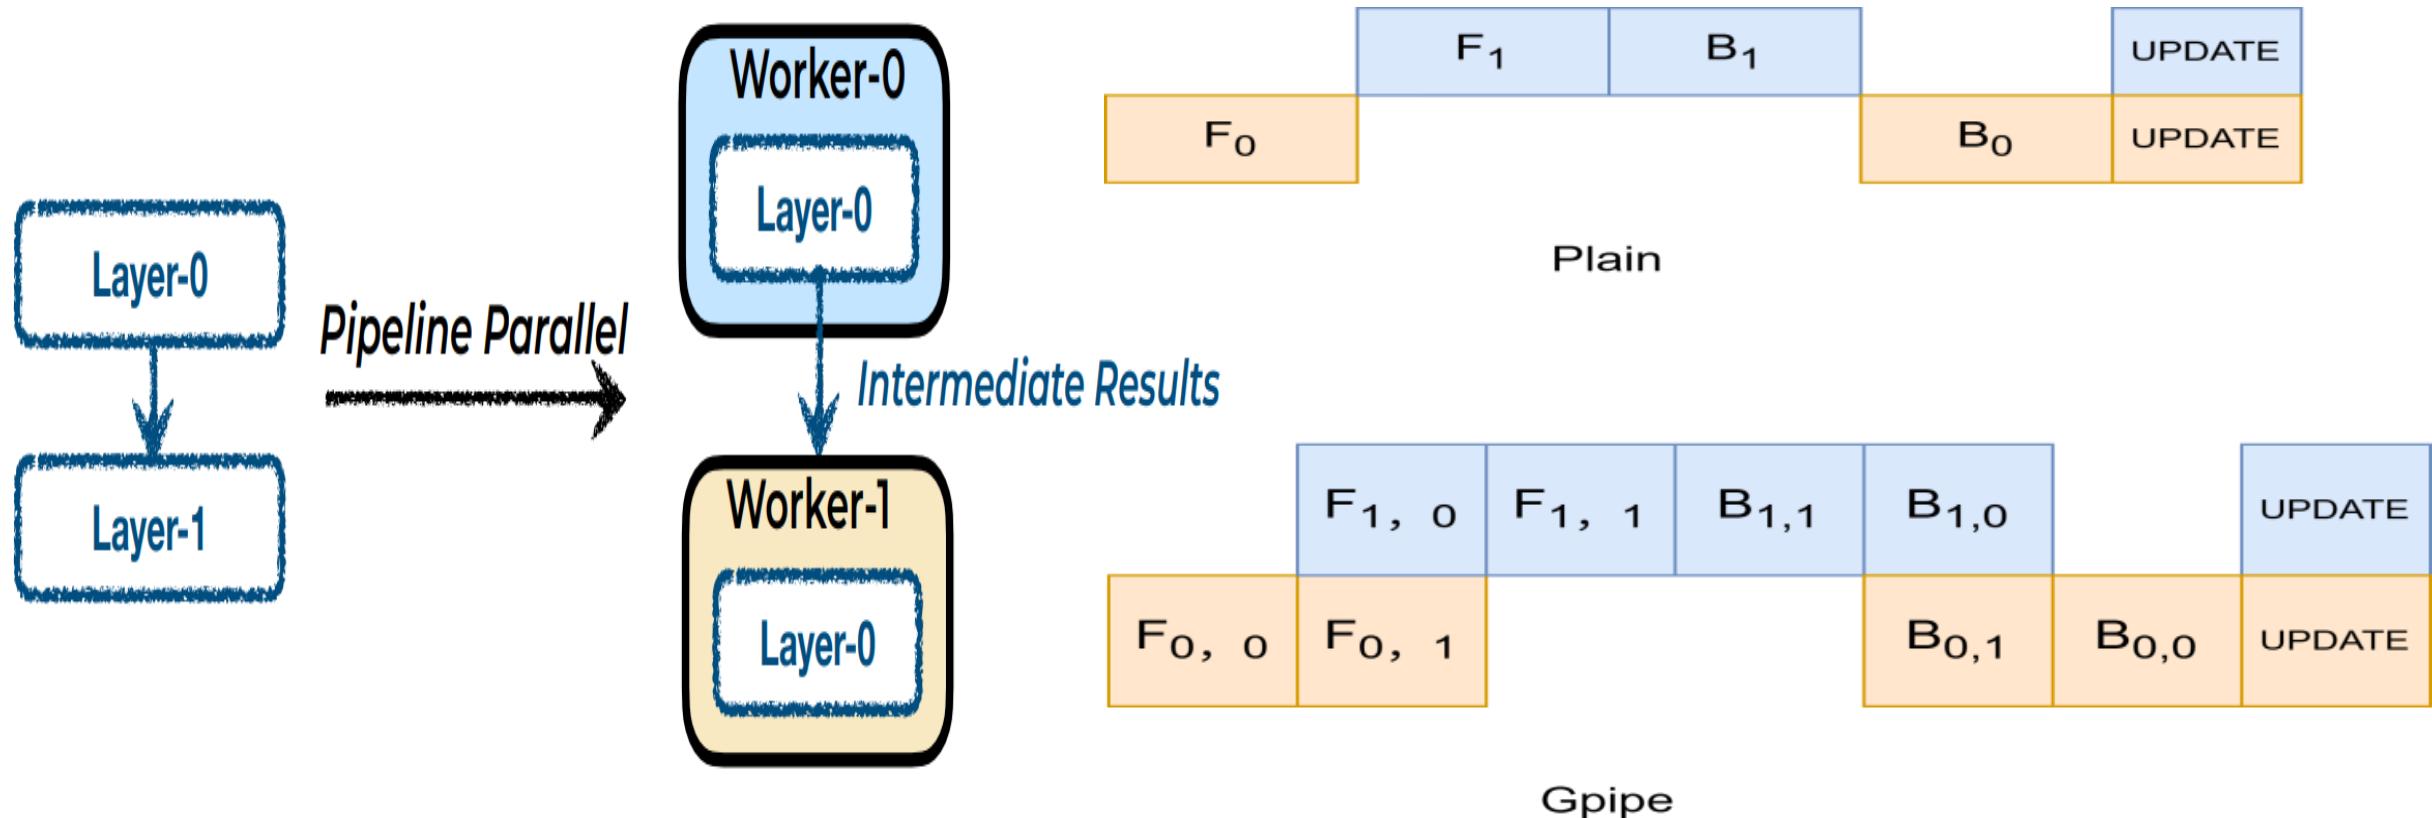
Background: Data Parallelism

Duplicate models to different workers



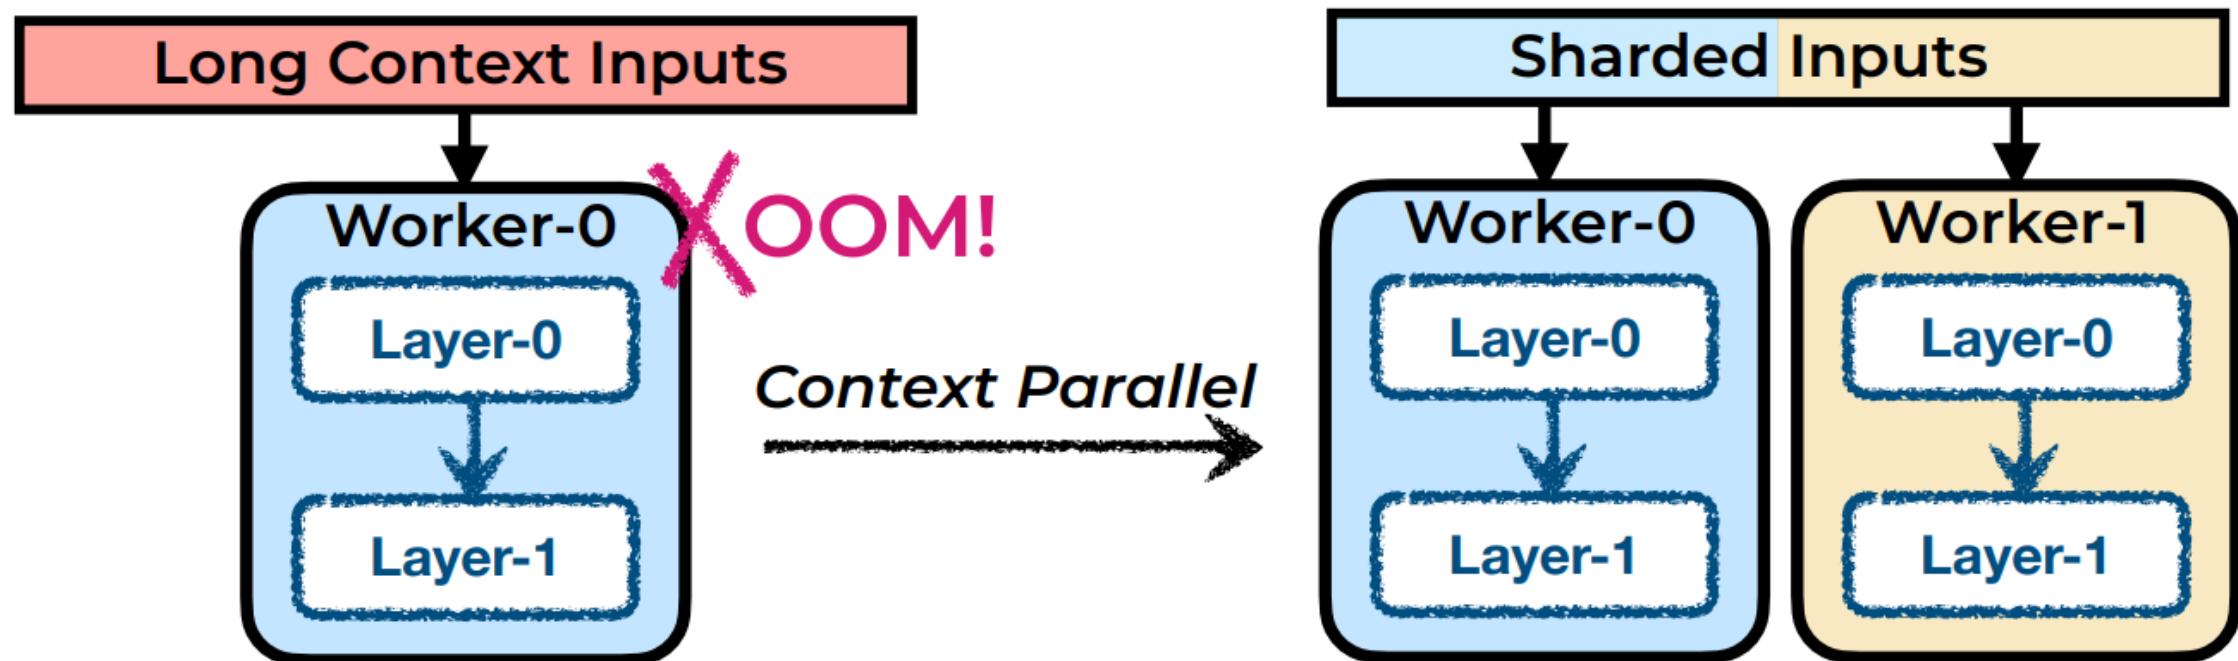
Background: Pipeline Parallelism

Split model in a layer-wise manner



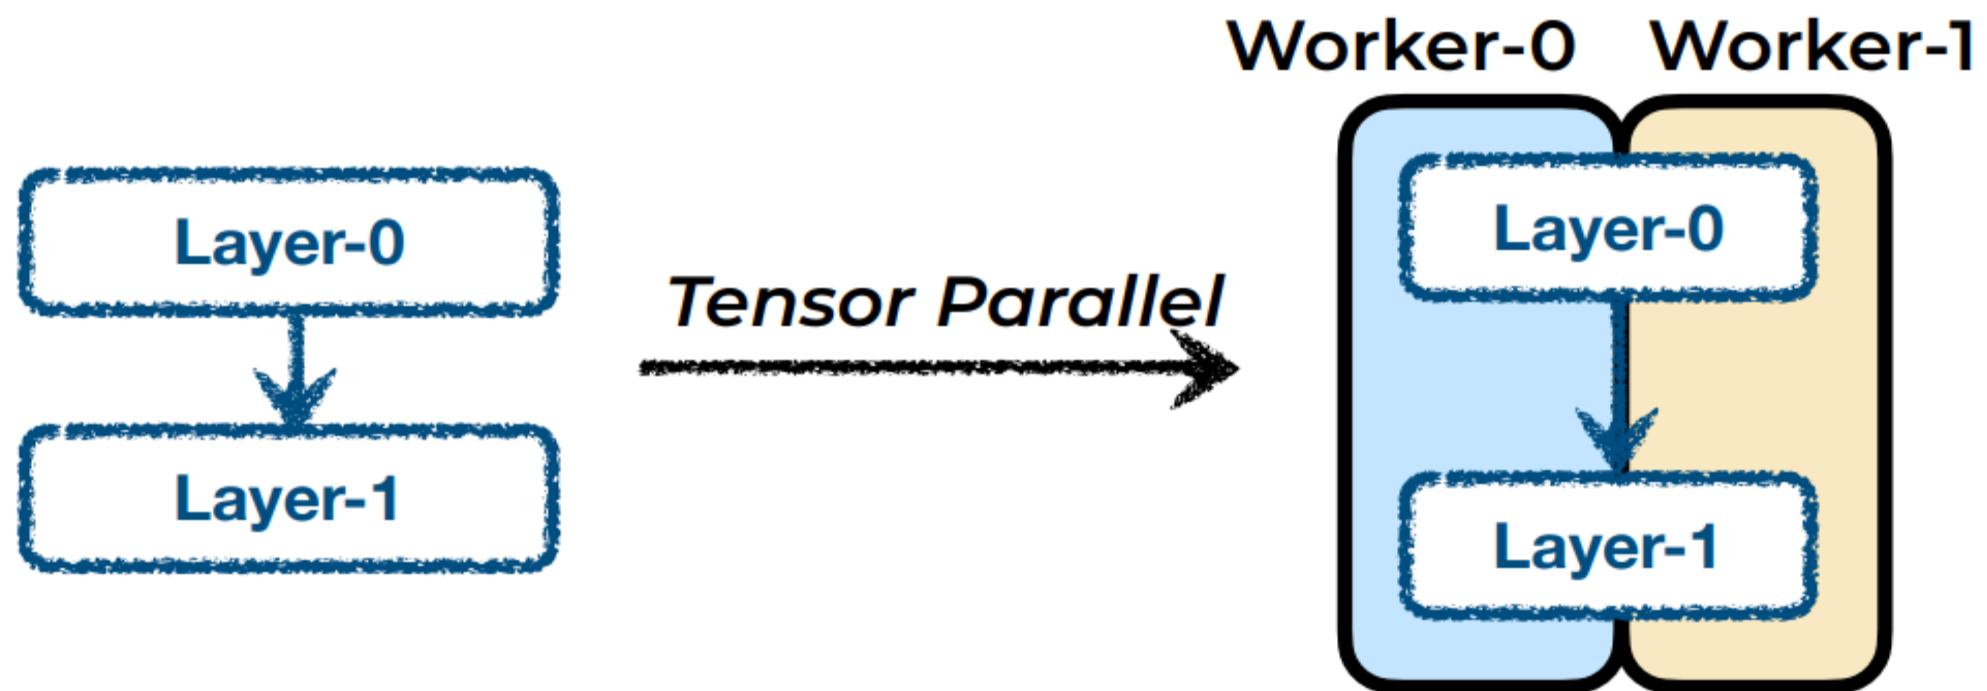
Background: Context Parallelism

Shard the long-context input along sequence length dimension



Background: Tensor Parallelism

Split model parameters within a layer

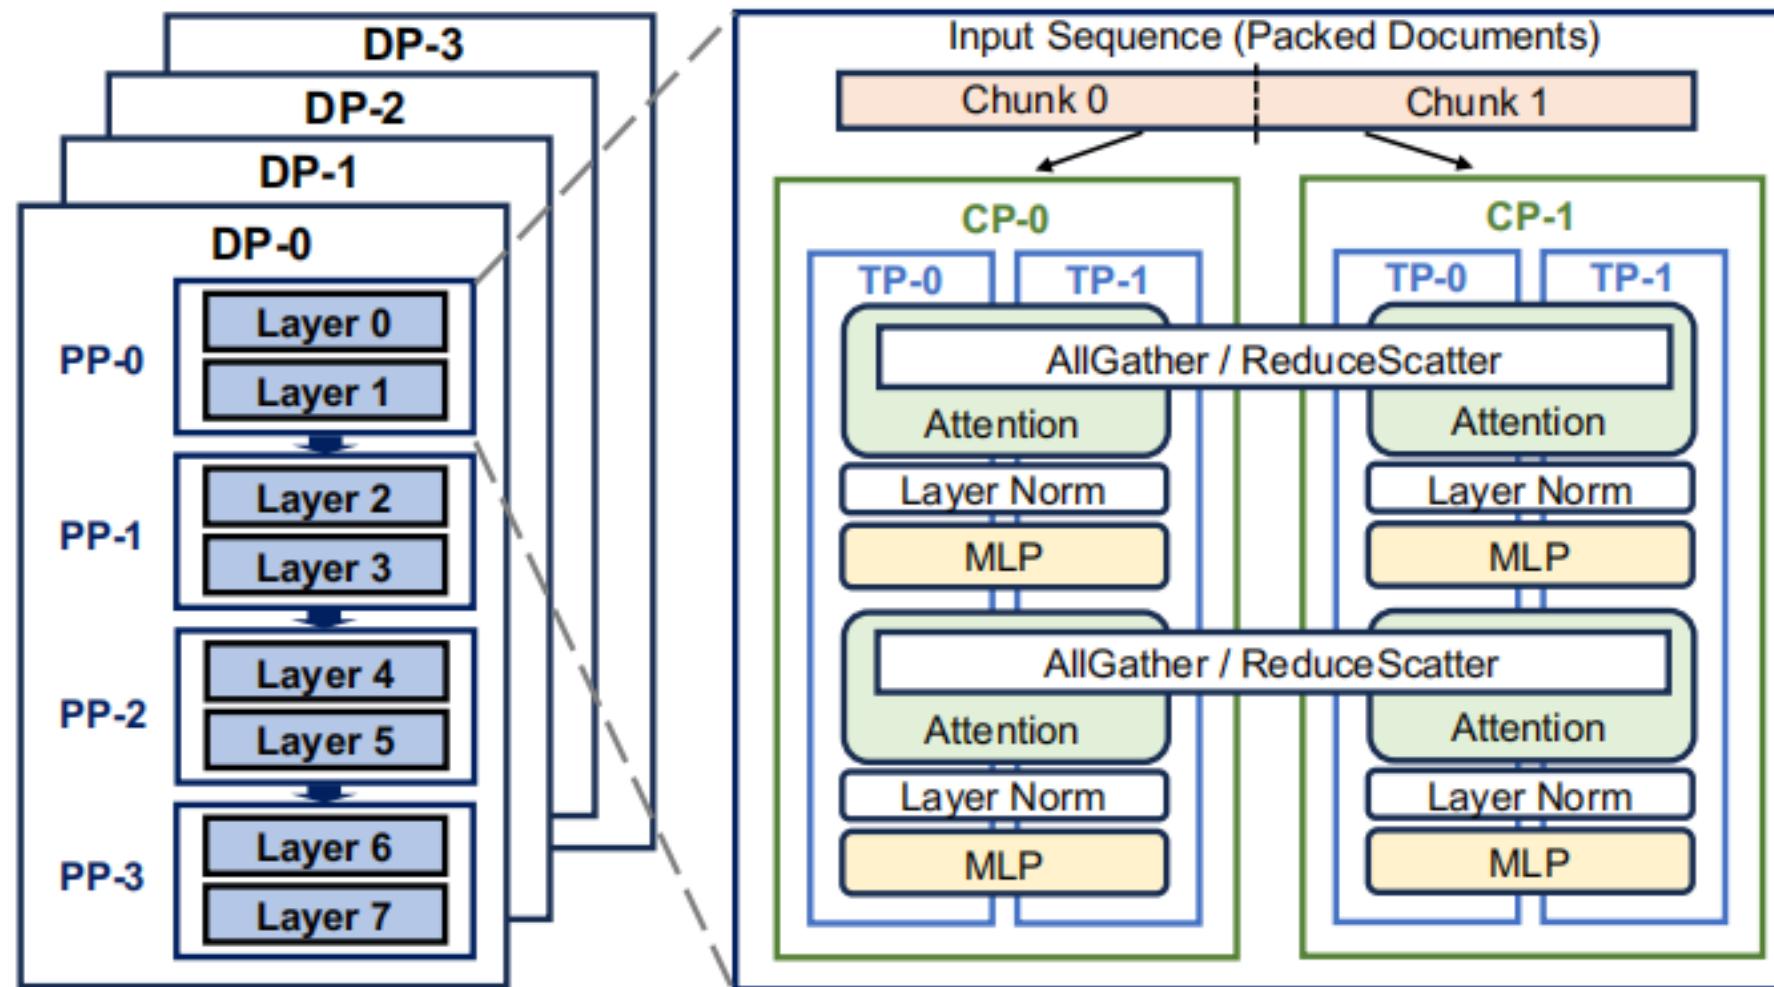


Background: 4D parallelism

Training is costly

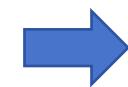


multi level parallelism

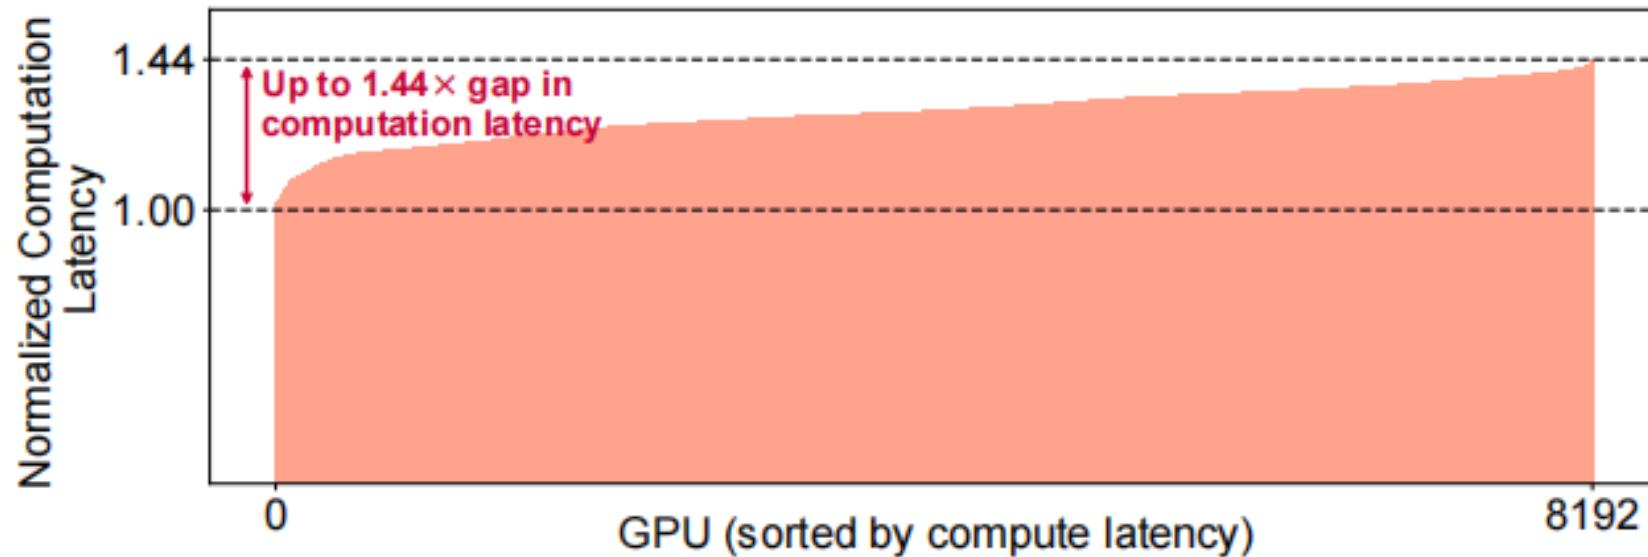


Background: workload imbalance

not fully utilizing GPUs



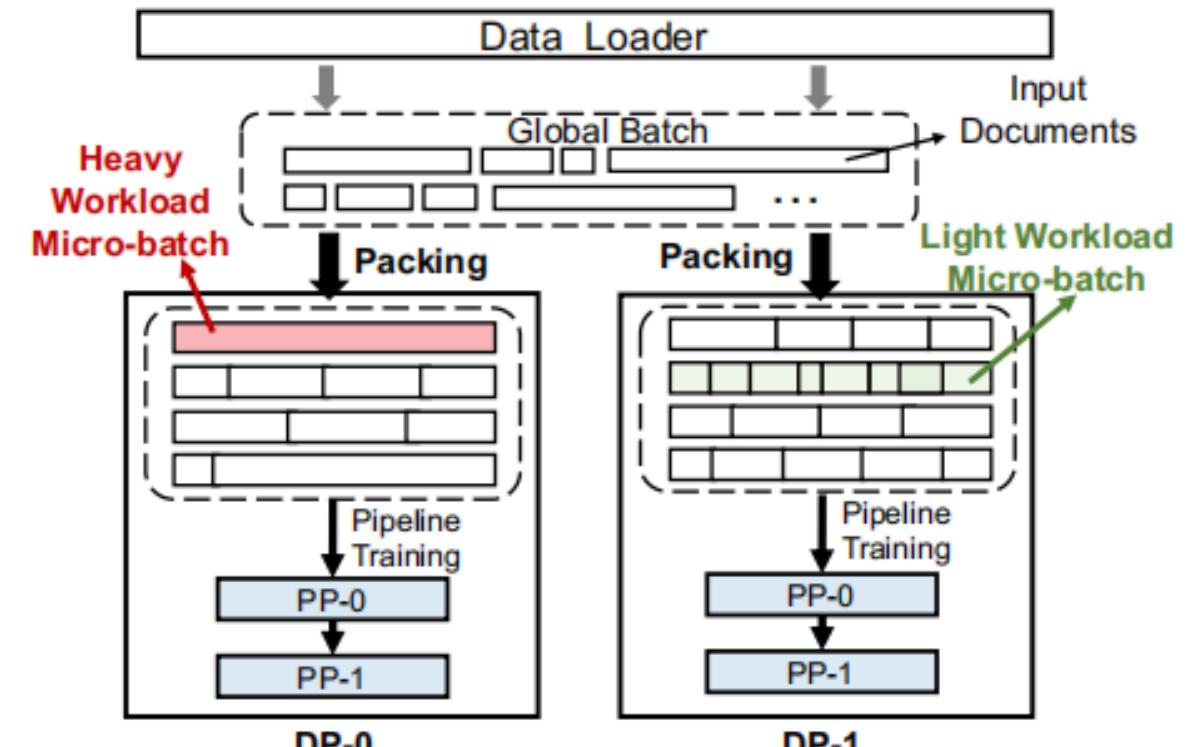
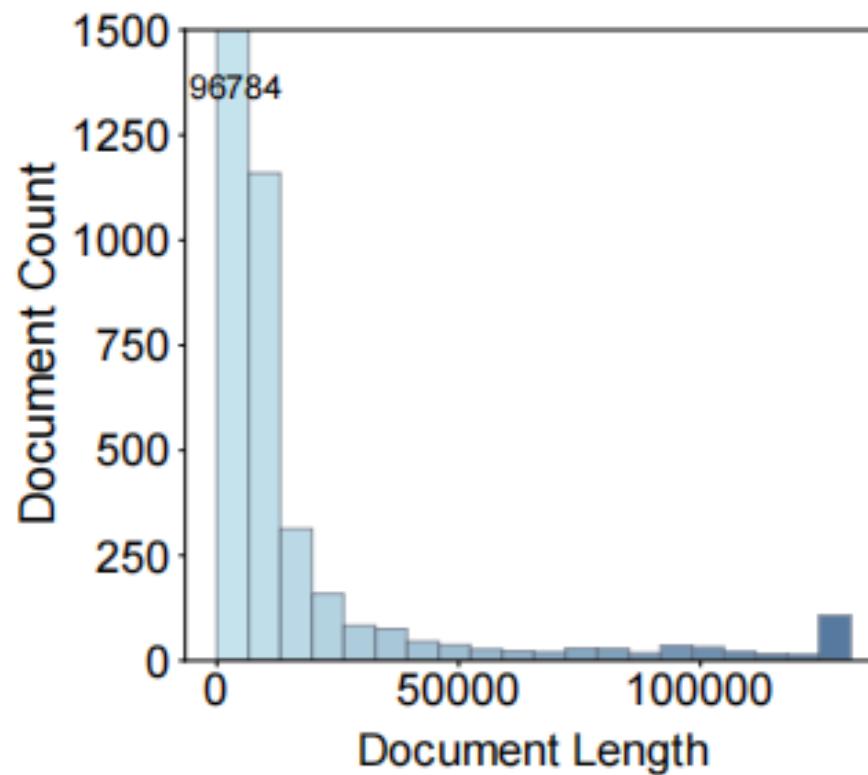
1. the causes of workload imbalance
2. how to resolve the imbalance



(a) Normalized computation latency in an 8K-GPU LLM training job.

Background: PP level imbalance

varing input document length → imbalance in parallelism



(1) Imbalance from PP-level Document Packing

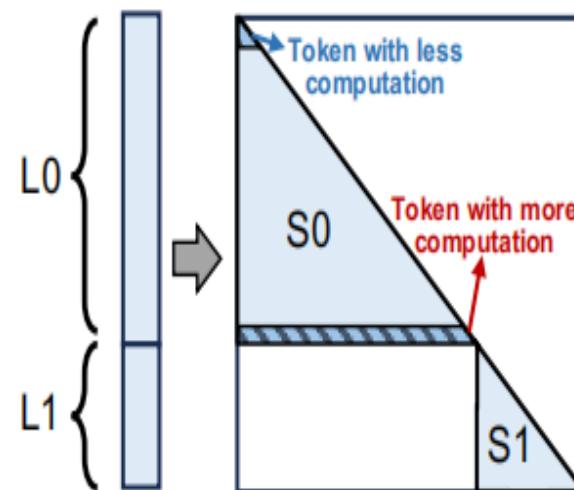
Background: reason of imbalance

varing computation



imbalance across PP workers

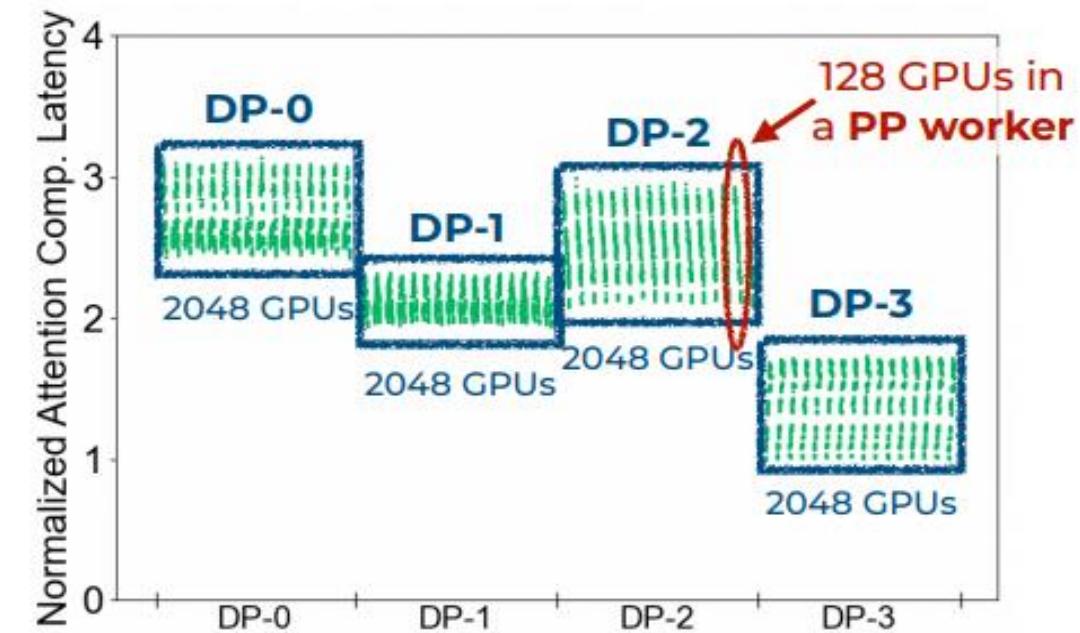
Packed Docs Attention Mask



Document lengths: $L_0 + L_1 = L_2 + L_3 + L_4 + L_5$

Computation (triangle areas): $S_0 + S_1 \gg S_2 + S_3 + S_4 + S_5$

Normalized Attention Computation Latency
(8K GPU: DP-4, PP-16, CP-16, TP-8)

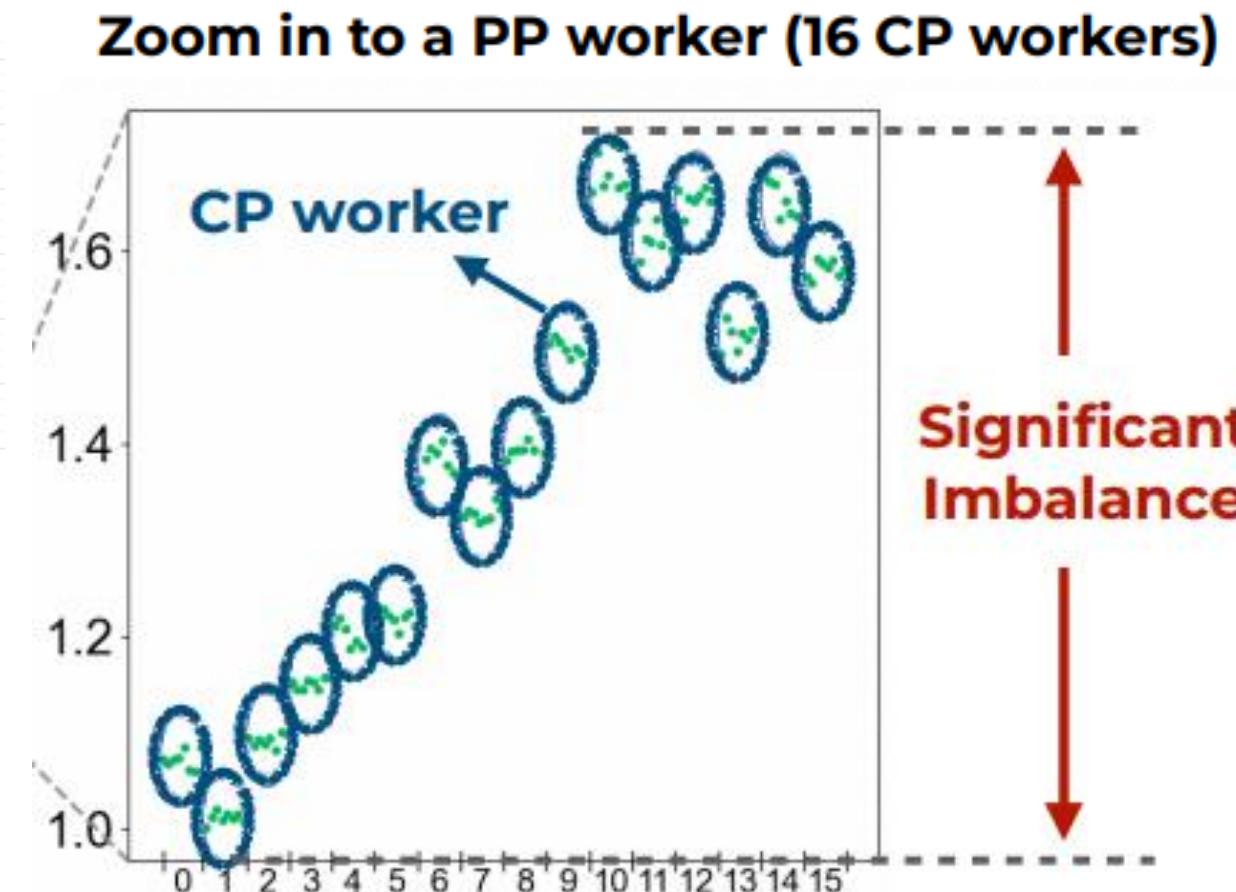
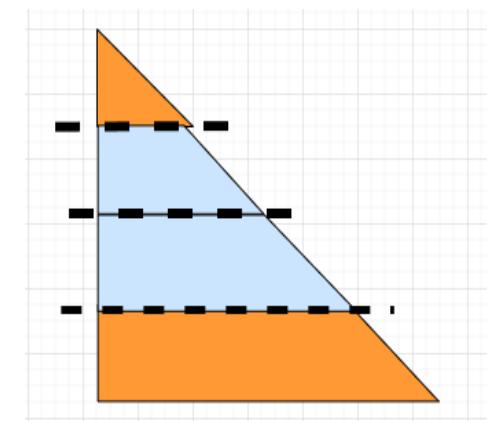
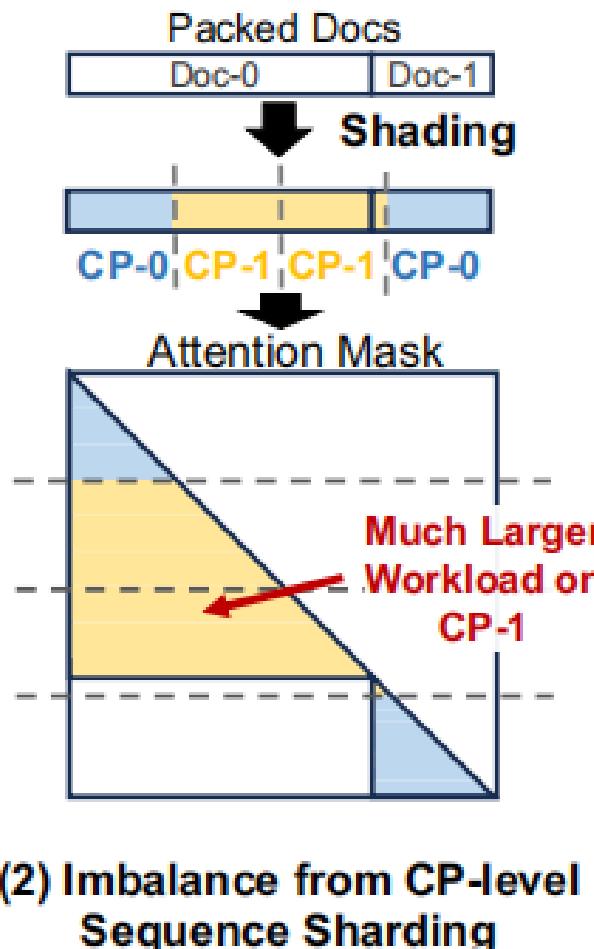


Background: CP level imbalance

sequence sharding

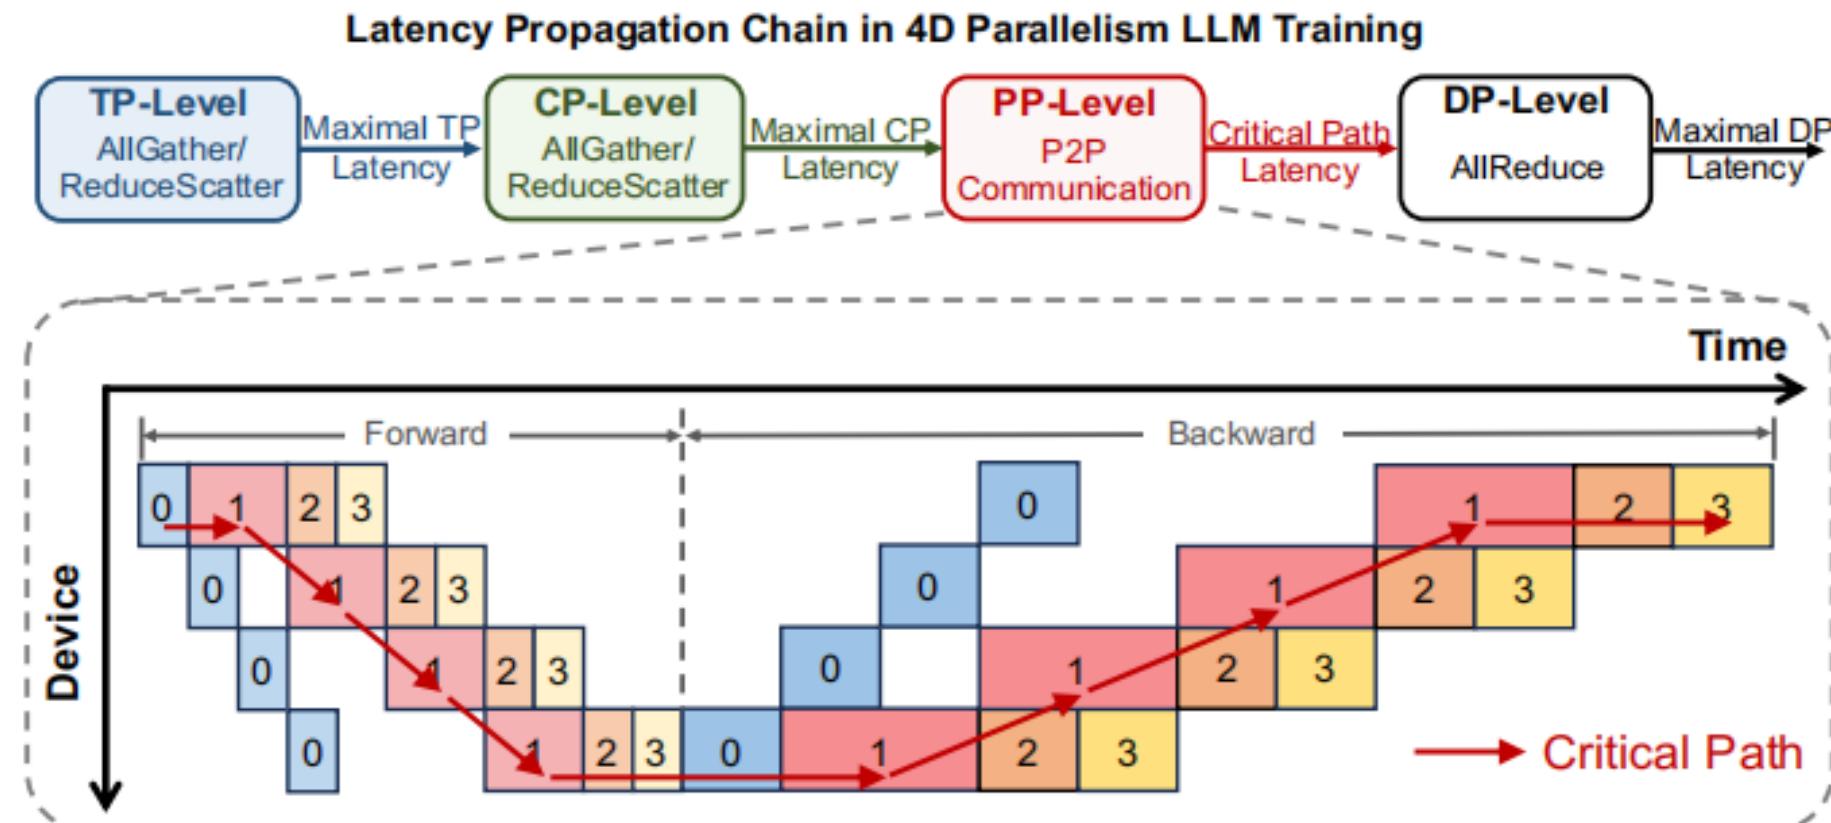


imbalance across CP workers



Background: influence of imbalance

imbalance will be accumulated → end-to-end training latency



Design:WLB-LLM

PP-level



- 1.variance-length packing
- 2.heuristic outlier document delay optimization

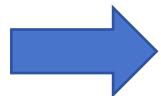
CP-level



fine-grained and adaptive sharding optimization

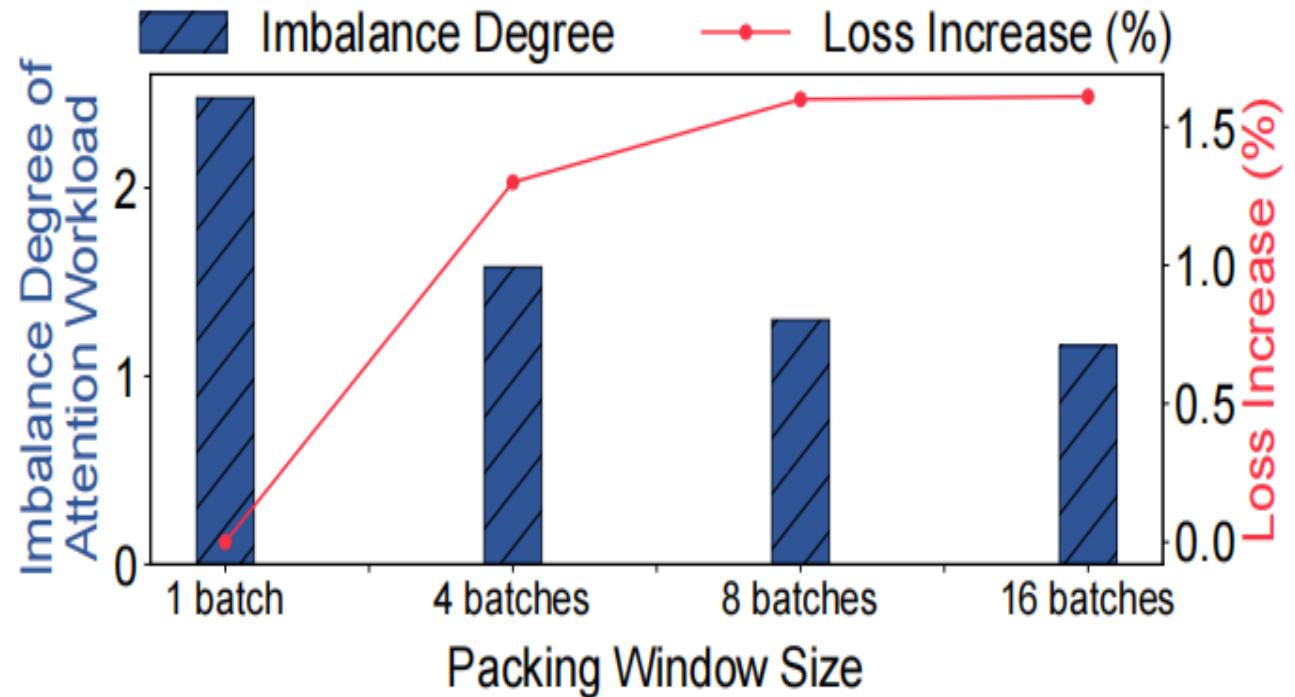
Design:Fixed-length

Fixed-length packing



Loss increase

$$\begin{aligned} \text{minimize} \quad & \max\left(\sum_{i=1}^N x_{ij} \cdot d_i^2\right), j = 1, \dots, M \\ \text{subject to} \quad & \sum_{j=1}^M x_{ij} = 1, \quad i = 1, \dots, N \\ & \sum_{i=1}^N x_{ij} \cdot d_i \leq L, \quad j = 1, \dots, M \\ & x_{ij} \in \{0, 1\} \end{aligned}$$



Design: Variable-length Packing

Fixed context length

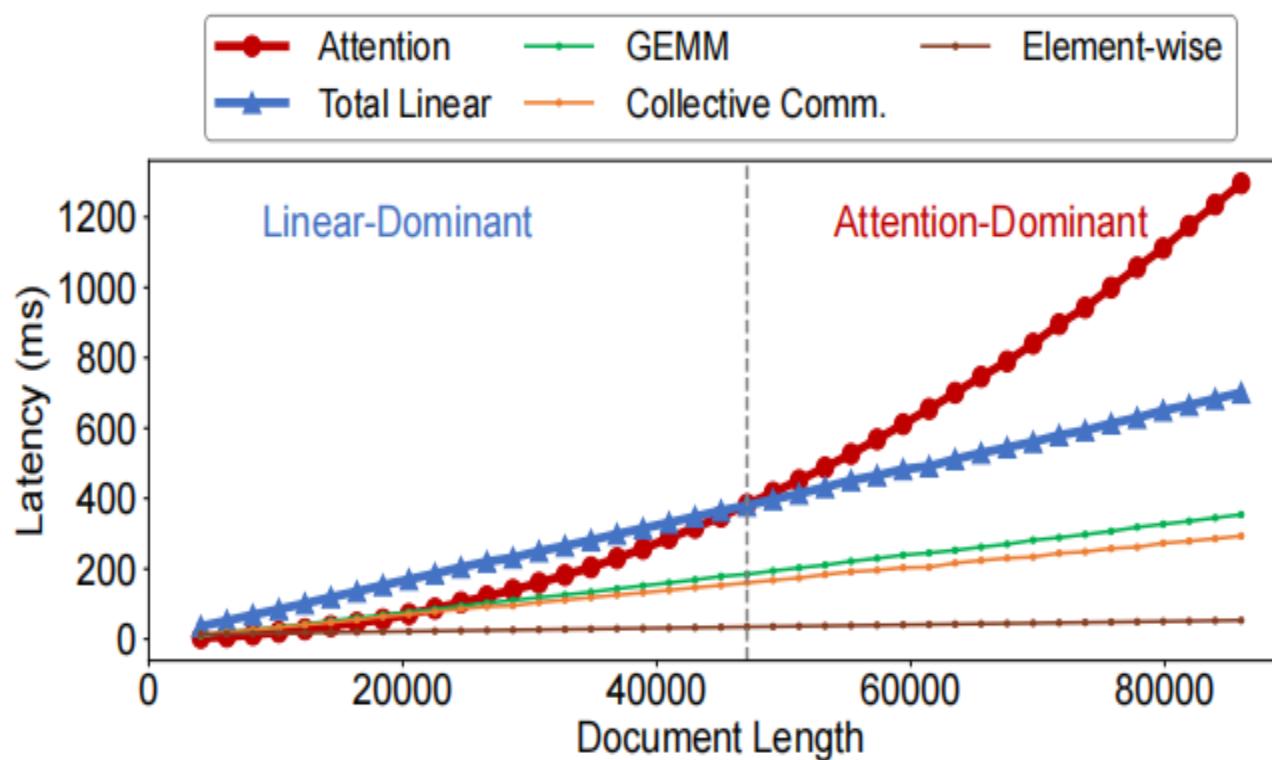


variable length

balance attention workload



balance attention and linear



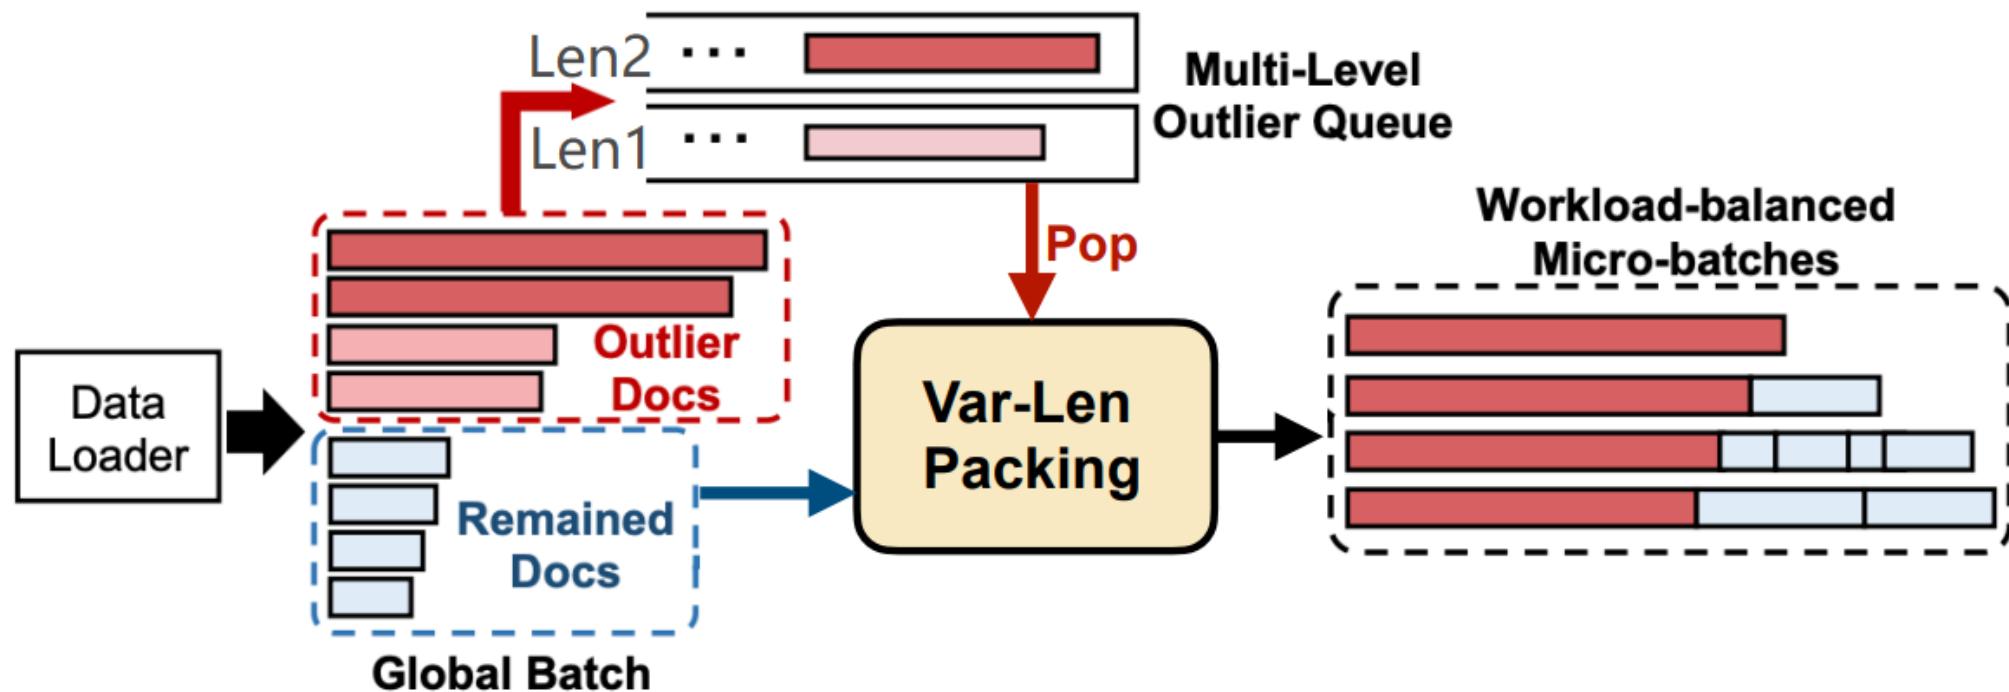
$$\begin{aligned} \text{minimize} \quad & \max\left(\sum_{i=1}^N (W_a(x_{ij} \cdot d_i) + W_l(x_{ij} \cdot d_i))\right), \\ \text{subject to} \quad & \sum_{j=1}^M x_{ij} = 1, \quad i = 1, \dots, N \\ & \sum_{i=1}^N x_{ij} \cdot d_i \leq L_{max}, \quad j = 1, \dots, M \\ & x_{ij} \in \{0, 1\} \end{aligned}$$

Design:Outlier document delay

short document insufficient



delay outlier document



Design: Heuristic Packing Algorithm

ILP solver takes a long time → heuristic packing algorithm

Algorithm 1: Heuristic Var-length Packing Algorithm

```
input : Dataloader:  $D$ , Waiting Queues:  $Q$ ,  
       #Micro-Batch per Iteration:  $N$ ,  
       Sequence Length Upper bound:  $L_{max}$   
output: Packed Micro-Batches for Training:  $B$   
1  $Remained\_Doc = []$ ;  
2 for  $Cur\_Batch$  in  $D$  do  
3    $Doc\_Set = Remained\_Doc$ ;  
4   for  $Doc$  in  $Cur\_Batch$  do  
5     /* Delay the execution of outlier documents. */  
6     if  $Doc.Is\_Outlier()$  then  
7       |  $Q.Add(Doc)$ ;  
8     else  
9       |  $Doc\_Set.Push(Doc)$ ;  
10    end  
11  end  
12  for  $q$  in  $Q$  do  
13    if  $len(q) \geq N$  then  
14      /* Pop outlier documents for the current batch. */  
15      |  $Doc\_Set.Push(q.Pop(N))$ ;  
16    end  
17  end  
18  /* Sort the documents in descending order by length.  
19   $Doc\_Set.Sort\_by\_Length()$ ;
```

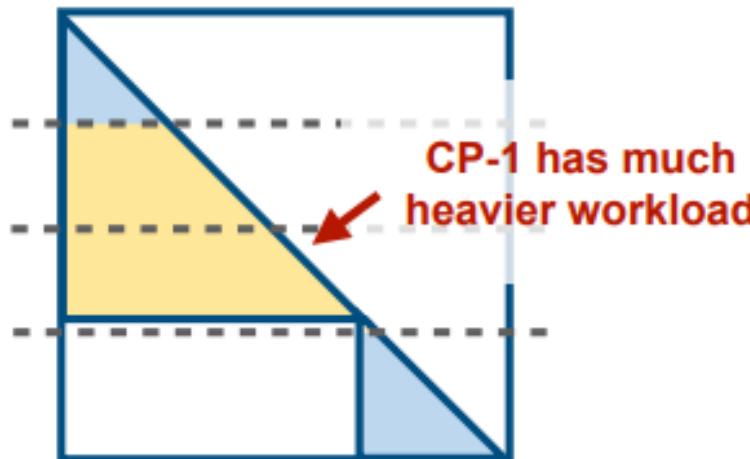
```
17  /* Start packing. */  
18   $New\_Batch = Create\_Batch(N)$ ;  
19  for  $Doc$  in  $Doc\_Set$  do  
20    /* Get micro-batches with minimum workload/length. */  
21     $W\_idx = New\_Batch.Get\_Min\_Workload()$ ;  
22     $L\_idx = New\_Batch.Get\_Min\_Length()$ ;  
23    if  $New\_Batch[W\_idx].Len() + Doc.Len() < L_{max}$  then  
24      |  $New\_Batch[W\_idx].Push(Doc)$ ;  
25    else  
26      |  $New\_Batch[L\_idx].Len() + Doc.Len() < L_{max}$  then  
27        | |  $New\_Batch[L\_idx].Push(Doc)$ ;  
28      else  
29        | |  $Remained\_Doc.Push(Doc)$ ;  
30      end  
31   $B.Push(New\_Batch)$ ;  
32 end
```

Design:per-document sharding

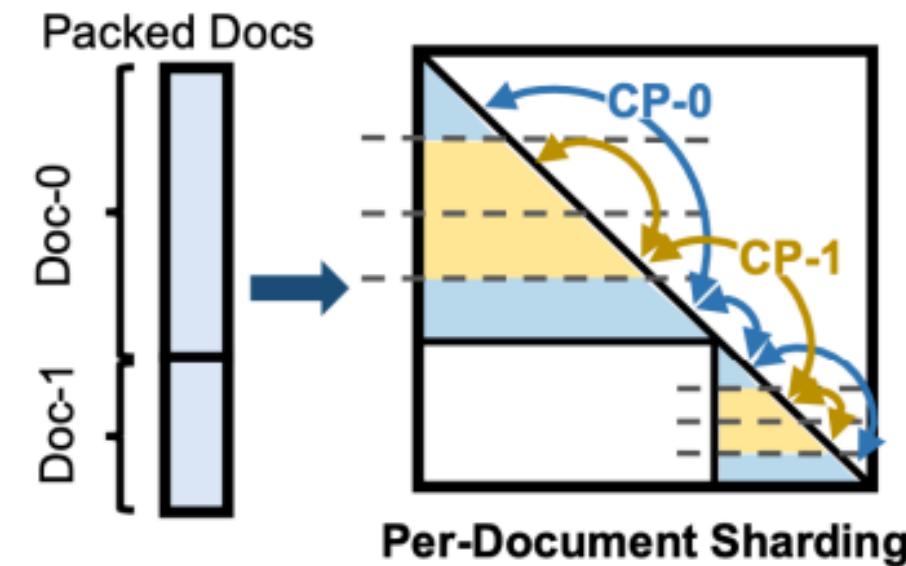
per-sequence sharding



per-document sharding



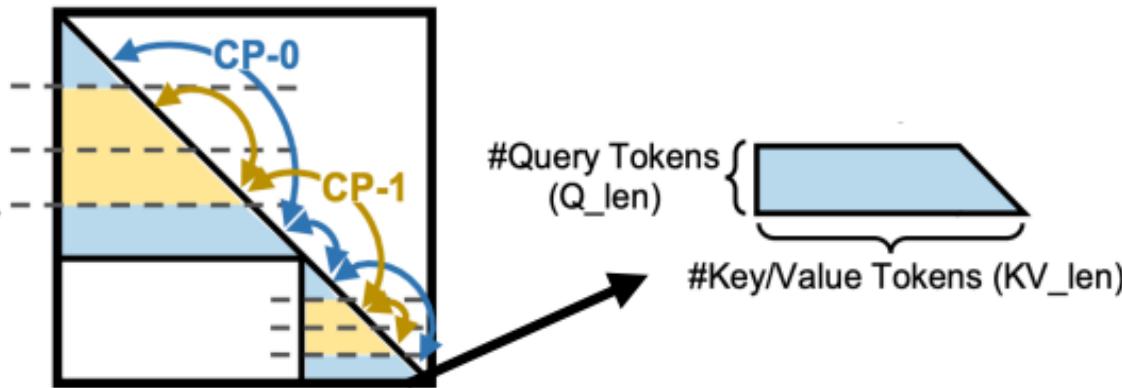
Per-Sequence Sharding



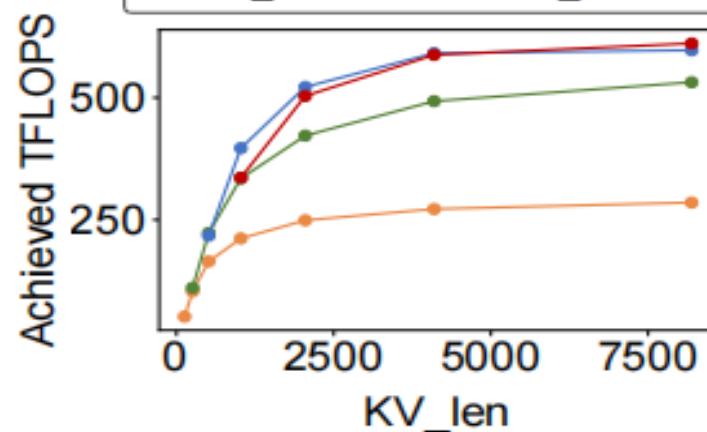
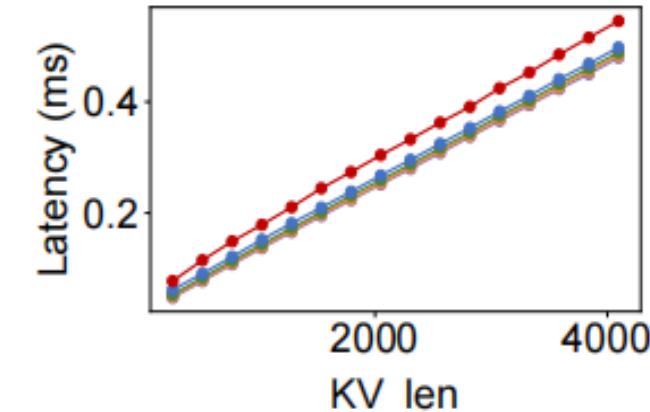
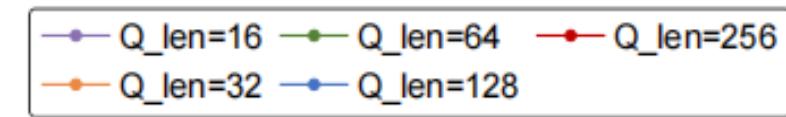
Per-Document Sharding

Design: per-document sharding

sharding balance vs. kernel efficiency

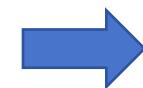


Fine-grained sharding may leads to small document shards

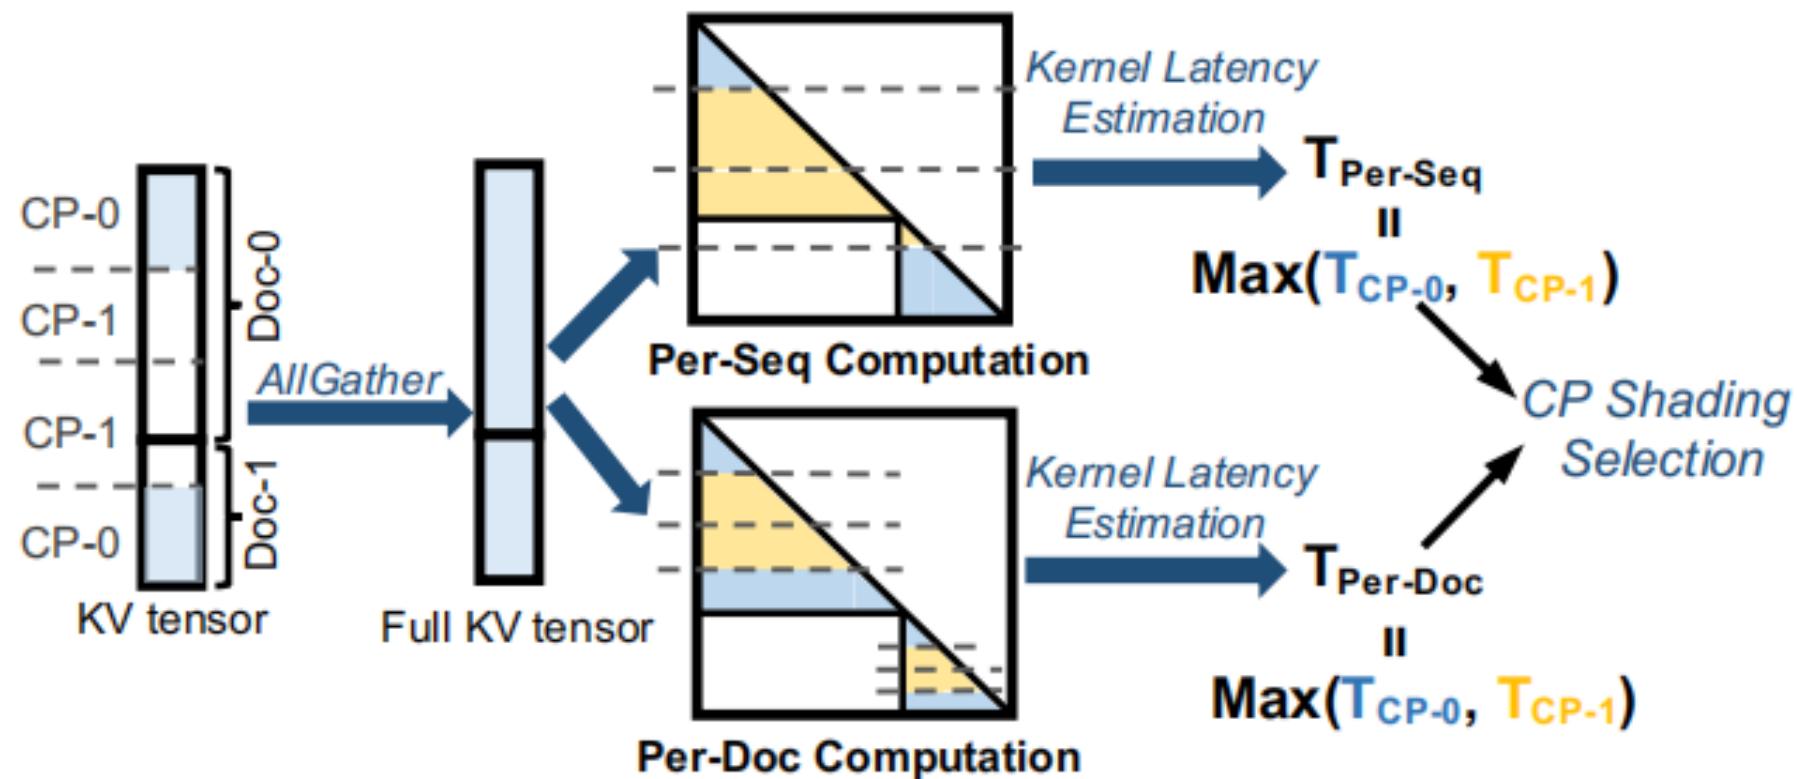


Design:adaptive sharding

trade-off



adaptive sharding



Evaluation: setting

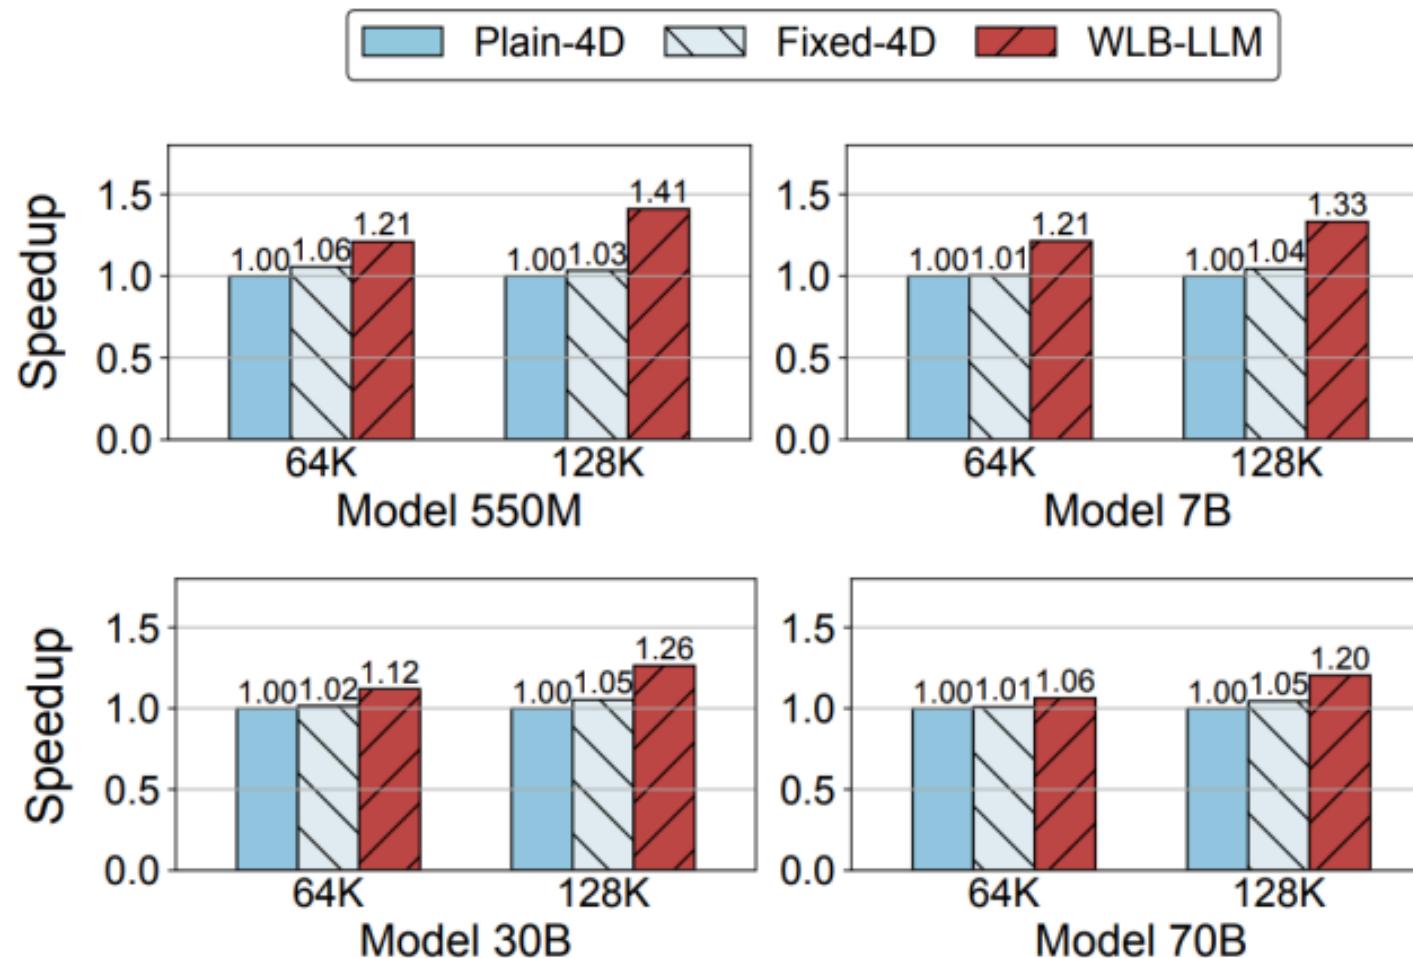
model and parallelism config

Model Size	Context Window	#GPU	4D Parallelism Configs (TP, CP, PP, DP)
550M	64K	32	(2, 2, 4, 2)
	128K	32	(2, 4, 4, 1)
7B	64K	32	(4, 2, 4, 1)
	128K	64	(8, 2, 4, 1)
30B	64K	64	(8, 2, 4, 1)
	128K	128	(8, 4, 4, 1)
70B	64K	256	(16, 4, 4, 1)
	128K	256	(16, 4, 4, 1)

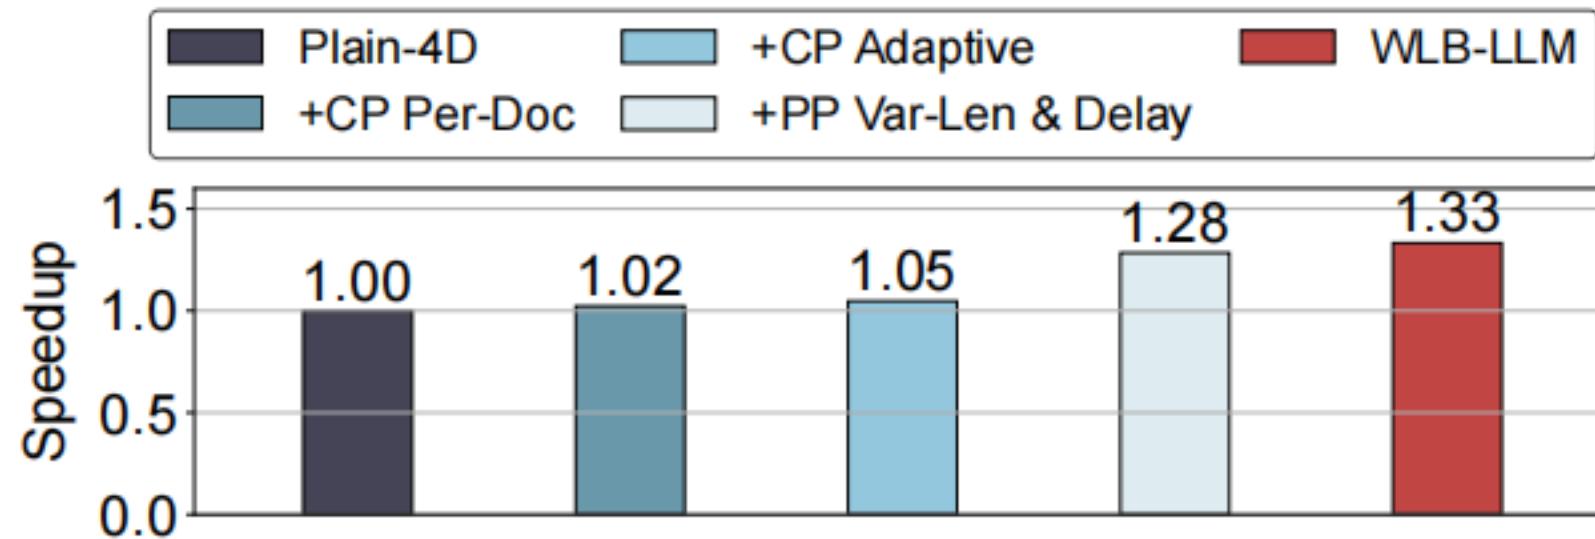
Baseline

- plain-4D:
 - no optimizing packing
 - per-seq sharding
- fixed-4D:
 - fixed-len packing
 - better sharding

Evaluation



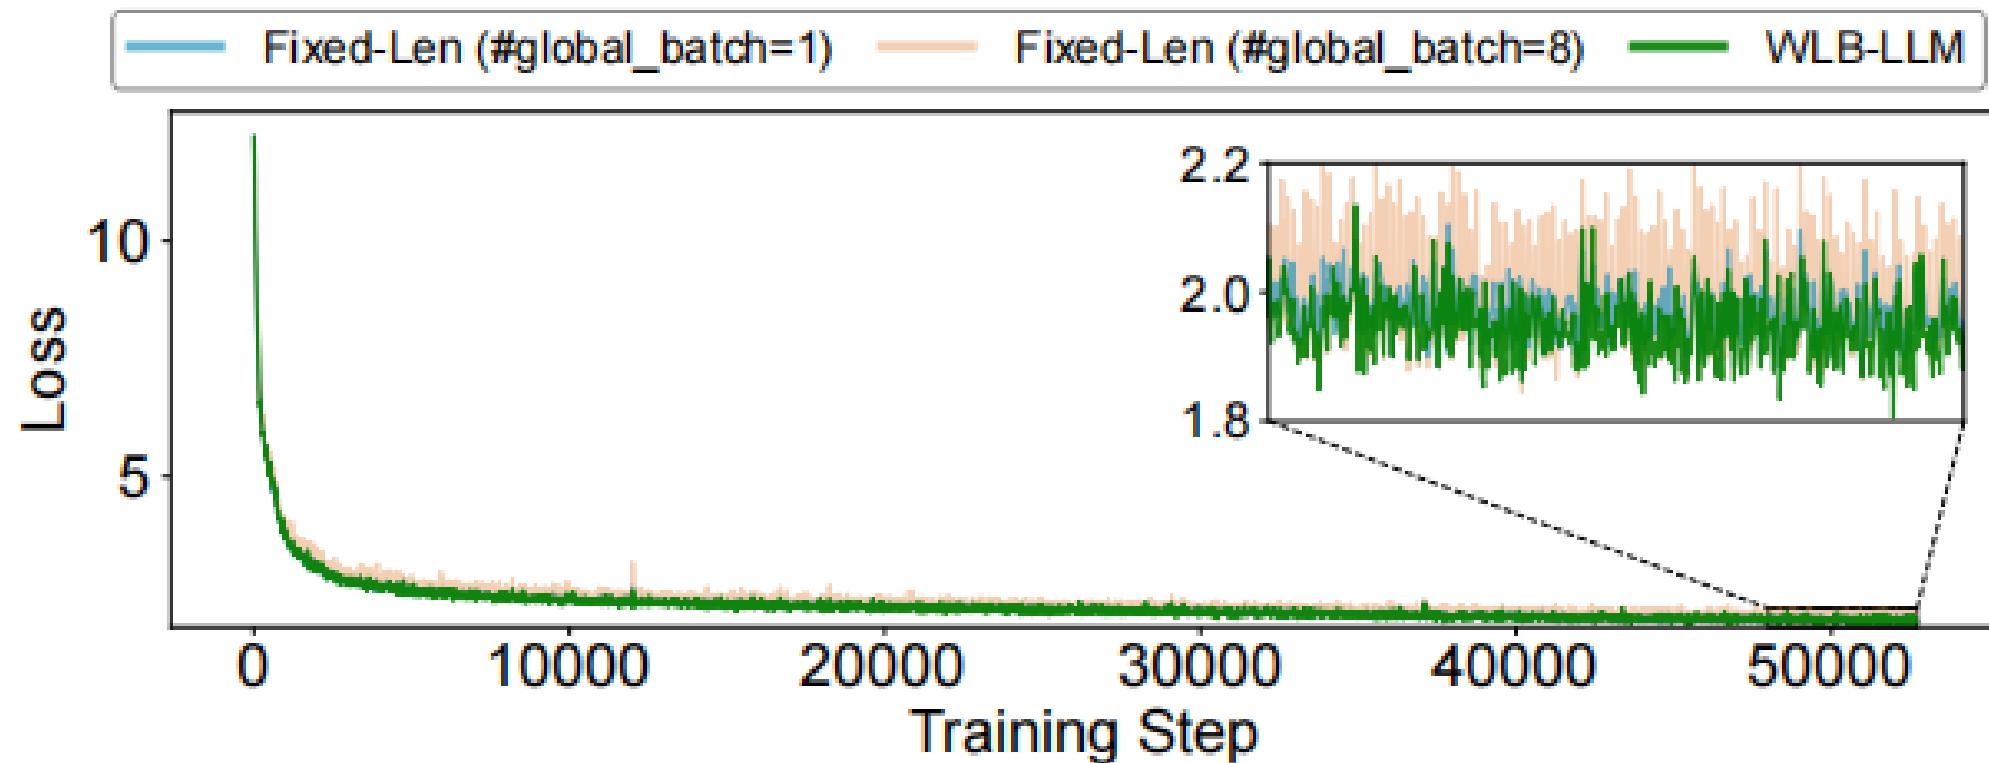
Evaluation: optimization analysis



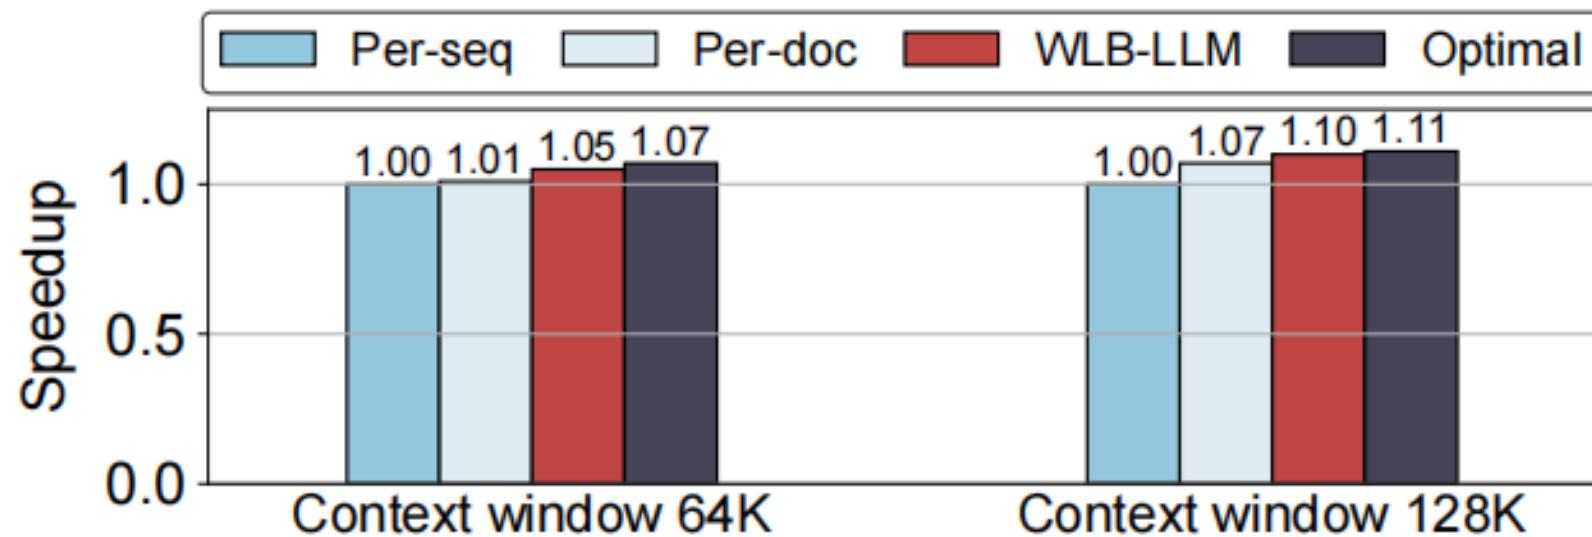
Evaluation: packing

Packing Method		Imbalance Degree	Packing Overhead (ms)
Method	Config		
<i>Original Packing</i>	/	1.44	0
<i>Fixed-Len Greedy</i>	#global batch=1	1.41	4
	#global batch=2	1.22	5
	#global batch=4	1.11	5
	#global batch=8	1.08	5
<i>Fixed-Len Solver</i>	#global batch=1	1.40	467
	#global batch=2	1.18	1488
	#global batch=4	1.09	25313
<i>WLB-LLM</i>	#queue=1	1.24	8
	#queue=2	1.05	20
	#queue=3	1.05	23

Evaluation: model convergence



Evaluation: sharding



Thinking

- 文章有什么疑问
在考虑本文方法对模型效果的影响时，其基线对比的是用固定长度 packing方法的效果，而不是和不加任何优化的模型训练作对比，其影响程度未知
- 在此之上：
文中的packing算法的针对异常文档的队列的长度阈值是预设的，是否可以根据数据的分布自适应的确定长度阈值
- 于此平行：
本文是优化训练过程中的负载不均衡提升性能，对于推理服务系统能否使用相同方法，优化推理延迟

Thanks