



東南大學  
SOUTHEAST UNIVERSITY



计算机科学与工程学院  
School of computer science and engineering

# Pruner: A Draft-then-Verify Exploration Mechanism to Accelerate Tensor Program Tuning

## 一种通过草稿-验证探索方式实现加速张量程序调优的机制

Liang Qiao\*  
University of Science and  
Technology of China  
Hefei, China  
ql1an9@mail.ustc.edu.cn

Jun Shi\*  
University of Science and  
Technology of China  
Hefei, China  
shijun18@mail.ustc.edu.cn

Xiaoyu Hao  
University of Science and  
Technology of China  
Hefei, China  
hxy2018@mail.ustc.edu.cn

Xi Fang  
University of Science and  
Technology of China  
Hefei, China  
fangxi@mail.ustc.edu.cn

Sen Zhang  
University of Science and  
Technology of China  
Hefei, China  
sen@mail.ustc.edu.cn

Minfan Zhao  
University of Science and  
Technology of China  
Hefei, China  
zmf@mail.ustc.edu.cn

Ziqi Zhu  
University of Science and  
Technology of China  
Hefei, China  
ta1ly@mail.ustc.edu.cn

Junshi Chen†  
University of Science and  
Technology of China  
Hefei, China  
cjuns@ustc.edu.cn

Hong An‡†  
University of Science and  
Technology of China  
Hefei, China  
han@ustc.edu.cn

Xulong Tang  
University of Pittsburgh  
Pittsburgh, USA  
tax6@pitt.edu

Bing Li  
NIO  
Shanghai, China  
libing475211023@sjtu.edu.cn

Honghui Yuan  
Xinyang Wang  
ethan.song@nio.com  
12307130155@fudan.edu.cn  
NIO  
Shanghai, China

<https://github.com/qiaolian9/Pruner>

ASPLOS'25

汇报人：王维龙 2025年11月24日



**安虹**, 博士, 教授, 博导, 省教学名师



**石军**, 中国科学技术大学博士后研究员

[1] PWDFT-SW: Extending the Limit of Plane-Wave DFT Calculations to 16K Atoms on the New Sunway Supercomputer (**TPDS 2025**) CCF A

[2] Pruner: A Draft-then-Verify Exploration Mechanism to Accelerate Tensor Program Tuning (**ASPLOS 2025**) CCF A

[3] NDFT: Accelerating Density Functional Theory Calculations via Hardware/Software Co-Design on Near-Data Computing System (**DAC 2025**) CCF A

[4] Predictive Accuracy-Based Active Learning for Medical Image Segmentation (**IJCAI 2024**) CCF A


# 提纲

- 1 研究背景
- 2 相关工作
- 3 研究内容
- 4 实验评估


# 提纲

- 1 研究背景
- 2 相关工作
- 3 研究内容
- 4 实验评估

# AI模型的应用周期




# AI模型的应用周期




1. 直接部署
2. 编译优化后部署


# AI模型的应用周期




1. 直接部署
2. 编译优化后部署





# AI模型的编译优化流程




1. 直接部署
2. 编译优化后部署




## 例子：GPU中的矩阵计算（通用矩阵乘GEMM-ReLU）




# 例子：GPU中的矩阵计算（通用矩阵乘GEMM-ReLU）



# AI模型的编译优化流程




# AI模型的编译优化流程



# 提纲


- 1 研究背景
- 2 相关工作
- 3 研究内容
- 4 实验评估

# L1: 调度方案搜索成本高



- [1] Zheng L, Liu R, Shao J, et al. Tensem: A large-scale program performance dataset for learned tensor compilers[C]//Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1). 2021.
- [2] Zheng L, Jia C, Sun M, et al. Ansor: Generating {High-Performance} tensor programs for deep learning[C]//14th USENIX symposium on operating systems design and implementation (OSDI 20). 2020: 863-879.
- [3] Zhai Y, Zhang Y, Liu S, et al. Tlp: A deep learning-based cost model for tensor program tuning[C]//Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2. 2023: 833-845.
- [4] Zhao Y, Sharif H, Adve V, et al. Felix: Optimizing tensor programs with gradient descent[C]//Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3. 2024: 367-381.

# L1: 调度方案搜索成本高




- [1] Zheng L, Liu R, Shao J, et al. Tensem: A large-scale program performance dataset for learned tensor compilers[C]//Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1). 2021.
- [2] Zheng L, Jia C, Sun M, et al. Ansor: Generating {High-Performance} tensor programs for deep learning[C]//14th USENIX symposium on operating systems design and implementation (OSDI 20). 2020: 863-879.

局限性1: 基于深度学习模型的成本模型, 导致探索成本高昂

- [4] Zhao Y, Sharif H, Adve V, et al. Felix: Optimizing tensor programs with gradient descent[C]//Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3. 2024: 367-381.


## L2: 不同硬件设备的方案难迁移



成本模型


V100

## L2: 不同硬件设备的方案难迁移



成本模型

V100



Jetson Orin Nano

局限性2: 不同硬件设备的方案间迁移效果不佳，甚至能力崩溃


# O1：探索机制可以更加高效

**Table 1.** Tuning costs (min) for Ansor with 2,000 trials on Orin, which means the time cost for space exploration, model training, and kernel hardware measurement, respectively.


Ansor	R50 [18]	DeTR [8]	I-V3 [35]
Exploration	35	30.31	41.8
Training	5.4	5.6	5.5
Measurement	44.4	50.61	49.4

发现：使用学习成本模型进行空间探索所占时间接近40%。当应用更复杂的成本模型时，所占时间会增加。

## O2：构建成本模型困难




## O2：构建成本模型困难




发现：离线训练的成本模型迁移困难，在线的成本模型需要额外微调甚至重训练

- 1 研究背景
- 2 相关工作
- 3 研究内容
- 4 实验评估

# Overview




# Overview



TVM中Ansan  
自动或集成的  
工具实现

# Draft: Latent Schedule Explorer




# Hardware-aware Symbols

Mem	Symbol
L0	S1-L0MemAlloc, S2-L0CompCount
L1	S3-L1MemAlloc, S4-L1ParaInfo
L2	S5-L2MemFootprint, S6-L2ParaInfo S7-L2TransDim, S8-L2CompCount

L0 (2个) : 寄存器的“S1内存占用 + S2计算量”;  
L1 (2个) : 共享内存的“S3内存占用 + S4并行度”;  
L2 (4个) : 全局内存的“S5访问量 + S6并行度 + S7传输维度 + S8总计算量”。

# Hardware-aware Symbols

Mem	Symbol
L0	S1-L0MemAlloc, S2-L0CompCount
L1	S3-L1MemAlloc, S4-L1ParaInfo
L2	S5-L2MemFootprint, S6-L2ParaInfo S7-L2TransDim, S8-L2CompCount



**Schedule Template**

**Schedule Generation Rules**

- Split**(loop:i, extent:128, into:[I0, I1, I2, I3, I4])
- Split**(loop:j, extent:128, into:[J0, J1, J2, J3, J4])
- Split**(loop:k, extent:128, into:[K0, K1, K2])
- Reorder**(0,5,1,6,2,7,10,11,3,8,12,4,9)
- CacheRead, ComputeAt, Annotation,
- .....

**Transformed Program:**

```

block_parallel I_J_0 in (0, I0*I0):
thread_parallel I_J_1 in (0, I1*I1):
  for v in (0, V): // vthread
    for k0 in (0, K0):
      A.shared = A;
      B.shared = B;
      for k1, i3, j3 in grid(K1, I3, J3):
        for k2, i4, j4 in grid(K2, I4, J4):
          C.local += A.shared * B.shared;
      for i5, j5 in grid(I3*I4, J3*J4):
        D = max(C, 0.0)
  
```

Hardware-aware Symbol
<b>Symbol 3</b>
Prod(L1_A, [I1,...,I4,K1,K2])
Prod(L1_B, [J1,...,J4,K1,K2])
Sum(L1MemAlloc, [L1_A, L1_B])
<b>Symbol 5</b>
Prod(L2_A_traffic, [I0,...,I4,J0,K0,...,K2])
Prod(L2_B_traffic, [I0,J0,...,J4,K0,...,K2])
Sum(L2MemFootprint, [L2_A_traffic, L2_B_traffic])
<b>Symbol 4</b>
Prod(L1ParaInfo, [I1, J1])
<b>Symbol 1</b>
Prod(L0_C, [I2,...,I4,J2,...,J4])
Prod(L0_A, [I2,...,I4])
Prod(L0_B, [J2,...,J4])
Sum(L0MemAlloc, [L0_C, L0_A, L0_B])
<b>Symbol 2</b>
Prod(L0CompC, [I2,...,I4,J2,...,J4,K0,...,K2])
<b>Symbol *</b>

全局扫描  
然后生成  
Symbols表示

# Symbol-based Analyzer

本质是一个轻量级经验公式代价模型

计算相关延迟

$$L_c^i = \frac{S8}{U_p}, \quad L_m^i = \frac{S5}{U_m}, \quad L_{total} = \sum_i (L_c^i + L_m^i) \quad (1)$$

内存相关延迟

# Symbol-based Analyzer

本质是一个**轻量级经验公式代价模型**

计算相关延迟

$$L_c^i = \frac{S8}{U_p}, \quad L_m^i = \frac{S5}{U_m}, \quad L_{total} = \sum_i (L_c^i + L_m^i) \quad (1)$$

内存相关延迟


数学公式简单计算

深度学习模型


高效筛选有  
潜力候选集  
(Schedule)

低效筛选有  
潜力候选集  
(Schedule)

# Verify: Pattern-aware Cost Model



# Temporal Dataflow Pattern




平铺 (tiling) 表征数据的流动过程 (时序) , 解决传统特征忽略数据流依赖的问题

Statement-level features  
(语句级特征) 是传统张量程序代价模型重点关注的特征类型

# Pattern-aware Transformer

**Transformed Program:**

```
block_parallel I_J_0 in (0, I0*J0):
    thread_parallel I_J_1 in (0, I1*J1):
        for v in (0, V): // vthread
            C.local = 0.0;
            for k0 in (0, K0):
                A.shared = A;
                B.shared = B;
                for k1, i3, j3 in grid(K1, I3, J3):
                    for k2, i4, j4 in grid(K2, I4, J4):
                        C.local += A.shared * B.shared;
                for i5, j5 in grid(I3*I4, J3*J4):
                    D = max(C.local, 0.0);
```



**Multi-Tiling Pattern**

```
#1: C.local = 0.0
#2 & 3: A.shared, B.shared ← A, B // L2 to L1
#4: C.local ←
    Compute(A.shared, B.shared) // L1 to L0
#5: D ← Elemwise(C.local) // L0 to L2
```

**Temporal Data Flow Feature**


```
(compute:1|mem access:21| alloc size:1)
C.local; A.global; ...
C.local: (2, 128, 128, 32, ...)
....;
D.globa: (0, 1024, ...) // Dim(10,23)
```

Statement-level features  
(语句级特征) 是传统张量程序代价模型  
重点关注的特征类型


任务的核心不是预测的延迟更准，而是排序/相对性能更准

以归一化延迟为目标、LambdaRank 为损失函数，通过线性层输出相对性能评分


# Verify: Pattern-aware Cost Model



# Momentum online Adaptation (MoA)-Pruner




# Momentum online Adaptation (MoA)-Pruner



相似架构的预训练的成本模型

1. Load Param: 是为了成本模型的初始权重不是随机的，而是具备一定知识。（可以类似CV领域的使用随机初始权重，和使用ImageNet预训练权重）


# Momentum online Adaptation (MoA)-Pruner



相似架构的预训练的成本模型

1. **Load Param**: 是为了成本模型的初始权重不是随机的，而是具备一定知识。（可以类似CV领域的使用随机初始权重，和使用ImageNet预训练权重）
2. **Online**: 与传统方案一致，采样+验证+微调

# Momentum online Adaptation (MoA)-Pruner



相似架构的预训练的成本模型

1. **Load Param**: 是为了成本模型的初始权重不是随机的，而是具备一定知识。（可以类似CV领域的使用随机初始权重，和使用ImageNet预训练权重）
2. **Online**: 与传统方案一致，采样+验证+微调
3. **Momentum update**: 反向更新Siamese Model，使其更具通用性，为后续其他设备迁移提供经验

- 1 研究背景
- 2 相关工作
- 3 研究内容
- 4 实验评估

# 实验设置

## • DNN workloads

**Table 3.** Evaluated DNN models in Pruner, with shapes and optimization precisions.

CNNs	Shape & Precision	Transformers	Shape & Precision
ResNet[18]	(1, 3, 224, 224) & (F)	Bert-B/T[15]	(1 & 4, 128) & (F, H)
WideResNet[43]	(1, 3, 224, 224) & (F)	GPT-2[30]	(1 & 4, 128) & (F, H)
Inception-V3[35]	(1, 3, 299, 299) & (F)	Llama[39]	(1 & 4, 128) & (F, H)
Densenet-121[19]	(1, 3, 224, 224) & (F)	OPT[46]	(1 & 4, 128) & (H)
Mobilenet-V2[32]	(1, 3, 224, 224) & (F)	Mistral[20]	(1 & 4, 128) & (H)
DCGAN[29]	(1, 100) & (F)	ViT[16]	(1, 3, 256, 256) & (F)
Deeplab-V3[9]	(1, 3, 224, 224) & (F)	DeTR[18]	(1, 3, 256, 256) & (F)

**Table 4.** Details for Transformer-based language models.

Model	layers	heads	hidden	intermediate
Bert-Tiny	6	8	512	2048
Bert-Base	12	12	768	3072
GPT-2	12	12	768	3072
OPT-1.3b	24	32	2048	8192
Llama	12	12	768	3072
Mistral-7b	32	32	4096	14336


## • Platforms and baselines

- A100 (服务器)
- Titan V (服务器)
- Jetson Orin-AGX (边缘设备)

## • Baselines

- 工业界: Ansol、MetaSchedule
- 学术界: TenSetMLP [NeurIPS'21]、TLP [ASPLOS'23]、Felix [ASPLOS'24]、Adatune [NeurIPS'20]、TLM [OSDI'24]。

# 最优推理延迟



所需方案推理延迟低

(Draft: Latent Schedule Explorer,

Verify: Pattern-aware Cost Model)

在线方案显著优于离线方案

设备适配能力强

(Momentum online Adaptation (MoA) Pruner)

Pruner


TenSetMLP

TLP

Pruner offline

横轴是搜索时间（基于在线/离线训练的成本模型），纵轴是筛选的最优方案的推理延迟

# 搜索效率




搜索时间显著减少

(Draft: Latent Schedule Explorer,  
Verify: Pattern-aware Cost Model)

横轴是不同模型，纵轴是搜索时间

# 最优推理延迟



横轴是不同模型，纵轴是归一化后的模型推理性能

具备更强的泛化性，不是着眼特定算子，而是重点关注设备特征计算  
**(Momentum online Adaptation (MoA)-Pruner)**

# 总结与思考

- **框架**: 提出一种新型的张量程序调度方案探索机制，即采用先草拟后验证”的范式加速搜索过程。
- **技术**: 引入了动态在线适应技术，解决了预训练成本模型在跨平台中难迁移的问题，使得在线成本模型能够高效适应任何平台。
- **实验**: 通过与现有最先进方法在三个基于GPU的平台上（包括CUDA核心和Tensor核心）的对比分析，突显了Pruner的普适性。全面的实验表明，Pruner在调优质量和效率之间取得了显著的平衡，其性能远超现有方法。并认为Pruner背后的核心思想与现有的基于搜索的方法相辅相成。

# 思考

## 1. 这个paper有什么问题，基于这个paper还能做什么？

- 本文重点关注GPU，重点关注了latency，是否可以将问题聚焦与边缘计算场景，并且还可以考虑energy，考虑多目标优化探索。

## 2. 这个paper提到的idea，能不能用在自己的方向/project上面？

- 本文的方案是一种特别适合和硬件设备相关的，以性能为导向的探索类问题。其理论上也可以用于神经网络超网搜索最优子网的问题，至少可以用于筛选部分探索候选集（无需真实推理即可排除不满足延迟约束的候选网络）。

## 3. 这个paper能不能泛化，需要较为熟悉这个小方向？

- 本文的面向的设备是GPU。从理论上可知，只要可以抽象出不同设备的共有特征，或者采用双重特征（1重特征表征什么类型设备，2重特征表征具体计算能力特征），是不是就可以泛化到更多的设备场景，甚至实现GPU→CPU, TPU等。



東南大學  
SOUTHEAST UNIVERSITY



计算机科学与工程学院  
School of computer science and engineering

感谢各位老师和同学!  
请大家提出宝贵意见!